



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# Relational Database Systems I

**Wolf-Tilo Balke**

**Simon Barthel**

Institut für Informationssysteme  
Technische Universität Braunschweig  
[www.ifis.cs.tu-bs.de](http://www.ifis.cs.tu-bs.de)



# 4 Overview


- View integration
- Resolving conceptual incompatibility
- Entity clustering for ER models
- Commercial dimension:  
The BEA story





# 4.1 Business Integration


- Business currently is a world of M&A
  - Companies need to diversify/enhance their portfolio
  - But it is expensive to develop necessary applications
    - Knowledge gathering costs time
    - Will the output be worth it?
  - Idea: Rely on people who **are already** knowledgeable
    - Acquire small, specialized, and promising companies
    - Merge with big players in the field





# 4.1 Business Integration

- Examples
  - The Daimler–Chrysler merger
  - The Oracle–Sun merger
  - Oracle buys PeopleSoft, Siebel Systems, BEA Systems, ...
    - Siebel Sales as CRM tool now part of Oracle's business intelligence suite





## 4.1 Business Integration

- Merged (parts of) businesses are administrated by
  - Different specialized software systems?
  - One company-wide system?
- Usually there is an **historical evolution** of separate tools and programs
  - Accounting, sales & marketing, development, ...
  - Based on individual requirements
- However, often a **unified view** is needed
  - Business intelligence? Warehousing...?!
    - Warehousing is also a **great lecture** at ifis...



## 4.1 View Integration

- Usually, there are **several conceptual schemas**
  - **Several designers** are part of the modeling process (modular software development)
  - **Different tasks** were modeled within the same organization (legacy systems)
  - **Several organizations** need to be integrated (business integration)







# 4.1 View Integration

- **View Integration**
  - Several conceptual schemas need to be combined into a unified global schema
  - All differences in perspective and terminology have to be resolved
  - All redundancy has to be removed
- **But,... what happens, if you don't integrate?!**






# 4.1 The Mars Desaster

*Detour*

- Example: Big NASA project **Mars Surveyor**
  - The 1998 missions investigated “Volatiles and Climate History” on Mars
    - Characterization of climate change and its evolving impact on the distribution of water
    - Idea: explore the **polar ice caps** of Mars and see whether there is water ice
    - Planned were about three months of experiments on Mars









## 4.1 The Mars Desaster

*Detour*

- Two vehicles:

### **Mars climate orbiter** and **Mars polar lander**

- The lander was supposed to probe the layers of ice and dust on the polar ice caps to investigate changes
- The orbiter was built to monitor the daily weather and record changes in water vapor and dust in the atmosphere





## 4.1 The Mars Desaster

*Detour*

- Catastrophic failure
  - The Mars climate orbiter approached Mars up to 57 km instead of 150 km, and was **destroyed** in the atmosphere on September 23, 1999
  - The Mars polar lander **crashed** during its attempted landing on Mars, December 3, 1999
  - **\$327.6 million** in total for both
    - \$193.1 million for spacecraft development
    - \$ 91.7 million for launch
    - \$ 42.8 million for mission operations






# 4.1 The Mars Desaster

*Detour*

- Why did the **climate orbiter** come too close to Mars' atmosphere?
  - **Many organizations** were involved in the development
  - There was no global schema
    - Navigation software produced by Lockheed Martin used **non-metric** units (e.g., inches, feet, and pounds)
    - NASA uses **metric** units
    - A small correction of the course led to the fatally low orbit...






## 4.1 The Mars Desaster

*Detour*

- Happy ending!
  - The next try was the successful **2001 Mars Odyssey**
  - The measurements pointed to **water ice on Mars**
  - Confirmed by the Mars Express (ESA) in 2004
    - The polar caps consist of 85% carbon dioxide (CO<sub>2</sub>) ice and 15% water ice





## 4.2 Resolving Incompatibilities

- **Schema diversity** occurs when different users develop their own understanding of the world
  - The same reality is not always modeled in the same way due to different information needs or workflows
- **Common principle** for schema integration
  - Identify the parts of the input schemas that represent the same reality
  - Unify their representations








## 4.2 Basic Steps

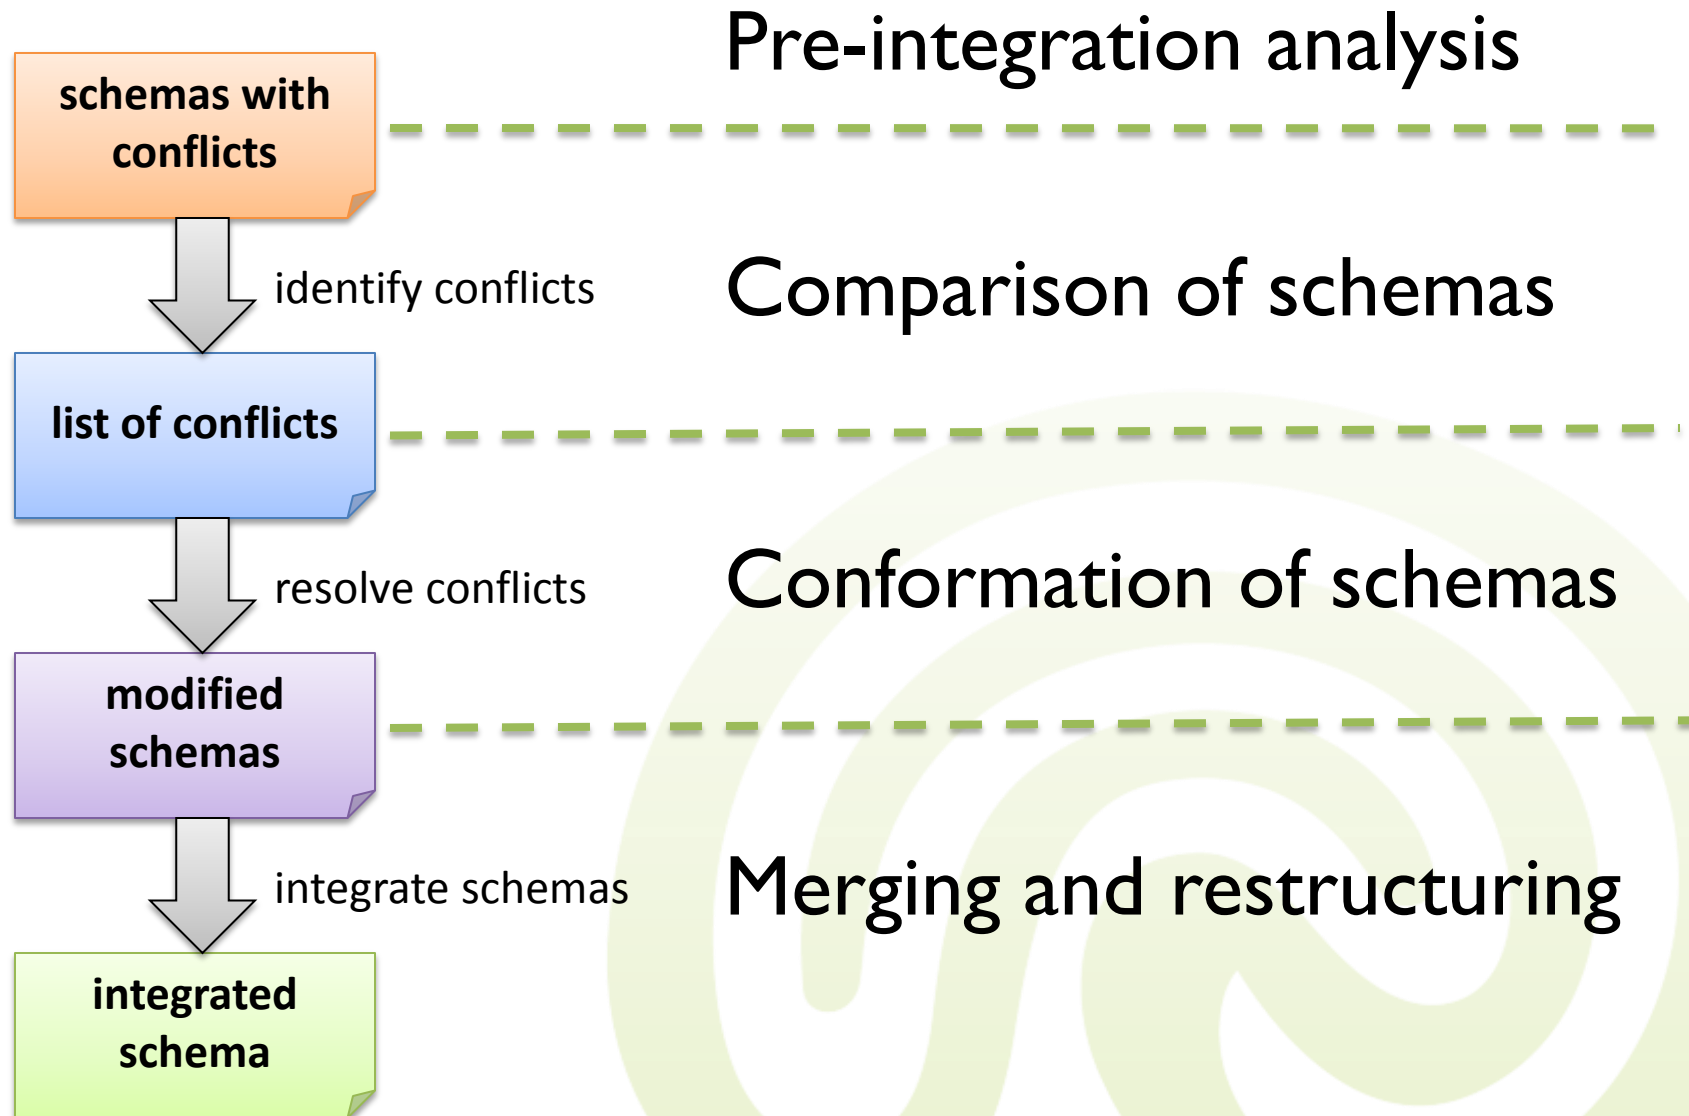
- There are **four basic steps** needed for conceptual schema integration
  1. Pre-integration analysis
  2. Comparison of schemas
  3. Conformation of schemas
  4. Merging and restructuring of schemas
- The integration process needs continual refinement and reevaluation







## 4.2 Schematic View





## 4.2 Pre-integration Analysis

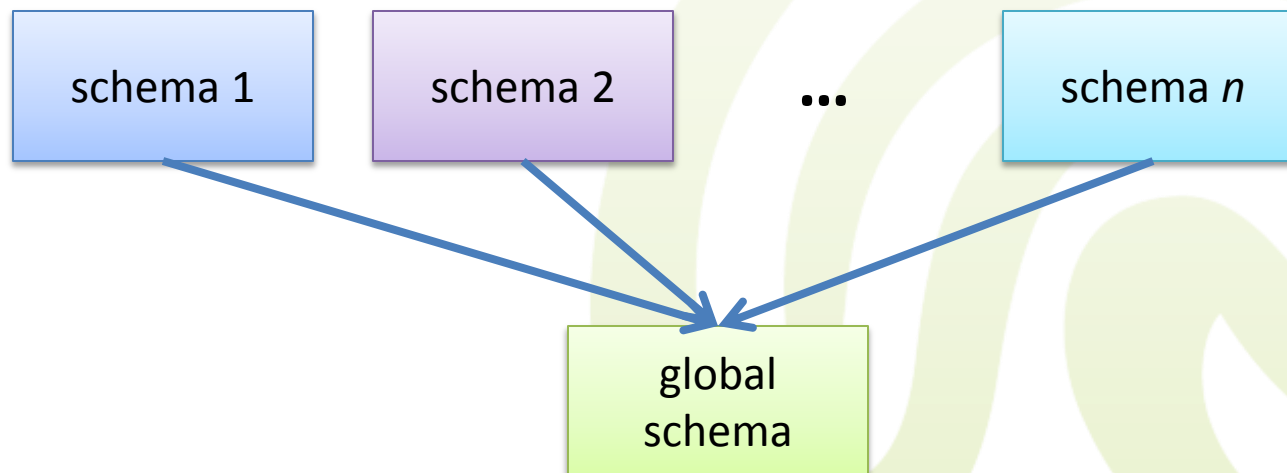
- **Pre-integration analysis** takes a close look on the individual conceptual schemas to decide for an **adequate integration strategy**
  - The larger the number of constructs, the more important is modularization
  - Is it really sensible/possible to integrate all schemas?





## 4.2 Pre-integration Analysis

- First, an **integration strategy** has to be chosen
- Schema integration can either be performed **many at a time**, ...
  - Requires only one consistent merge
  - Conflict analysis from many schemas is difficult

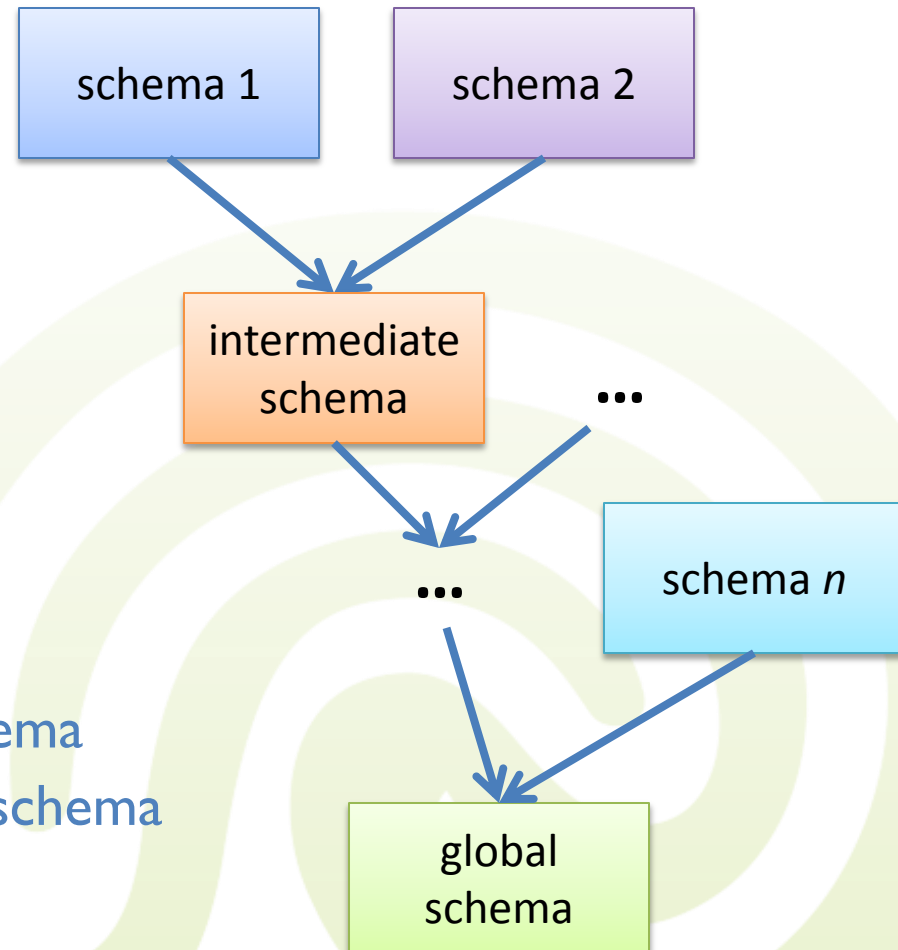




## 4.2 Pre-integration Analysis

- ... or can be performed **two at a time**

- Results are step by step accumulated into a single schema
- How to choose **the right order** of schema comparisons
  - The order can influence the final result
- Selecting the first schema
  - Mixed strategy: skeleton schema
  - Otherwise: most important schema





## 4.2 Comparison of Schemas

- The **resolution of conflicts** needs a thorough comparison of schemas
  - General question: How do entities **correspond**?
  - **Naming conflicts** can be detected, e.g., by scanning the data dictionary
  - **Structural conflicts** regarding semantics have to be resolved:
    - Different cardinalities in relationships
    - Key conflicts
    - ...





## 4.2 Comparison of Schemas

- The individual **perspective of the world** and the **level of abstraction** are major reasons for conceptual incompatibilities
  - Example: A **product** is a **unit of sale** for the marketing department, but consists of **parts** in the view of the engineering department







## 4.2 Comparison of Schemas


- The **level of abstraction** directly influences the schema design
- **Simple example:** A customer buys a product
- The **marketing view** focuses on how many people buy some product, e.g., for advertising
  - Only the characteristics of the customer and product and the connection are needed





## 4.2 Comparison of Schemas


- The **accounting view** also needs the exact order number for identification of individual customer transactions
  - Focus is on the purchase, but individual orders have to be distinguished





## 4.2 Comparison of Schemas

- The **sales view** needs all individual order details, e.g., for troubleshooting or CRM
  - Focus is on orders  
(which provide the basis for purchase contracts)






## 4.2 Comparison of Schemas


- Marketing



- Accounting




- Sales





## 4.2 Comparison of Schemas

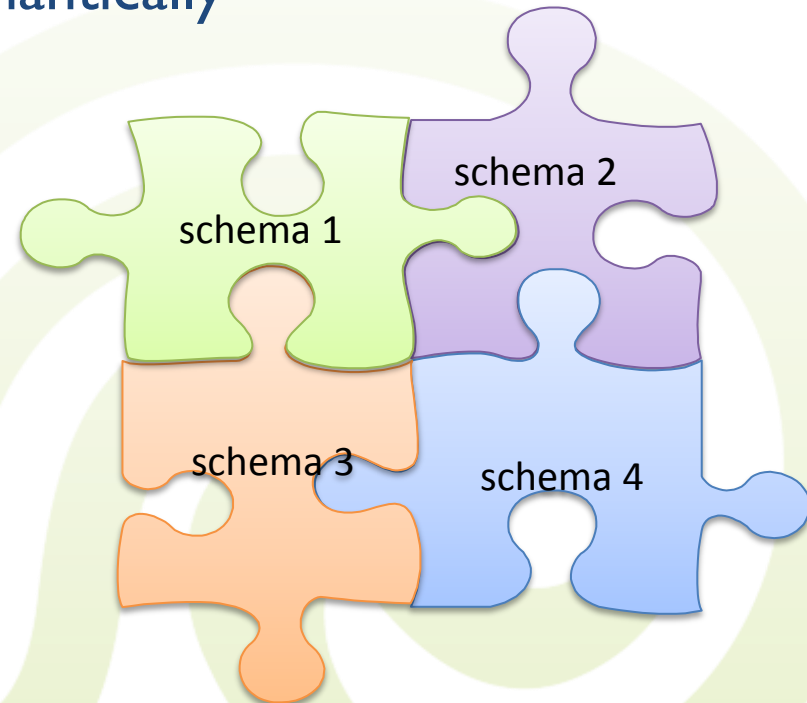
- Different user groups use different names to refer to the same entities (differing terminology)
  - **Synonyms:** Two terms for the same entity
  - **Homonym:** The same term for different entities
- **Rule of thumb:**  
eliminate synonyms,  
rename homonyms!





## 4.2 Conformation of Schemas

- The main goal is to make schemas **compatible** for integration
- Conformation usually needs **manual** interaction
  - Conflicts need to be resolved semantically
  - Rename entities/attributes
  - Convert differing types, e.g., convert an entity to an attribute or a relationship
  - Align cardinalities/functionalities
  - Align different data types







## 4.2 Conformation of Schemas

- Besides renaming and type conversions, **abstraction** can be useful
  - Generalization and aggregation allows to create new supertypes or subtypes
- Also **assertions and constraints** must be generalized or distributed among the type hierarchy
  - For example, checking accounts and saving accounts are both types of accounts, but may differ with respect to the minimum balance constraint



## 4.2 Conformation of Schemas

- **Example:** Resolving differing terminology
- **Homonyms:** Schema 1 and 2 both contain the term “jaguar,” but mean different entities
  - Rename to “**jaguar\_car**” and “**jaguar\_animal**”
- **Synonyms:** Schema 1 contains the term “jaguar,” whereas schema 2 contains the term “panther”
  - Global schema should model “**panther**” **is\_a** “**jaguar**”
  - Constraint on the black color should be added





## 4.2 Merging and Restructuring

- How to merge schemas into a **global schema**?
  - Copy all distinct **entities** from the individual schemas
  - Apply renaming, overlapping entity integration, abstraction, attribute type conversions, etc.
  - Put in the distinct **relationships** from all schemas
  - Again use renaming, cardinality/functionality conversions, etc.
  - Restructure the resulting global schema



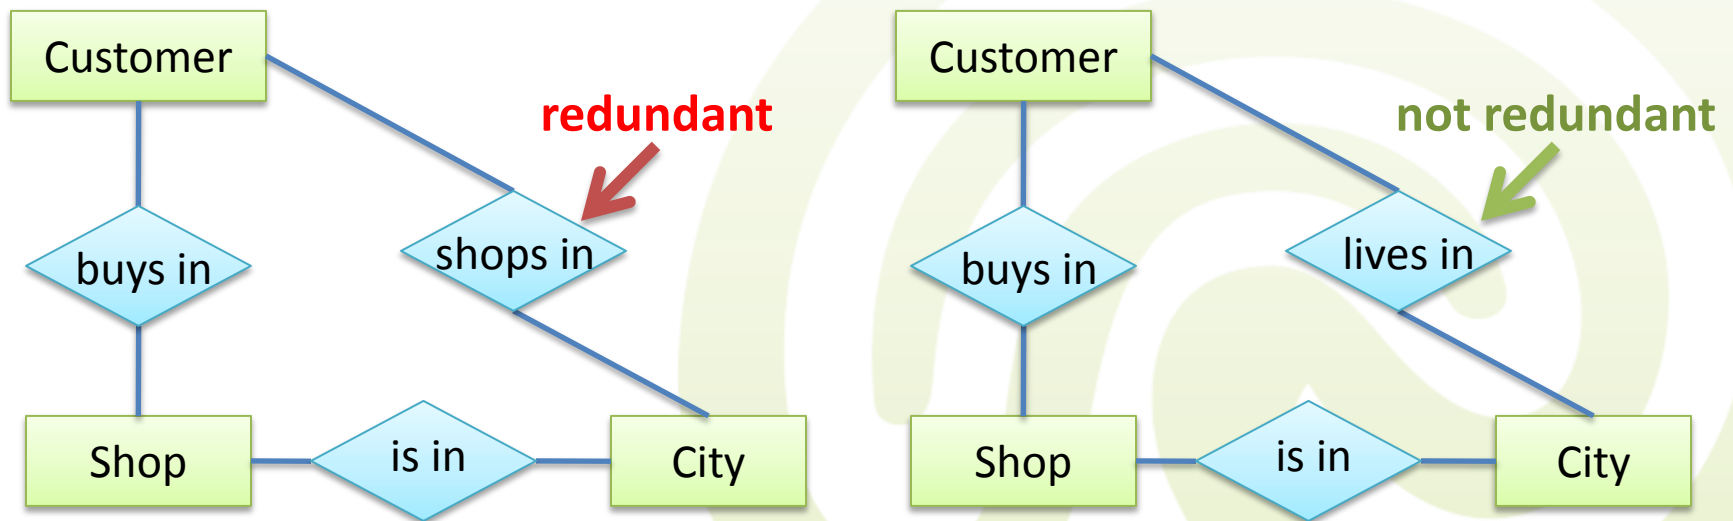
## 4.2 Merging and Restructuring

- The final **restructuring of the schema** is driven by the goal of completeness, minimality, and ease of understanding
  - **Completeness** mandates that all concepts in the global schema appear “semantically intact”
    - All different concepts of every individual schema are also part of the global schema
    - For each concept, there are no missing attributes, no constraints that cannot be met by all members of a type, etc.



## 4.2 Merging and Restructuring

- **Minimality** enforces to remove all redundant concepts from the global schema
  - For example, overlapping entities or redundant relationships
  - Often, the question of minimality can only be decided semantically





## 4.2 Merging and Restructuring

- **Ease of understanding** means that the global schema makes sense to the users
  - In particular, abstraction and fine granular levels for entities can be very confusing
  - Example: Subtype entities have to be clearly distinguishable, and should have only attributes that are not inherited from the supertype





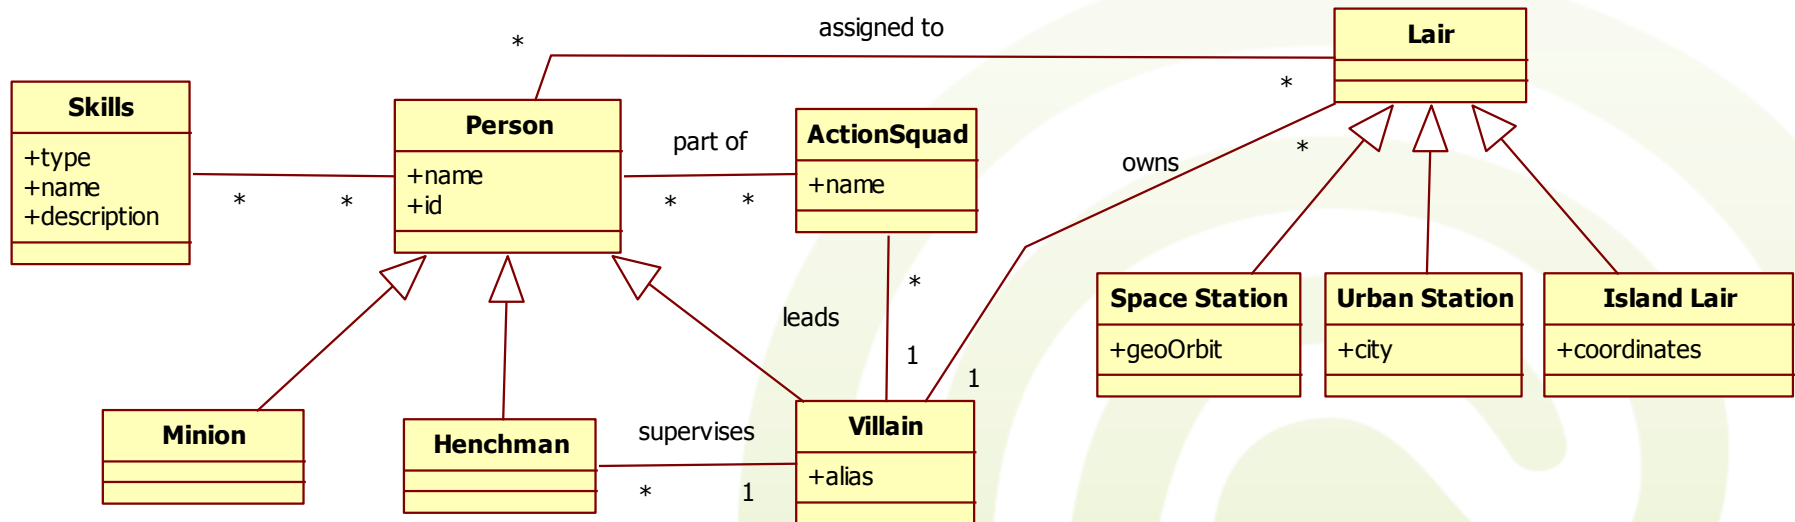


- **Doomsday Legion (DDL)**
  - Cooperation of villains from all over the world striving for global domination
  - Channeling resources, staff, experience and power for reaching their goals
  - Centralized and coordinated management of all shared assets
    - Lairs
    - Personnel
    - Assault squads
    - ...





- Doomsday Legion schema





- **Justice League (JL)**

- Federation of super-powered heroes fighting against global crime and villainousness

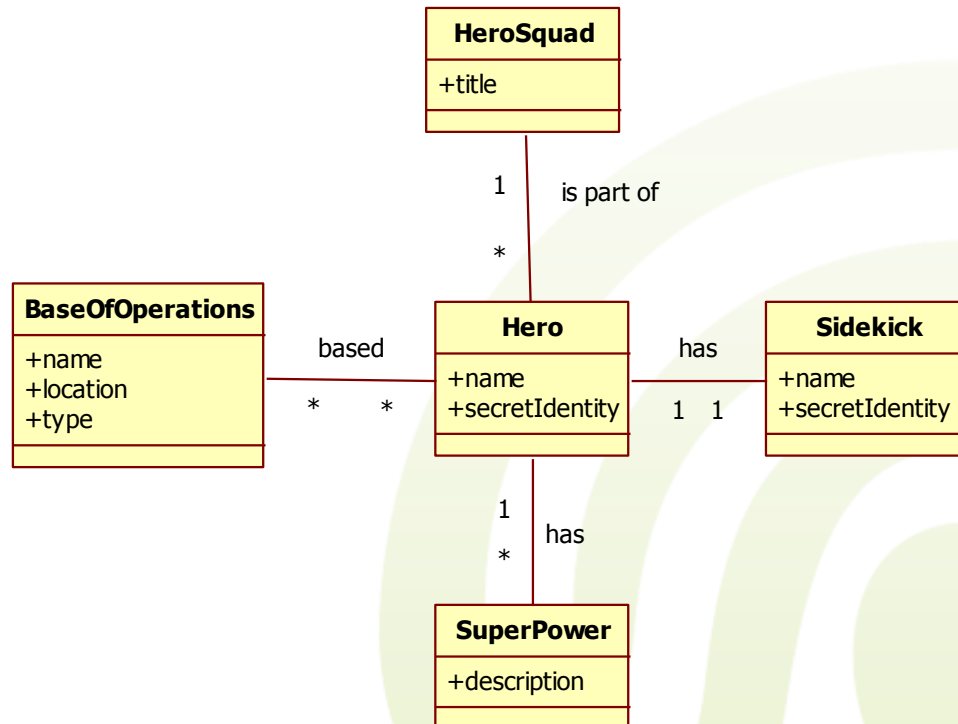
- In particular: Opposing the Doomsday Legion

- Central management of joint operations and resources





- Justice League schema



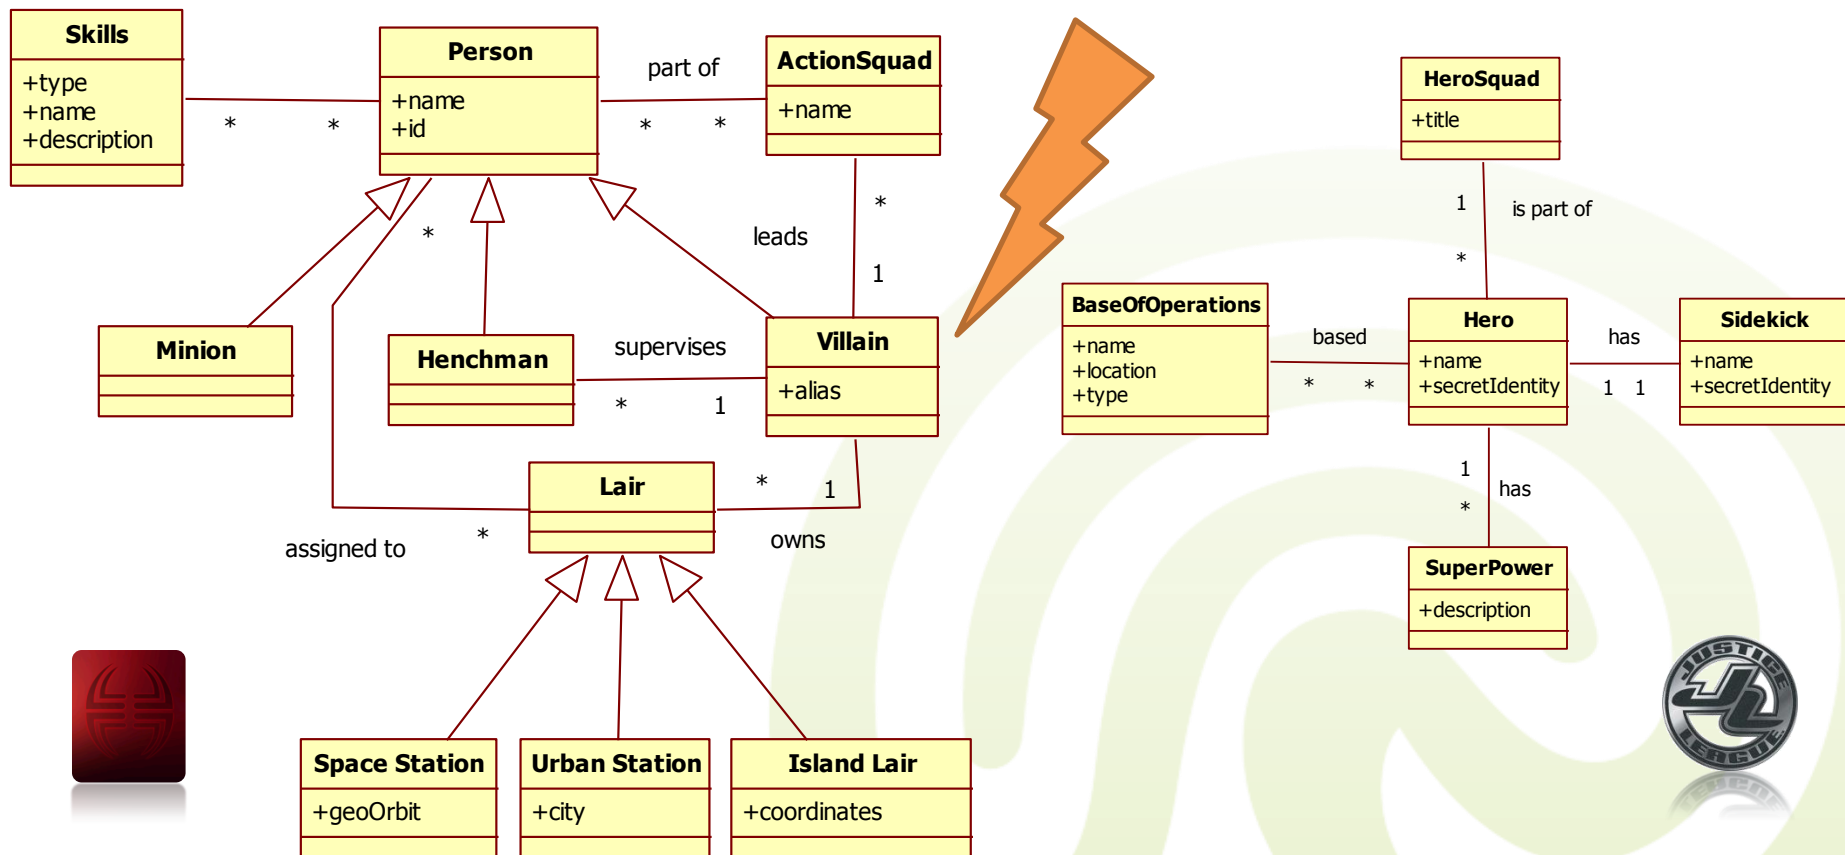


- “And then, strange and evil **aliens invade earth** without any obvious reason”
  - Justice League wants to save earth (that’s what they do)
  - Doomsday Legion wants to save earth (without people, global domination is no fun)
  - Great Idea: **Join Forces**
    - “**D**efenders of the **E**arth”
  - Great Problem: Joining large organizations is **not that easy**
    - Beside the problem of ignoring old hatred, the data **schemas need to be integrated** for central mission control and planning






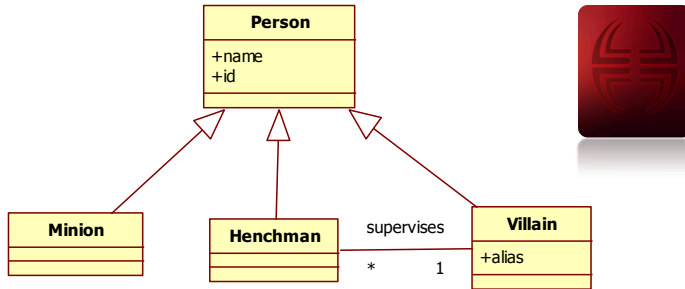
- How to integrate?







- Integrating the person models



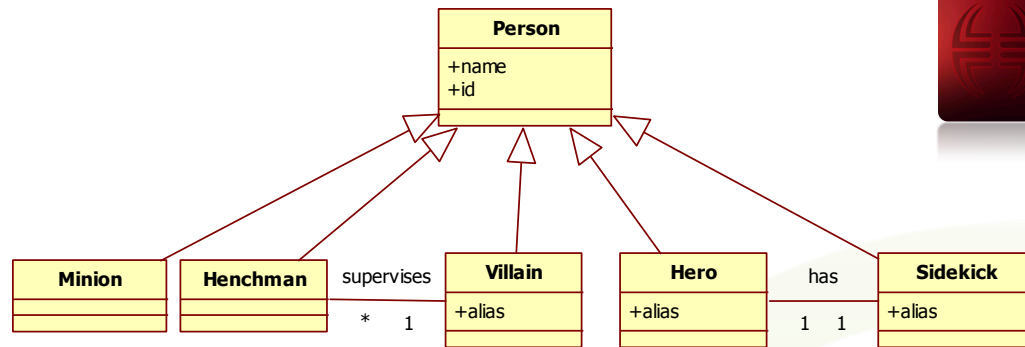
- Different structure

- DDL more general → Merge JL into DDL
  - Generalize Hero and Sidekick into Person
- But: **Attribute Homonyms!**
  - DDL uses the real name of “name,” the villain identity is “alias”
  - JL puts real name into “secret identity” and hero name into “name”
  - That is: (*name*: “Victor von Doom”, *alias* “Dr. Doom”)
  - (*name*: “Invisible Woman”, *secret identity* “Susan Storm”)
  - Attributes need to be renamed and transformed correctly





- Integrating the person models


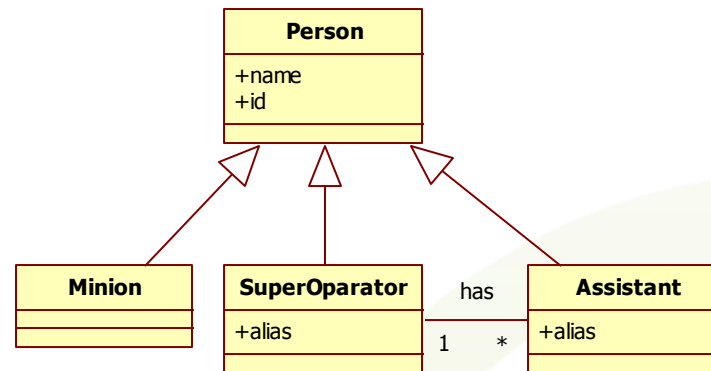


## – Semantic consolidation

- Hero and Villains are should be treated the same
  - Both are highly skilled and powerful super members of DotE
  - Merge classes into *SuperOperator* class
- Sidekicks and henchmen are close *assistants* of an operator
  - Heroes usually only have one sidekick
  - Use more general 1:n association to also capture henchmen



- Integrating the person models



- Contains heroes and villains, as well as their respective sidekicks or henchmen
- Heroes and sidekicks get an additional id



# Example: View Integration

# Detour

- Integrating bases



- Villains only have 3 types of bases, explicitly modeled
- Heroes may have any kind of base, given by the type attribute
- Two solutions

- Merge DDL into JL

- geoOrbit*, *city*, and *coordinates* become *location*
    - type* is given by subclass
    - Only possible in a lossless fashion because subtypes don't have additional attributes

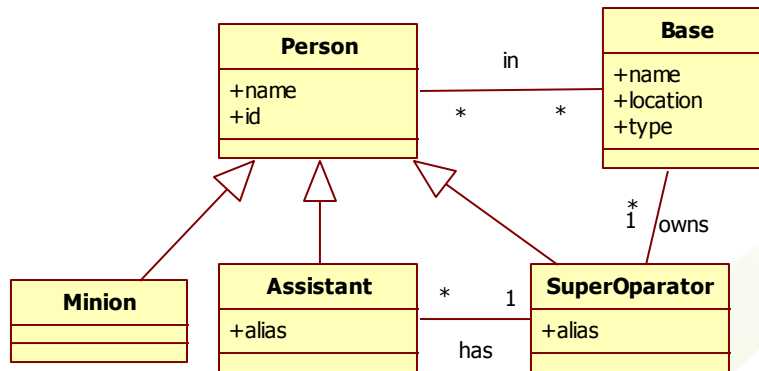
- Merge JL into DDL

- Depending on type, a base is assigned to one of the subclasses
    - 4<sup>th</sup> subclass necessary for all other types (could also be merged into superclass)





- Integrating bases



– Only villains owned lairs; no information of ownership for former hero bases



- Integrating skills and powers
  - JL only stores super powers
  - DDL stores all skills (including super powers)
    - More general
  - Merge JL into DDL
    - All old justice league super powers become skills of the *type* “super power”
    - *name* is either null or manually completed
    - No information on skills of sidekicks



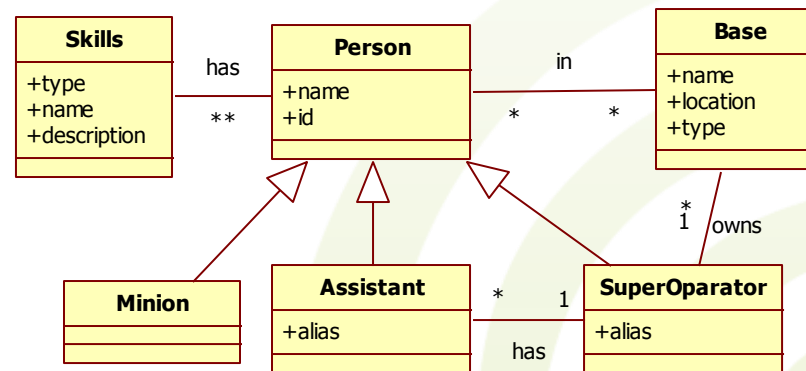
SuperPower
+description



Skills
+type
+name
+description



- Integrated schema





## 4.2 Outlook

- View integration is a **semantic process**
  - This usually means a lot of **manual work**
  - Computers can support the process by **matching** some (parts of) schemas
- There have been some approaches towards **(semi-)automatic matching** of schemas
  - Matching is a complex process and usually only focuses on simple constructs like  
“Are two entities semantically equivalent?”
  - The result is still rather error-prone...





## 4.2 Outlook

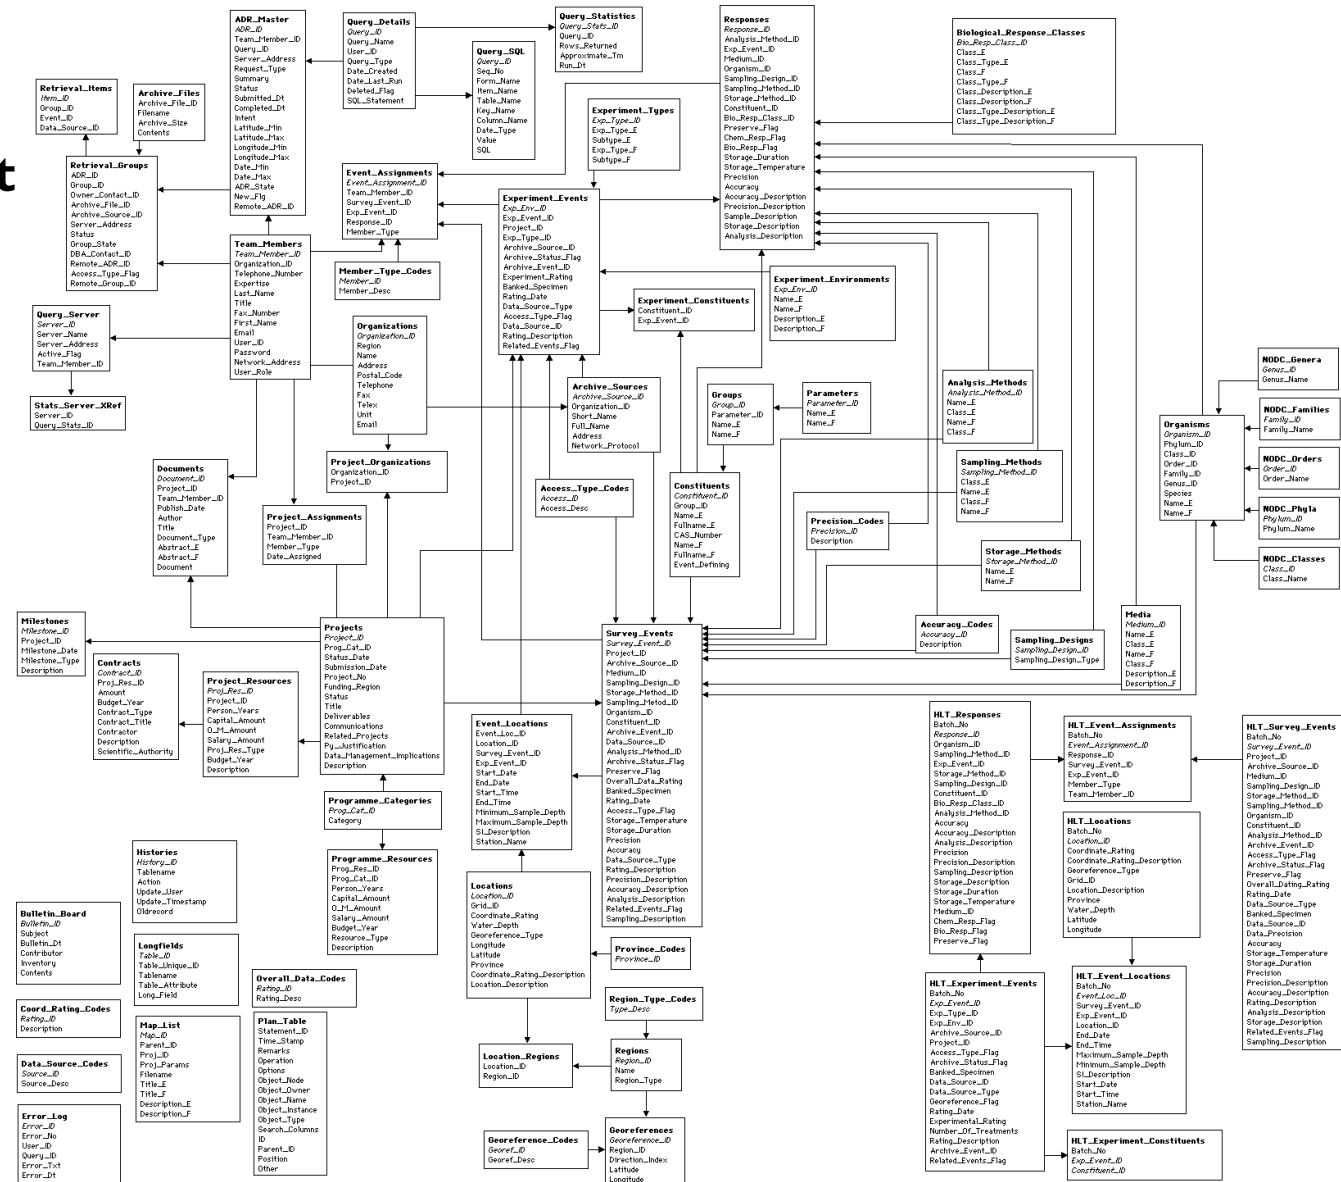
- **Basic methods** (that can of course be mixed freely)
  - **Label-based matching**
    - For each label in one schema, consider all labels of the other schema and every time gauge their semantic similarity
  - **Instance-based matching**
    - Looking at the instances (of entities or relationships) one can e.g., find correlations between attributes:
      - “Are there duplicate tuples?” or
      - “Are the data distributions in their respective domains similar?”
  - **Structure-based matching**
    - Abstracting from the actual labels, only the structure of the schema is evaluated, e.g., regarding element types, depths in hierarchies, number and type of relationships, etc.



## 4.2 Outlook

- Sometimes schema integration is **query-driven**
  - The integration is only needed in order to query several different information sources having different schemas
- In that case only a **schema mapping** is needed
  - Basically the mapping is a **list of correspondences** between equivalent entities or relationships of heterogeneous schemas
  - The query can then be **translated** for each different schema using the mapping
  - The mapping can be derived manually or automatically from a respective matching

- © Fisheries and Oceans, Canada





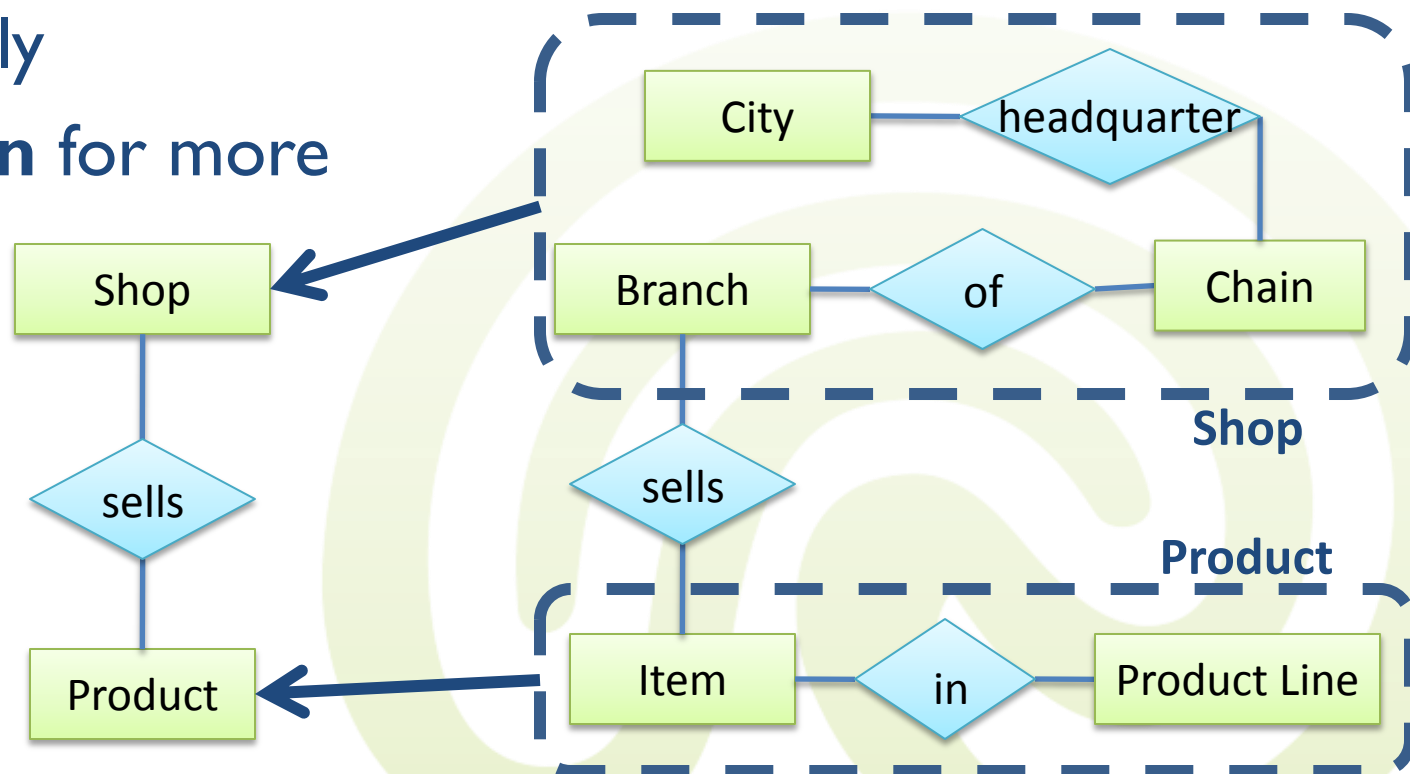
## 4.3 Entity Clustering

- When multiple schemas are merged, global schema can become **very large**
  - Many different entities and relationships between them
    - E.g., global view of a company with all its dependencies
  - But some parts from different views are entirely independent
    - E.g., accounting does not need technical specifications of products
- **Idea:** Cluster **semantically coherent parts** and abstract from their actual entities in the global schema



## 4.3 Entity Clustering

- Abstracting complex units allow showing the entire model on a **single sheet of paper**
  - Easy to get an overview and easier to integrate units separately
  - **Zoom in** for more details





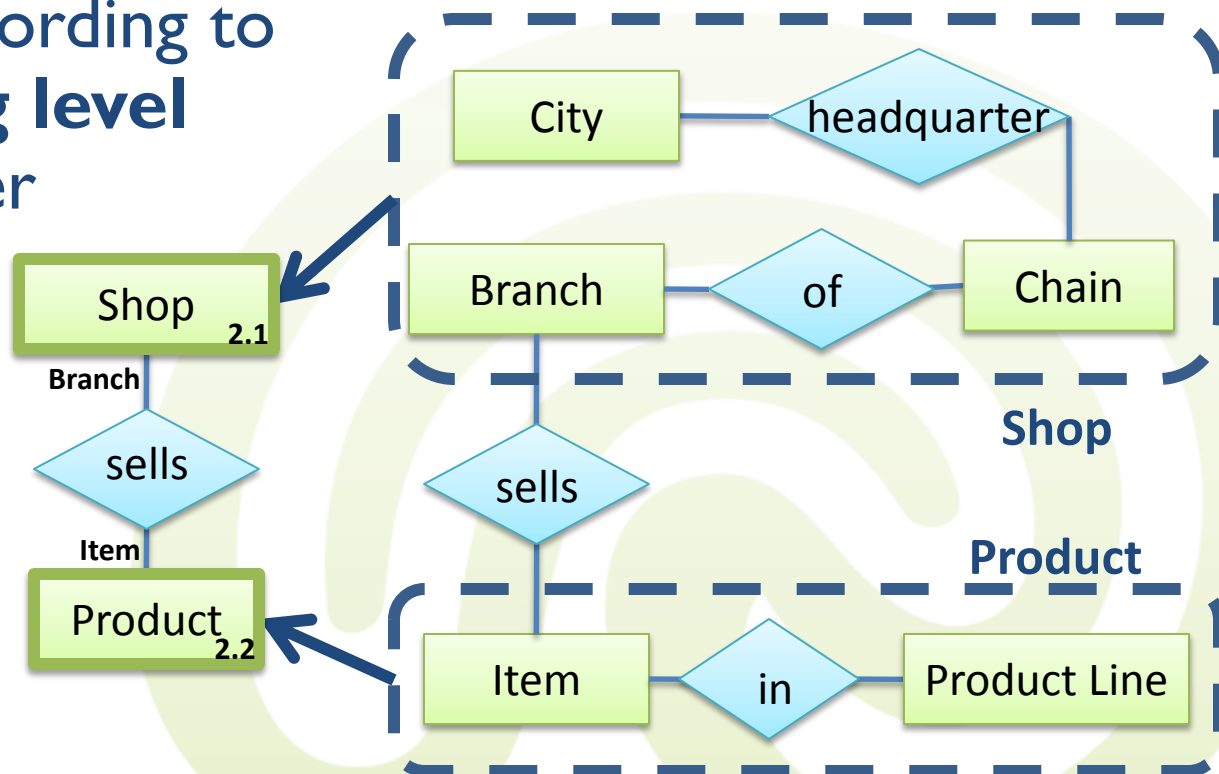
## 4.3 Clustering Concepts

- **Grouping** is an operation that combines entities and their relationships to form higher-level constructs
  - Groups are called **entity clusters**
  - Can also be performed **hierarchically** from the entire database (root entity cluster) over several levels down to the individual entities
    - All original entities are on clustering level 1



## 4.3 Clustering Concepts

- Usually, **entity clusters** are depicted similar to normal entities in ER diagrams
  - By a **dark-bordered** box
  - Numbered according to the **clustering level** and an identifier
  - Interfaces for **inter-cluster** relationships have to be annotated








## 4.3 Grouping Operations

- **Grouping operations** are the fundamental components of entity clustering
  - All operations are **heuristic** in nature
- **Often occurring** operations are
  - Dominance grouping
  - Abstraction grouping
  - Constraint grouping
  - Relationship grouping
- They can be applied **recursively** or in a variety of combinations to produce higher level clusters



## 4.3 Grouping Operations


- **Dominance grouping** focuses on semantically dominant entities in the ER diagrams
  - Hubs for **otherwise unconnected** or **weak** entities





## 4.3 Grouping Operations


- **Abstraction grouping** clusters entities of a specific super-type
  - Especially helpful, if sub-classes have no individual relationships





## 4.3 Grouping Operations


- **Constraint grouping** clusters entities related by the same constraint
  - E.g., integrity constraints such as XOR constraints





## 4.3 Grouping Operations

- **Relationship grouping** focuses on ternary or higher-degree relationships
  - The relationship is represented **as a whole**





## 4.3 Clustering Technique

- Identify all major functional areas and subareas in a top down analysis
  - Functional areas are often defined during the requirement analysis as important organizational units (e.g., HR or R&D) or business activities
  - Usually there will be a certain degree of overlap, for example employee data will be administrated by HR, but may also be needed in other departments



## 4.3 Clustering Technique

- The actual clustering has **four steps**
  - Define points of grouping within each functional area
    - Locate **dominant entities**, consider **abstraction**, find ***n*-ary** or **constrained relationships**, etc.
    - If such points do not exist, consider grouping the entire area
  - Form entity clusters
    - Use the **basic grouping operations** on elementary entities and their relationships to form higher level clusters
    - Since entities might belong to several clusters, **define priorities** like “always prefer abstraction grouping,” “avoid crossing boundaries of functional areas,” or “leave entities ungrouped, if they belong to two or more groups at the same level of precedence”





## 4.3 Clustering Technique

- Form higher level entity clusters
  - Apply the grouping operations **recursively** to any combination of elementary entities and entity clusters
  - Stop, if the diagram's complexity is **sufficiently low**:  
This defines the root entity cluster
- Validate the cluster diagram
  - Check for **consistency** of the interfaces (relationships) between entities or entity clusters at each level of the diagram
  - **Verify the meaning** of each level with the intended users



## 4.4 The BEA Story

*Detour*

- What happens, if you don't integrate properly?
  - Think about the Mars disaster...
- What happens, if you integrate?
  - Well, your processes are improved and you become more efficient...
- What happens, if you **help others** to integrate?
  - Short version: you found a company, get insanely rich and are finally bought by Oracle for 8.5 billion USD in 2008





## 4.4 The BEA Story

*Detour*

- **BEA Systems Inc.**

- Founded in 1995 in San José, CA, USA
- Before Oracle's takeover, the company had more than 4000 employees and about **one billion** in revenues



- **Product lines**

- Tuxedo for **distributed transaction processing** (1995)
- WebLogic provides a **J2EE enterprise infrastructure** (1998)
- AquaLogic provides a **service-oriented infrastructure** (2005)

- **Acquisitions** of some companies specializing in middleware and business process management

- WebLogic (1998), SolarMetric (2005), Plumtree Software (2006), Fuego (2006), ...



## 4.4 The BEA Story

*Detour*

- What are they actually doing?

„DekaBank



Finanzgruppe

- **Case study:**

- The **DekaBank Group** is the central asset manager of the Sparkasse Financial Group managing funds of around 90 billion EUR
- The Bank wanted an **access layer** to central data sources so that all data for the portfolio structure is available for fund management
- In 2006 DekaBank deployed the **BEA AquaLogic Data Services Platform**, which models the central data uniformly in a technical context and provides these business objects to local applications in real time



# 4.4 The BEA Story

*Detour*

## – Two **Challenges**

- Consolidation of various pieces of information from numerous channels
- Provide the information in different formats such that local applications can further process the data

– Finally, after a lot of integration, data is presented to the outside via a **standard access layer** in real time

– **Duration:** about five month

– **Costs:** ???





## 4.4 The BEA Story

*Detour*

- **BEA AquaLogic Data Services**

- Special Feature: **easy-to-use modeling**

- “In an **SOA environment**, a data model must be flexible so that it can represent any **complex entity** and rich enough to provide information about **data structure, relationships**, and services to read or update”
- “Data services are illustrated in **model diagrams** and can easily be shared with others in the enterprise for greater data consistency and reuse.”
- “Mappings and transformations can be designed in an easy-to-use **GUI tool** using a library of over 200 functions. For complex mappings and transformations, architects and developers can bypass the GUI tool and use an XQuery source code editor to define or edit services.”







## 4.4 The BEA Story

*Detour*

- What tools are actually given to support integration?
  - Data Translation Tool
    - Transforms binary data into XML
    - Transforms XML to binary data
  - Data Transformation Tool
    - Transforms an XML to another XML
  - Idea
    - Transform data to application specific XML
      - Transform to other application's XML or general schema
      - Transform back to binary
    - Note: the integration work still has to be done **manually**



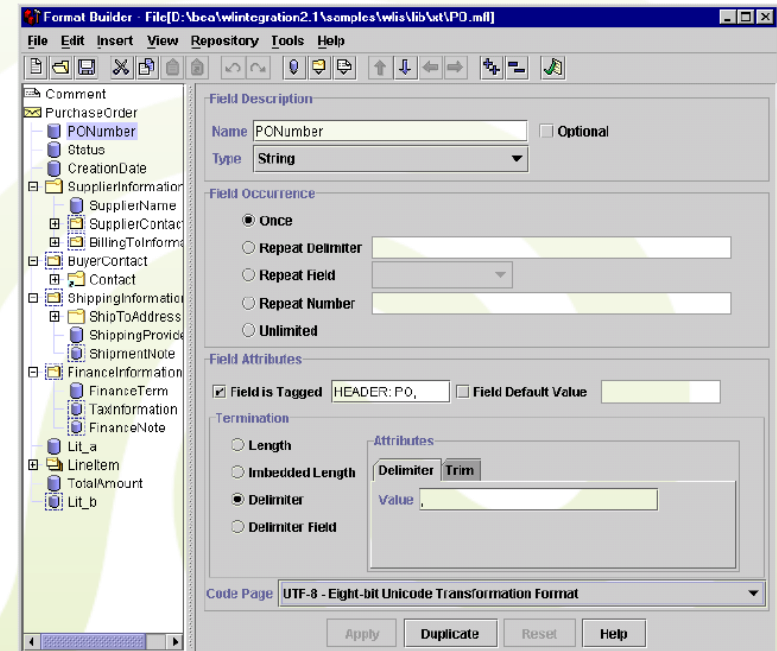
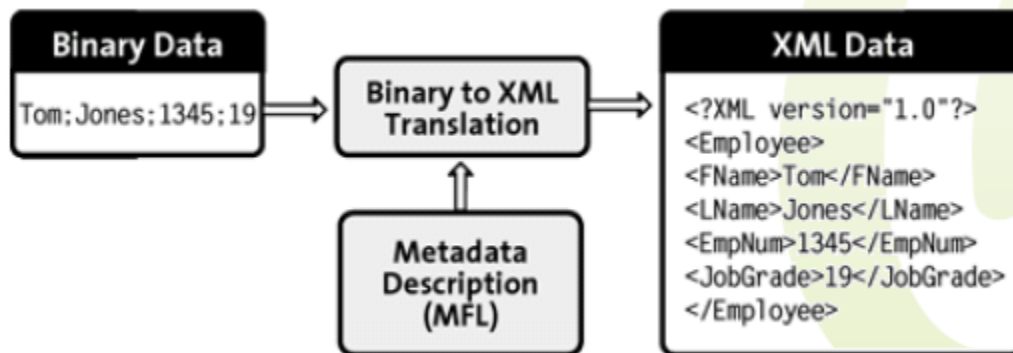




# 4.4 The BEA Story

# Detour

- Data Translation
  - Metadata Description Language describes how **binary data** (files, DB query results, etc.) is mapped to XML
  - BEA provides **editor** and execution **engine** for MFL
  - **Customer** creates MFL
    - Or alternatively: **Highly paid BEA consultants**

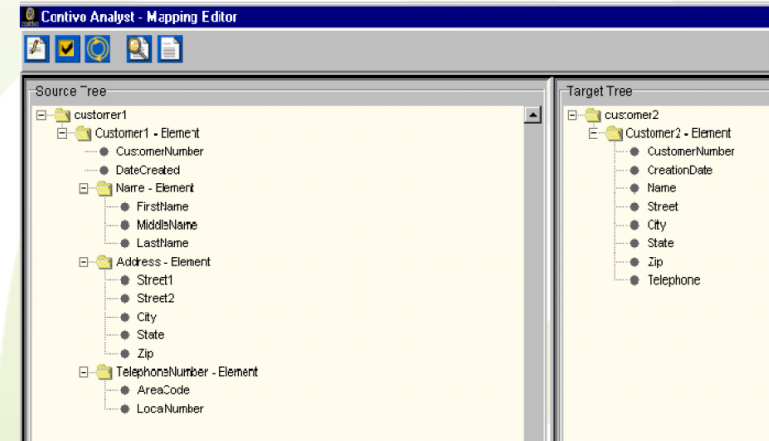
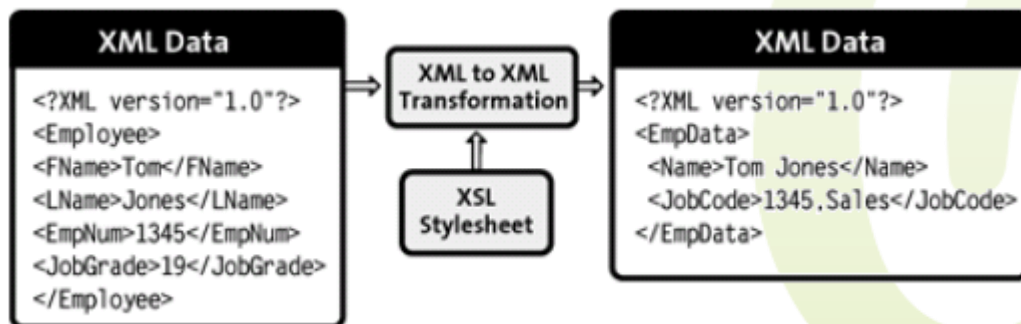




# 4.4 The BEA Story

*Detour*

- Data Transformation
  - XSL (eXtensible Stylesheet Language, a W3C standard...) is used to transform XML to different XML
  - BEA provides **editor** and execution **engine** for XSL
    - Both are rather...simple?
  - **Customer** manually creates XSL
    - Or alternatively:  
**Highly paid BEA consultants**





## 4.4 The BEA Story

*Detour*

- “I can’t afford expensive BEA consultants and the AquaLogic Integration Suite, what now??”
  - Do it completely **yourself**
    - Most used technologies can be found as open source projects (data mappers, XSL engines, XSL editors, etc.)
  - Do it **yourself** with **specialized tools**
    - Many companies and open source projects are specialized in developing data integration and transformation tools
      - CloverETL
      - Altova MapForce
      - BusinessObjects Data Integrator
      - etc...





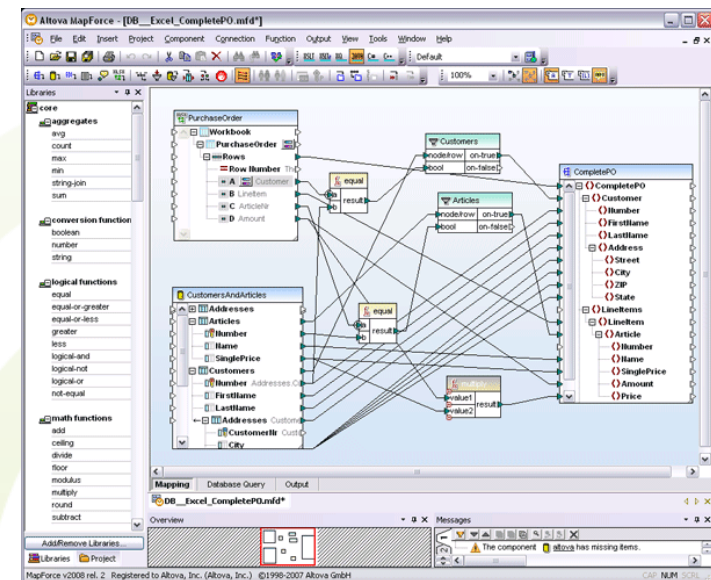
# 4.4 The BEA Story

# Detour

- **Altova MapForce**



- Same idea than BEA Integrator
  - Also based on XSLT and a data description language
- Editors for binary/DB to XML mapping
- Editor for XSL transformation
- Automatic generation of data sources, web-services, and transformation modules in Java, C#, C++





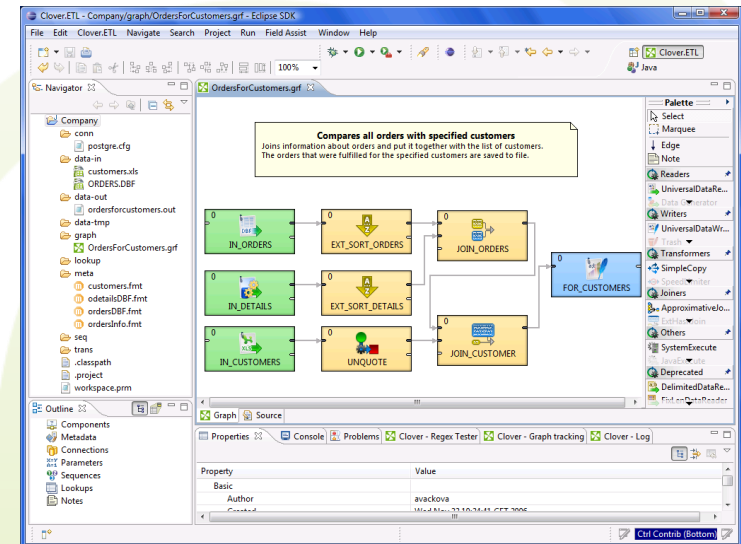
# 4.4 The BEA Story

# Detour

- **CloverETL**



- Based on own ETL transformation language
- Core tools are open source
  - Server and GUI tools are sold under commercial license
- Can read data from any database
- (Visually designed) ETL Script converts data into other data
  - XML
  - DB with different schema
  - etc





# Next Week

- Basic set theory
- Relational data model
- Transformation from ER
- Integrity Constraints
- From Theory to Practice

