



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# Relational Database Systems I

**Wolf-Tilo Balke**




**Simon Barthel**

Institut für Informationssysteme  
Technische Universität Braunschweig  
[www.ifis.cs.tu-bs.de](http://www.ifis.cs.tu-bs.de)



# Overview

- Basic set theory
- Relational data model
- Transformation from ER
- Integrity Constraints
- From Theory to Practice





# 5.1 Basic Set Theory

- **Set theory** is the foundation of mathematics
  - You probably all know these things from your math course, but repeating never hurts
  - The **relational model** is based on set theory; understanding the basic math will help a lot






## 5.1 Sets

- A **set** is a mathematical **primitive**, and thus has no formal definition
- A set is a **collection** of objects (called *members* or *elements* of the set)
  - Objects (or entities) have to be understood in a very broad sense, can be anything from physical objects, people, abstract concepts, other sets, ...
- Objects **belong** (or do **not belong**) to a set (alternatively, *are* or *are not* in the set)
- A set **consists** of all its elements



# 5.1 Sets


- Sets can be specified **extensionally**
  - List all its elements
  - Example:  $A = \{\text{IfIS}, 42, \text{Balke}, \text{Hurz!}\}$
- Sets can be specified **intensionally**
  - Provide a criterion deciding whether an object belongs to the set or not (**membership criterion**)
  - Examples:
    - $A = \{x \mid x > 4 \text{ and } x \in \mathbb{Z}\}$
    - $B = \{x \in \mathbb{N} \mid x < 7\}$
    - $C = \{\text{all facts about databases you should know}\}$
- Sets can be either **finite** or **infinite**
  - Set of all super villains is finite
  - Set of all numbers is infinite





## 5.1 Sets

- Sets are different, iff they have different members
  - $\{a, b, c\} = \{b, c, a\}$
  - Duplicates are not supported in standard set theory
    - $\{a, a, b, c\} = \{a, b, c\}$
- Sets can be empty (written as  $\{\}$  or  $\emptyset$ )
- Notations for set membership:
  - $a \in \{a, b, c\}$
  - $e \notin \{a, b, c\}$






## 5.1 Sets

- Defining a set by its **intension**:
  - Intension must be **well-defined** and **unambiguous**
  - There always is a clear **membership criterion** to determine whether an object belongs to the set (or not)
  - Not a valid definition (Russell's paradox):

In a small town, there is just one male barber.  
He shaves all and only those men in town  
who do not shave themselves.

- **Does the barber shave himself?**





## 5.1 Sets

- Still, the set's **extension** might be unknown (however, there is one)
- Example:
  - “All students in this room who are older than 22”
  - Well-defined, but not known to me ...
  - But (at least in principle) we can find out!
- As we will see later:
  - Intension  $\approx$  Database query
  - Extension  $\approx$  Result of a query







# 5.1 Sets





- For every set, there is an accompanying definition of **equality** (or equivalence)
  - Is  $x = y$ ?
  - If they are equal, they are actually just one element
- However, you could have two different descriptions of the same element
  - Example: The set of all 26 “standard” letters
    - ‘ö’ is not contained in this set
    - ‘m’ = ‘M’ and both reflect a single element of the set
      - ‘m’ and ‘M’ are different descriptions of the **same** object
  - Example: The set of all 59 letters and umlauts in German
    - ‘ö’ is element of the set
    - ‘m’  $\neq$  ‘M’ and are both elements of the set (two different objects)



## 5.1 Sets

- Sets have a cardinality (i.e., number of elements)
  - Denoted by  $|A|$
  - $|\{a, b, c\}| = 3$
- Set  $A$  is a **subset** of set  $B$ , denoted by  $A \subseteq B$ , iff every member of  $A$  is also a member of  $B$
- $B$  is a **superset** of  $A$ , denoted by  $B \supseteq A$ , iff  $A \subseteq B$





# 5.1 Tuples

- A **tuple** (or vector) is a sequence of objects
  - Length 1: Singleton
  - Length 2: Pair
  - Length 3: Triple
  - Length  $n$ :  $n$ -tuple
- In contrast to sets ...
  - Tuples can contain an object more than once
  - The objects appear in a certain order
  - The length of the tuple is finite
- Written as  $\langle a, b, c \rangle$  or  $(a, b, c)$






# 5.1 Tuples

- Hence:
  - $\langle a, b, c \rangle \neq \langle c, b, a \rangle$ , whereas  $\{a, b, c\} = \{c, b, a\}$
  - $\langle a_1, a_2 \rangle = \langle b_1, b_2 \rangle$  iff  $a_1 = b_1$  and  $a_2 = b_2$
- **$n$ -tuples** ( $n > 2$ ) can also be defined as a cascade of ordered pairs:
  - $\langle a, b, c, d \rangle = \langle a, \langle b, \langle c, d \rangle \rangle \rangle$



# 5.1 Set Operations


- Four binary **set operations**
  - Union, intersection, difference and Cartesian product
- Union:  $\cup$ 
  - Creates a new set containing all elements that are contained in (at least) one of two sets
  - $\{a, b\} \cup \{b, c\} = \{a, b, c\}$
- Intersection:  $\cap$ 
  - Creates a new set containing all elements that are contained in both sets
  - $\{a, b\} \cap \{b, c\} = \{b\}$





# 5.1 Set Operations


- Difference:  $\setminus$ 
  - Creates a set containing all elements of the first set without those also being in the second set
  - $\{a, b\} \setminus \{b, c\} = \{a\}$





# 5.1 Set Operations

- Cartesian product:  $\times$ 
  - The Cartesian product is an operation between two sets, creating a new set of pairs such that:
$$A \times B = \{ \langle a, b \rangle \mid a \in A \text{ and } b \in B \}$$
  - Named after René Descartes
- Example:
  - $\{a, b\} \times \{b, c\} = \{ \langle a, b \rangle, \langle a, c \rangle, \langle b, b \rangle, \langle b, c \rangle \}$
  - Cleverness = { genius, dumb }
  - Character = { hero, villain }
  - $\text{Cleverness} \times \text{Character} = \{ \langle \text{genius, hero} \rangle, \langle \text{dumb, hero} \rangle, \langle \text{genius, villain} \rangle, \langle \text{dumb, villain} \rangle \}$
- The Cartesian product can easily be extended to higher dimensionalities:  $A \times B \times C$  is a set of triples





# 5.1 Relations

- A **relation**  $R$  over some sets  $D_1, \dots, D_n$  is a **subset** of their **Cartesian** product
  - $R \subseteq D_1 \times \dots \times D_n$
  - The elements of a relation are **tuples**
  - The  $D_i$  are called **domains**
  - Each  $D_i$  corresponds to an **attribute** of a tuple
    - $n=1$ : Unary relation or **property**
    - $n=2$ : Binary relation
    - $n=3$ : Ternary relation
    - ...






# 5.1 Relations

- Some important properties:
  - Relations are sets in the mathematical sense, thus **no duplicate tuples** are allowed
  - The **list of tuples** is **unordered**
  - The **list of domains** is **ordered**
  - Relations can be modified by...
    - **inserting** new tuples,
    - **deleting** existing tuples, and
    - **updating** (that is, modifying) existing tuples.



# 5.1 Relations

- A special case: Binary relations
  - $R \subseteq D_1 \times D_2$ 
    - $D_1$  is called **domain**,  $D_2$  is called **co-domain** (range, target)
  - Relates objects of two different sets to each other
  - $R$  is just a set of ordered pairs
  - $R = \{ \langle a, 1 \rangle, \langle c, 1 \rangle, \langle d, 4 \rangle, \langle e, 5 \rangle, \langle e, 6 \rangle \}$ 
    - Can also be written as  $aR1, cR1, dR4, \dots$
  - Imagine **Likes**  $\subseteq$  **Person**  $\times$  **Beverage**
    - Joachim Likes Coffee, Tilo Likes Tea, ...





## 5.1 Relations

- Example:

- accessory = {spikes, butterfly helmet}
- material = {silk, armor plates}
- color = {pink, black}

**color × material × accessory =**  
{<pink, silk, butterfly helmet>,  
<pink, silk, spikes>,  
<pink, armor plates, butterfly helmet>,  
<pink, armor plates, spikes>,  
<black, silk, butterfly helmet>,  
<black, silk, spikes>,  
<black, armor plates, butterfly helmet>,  
<black, armor plates, spikes>}



## 5.1 Relations

- Relation **FamousHeroCostumes**  
 $\subseteq \text{color} \times \text{material} \times \text{accessory}$


**FamousHeroCostumes** =  
{<pink, silk, butterfly helmet>,  
<black, armor plates, spikes>}





# 5.1 Functions

- **Functions** are special case of binary relations
  - **Partial function:**  
Each element of the domain is related to **at most one** element in the co-domain
  - **Total function:**  
Each element in the domain is related to **exactly one** element in the co-domain





# 5.1 Functions

- Functions can be used to **abstract** from the exact order of domains in a relation
  - Alternative definition of relations:  
**A relation is a set of functions**
  - Every tuple in the relation is considered as a function of the type  $\{A_1, \dots, A_n\} \rightarrow D_1 \cup \dots \cup D_n$ 
    - That means, every tuple maps each attribute to some value



# 5.1 Functions

- Example:
  - color = {pink, black}
  - material = {silk, armor plates}
  - accessory = {spikes, butterfly helmet}
  - The tuple <pink, silk, butterfly helmet> can also be represented as the following function  $t$ :
    - $t(\text{color}) = \text{pink}$
    - $t(\text{material}) = \text{silk}$
    - $t(\text{accessory}) = \text{butterfly helmet}$
  - Usually, one writes  $t[\text{color}]$  instead of  $t(\text{color})$



## 5.2 Relational Model

- Well, that's all nice to know... but:  
we are here to learn about **databases!**
  - Where is the connection?
- **Here it is...**
  - A **database schema** is a description of concepts in terms of attributes and domains
  - A **database instance** is a set of objects having certain attribute values







## 5.2 Relational Model

- OK, then...
  - Designing a database schema (e.g., by ER modeling) determines entities and relationships, as well as their corresponding **sets of attributes** and associated **domains**
  - The **Cartesian product** of the respective domains is the set of all possible instances (of each entity type or relationship type)
  - A **relation** formalizes the **actually existing** subset of all possible instances



## 5.2 Relational Model

- Database schemas are described by **relation schemas**, denoted by  $R(A_1:D_1, \dots, A_n:D_n)$
- The actual database instance is given by a set of matching **relations**
- Example
  - Relation schema:  
**CATS**(name : varchar(10), age : integer)
  - A matching relation:  
 $\{ (Blackie, 10), (Pussy, 5), (Fluffy, 12) \}$





## 5.2 Relational Model

- Relations can be written as **tables**:

Diagram illustrating a relation (table) with annotations:

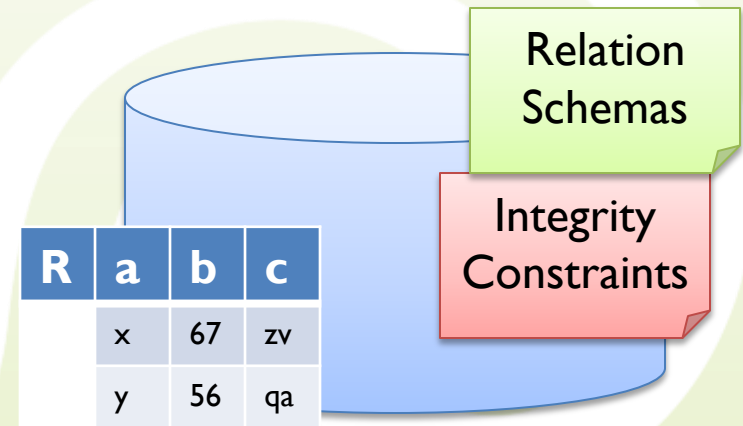
- relation name**: Points to the table name **PERSON**.
- attributes**: Points to the column headers **firstName**, **lastName**, and **sex**.
- tuples**: Points to the rows of data.
- domain values**: Points to the values in the **sex** column.

PERSON	firstName	lastName	sex
	Clark Joseph	Kent	m
	Louise	Lane	f
	Lex	Luthor	m
	Charles	Xavier	m
	Erik	Magnus	m
	Jeanne	Gray	f
...	Ororo	Munroe	f
	Tony Edward	Stark	m
	Matt	Murdock	m
	Raven	Wagner	f
	Robert Bruce	Banner	m



## 5.2 Relational Model

- A **relational database schema** consists of
  - a set of relation schemas
  - a set of integrity constraints
- A **relational database instance** (or state) is
  - A set of relations adhering to the respective schemas and respecting all integrity constraints





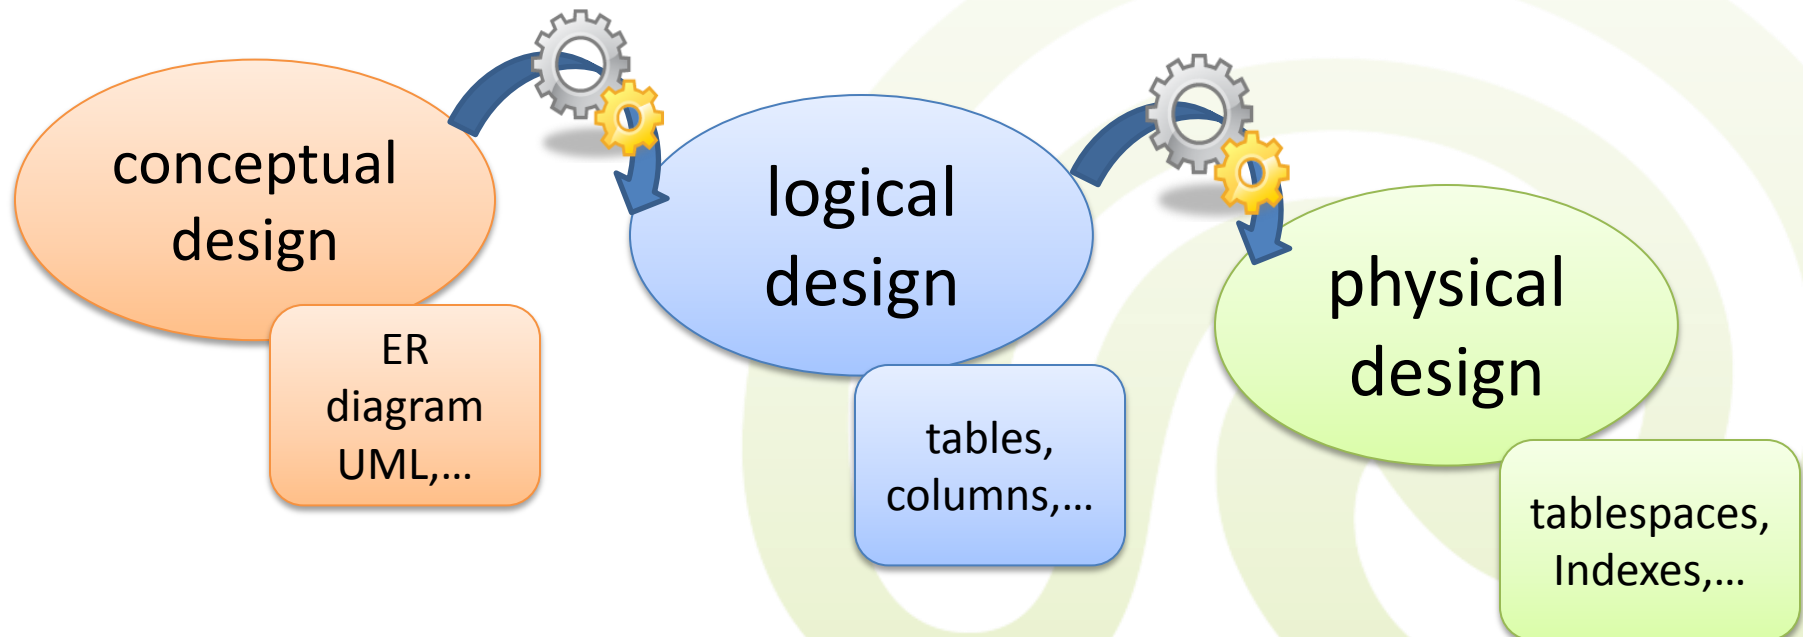
## 5.2 Relational Model

- Every relational DBMS needs a language to define its relation schemas (and integrity constraints)
  - **Data definition language (DDL)**
  - Typically, it is difficult to formalize all possible integrity constraints, since they tend to be complex and vague
- A relational DBMS also needs a language to handle tuples
  - **Data manipulation language (DML)**
- Today's RDBMS use **SQL** as both DDL and DML



## 5.3 Conversion from ER *Detour*

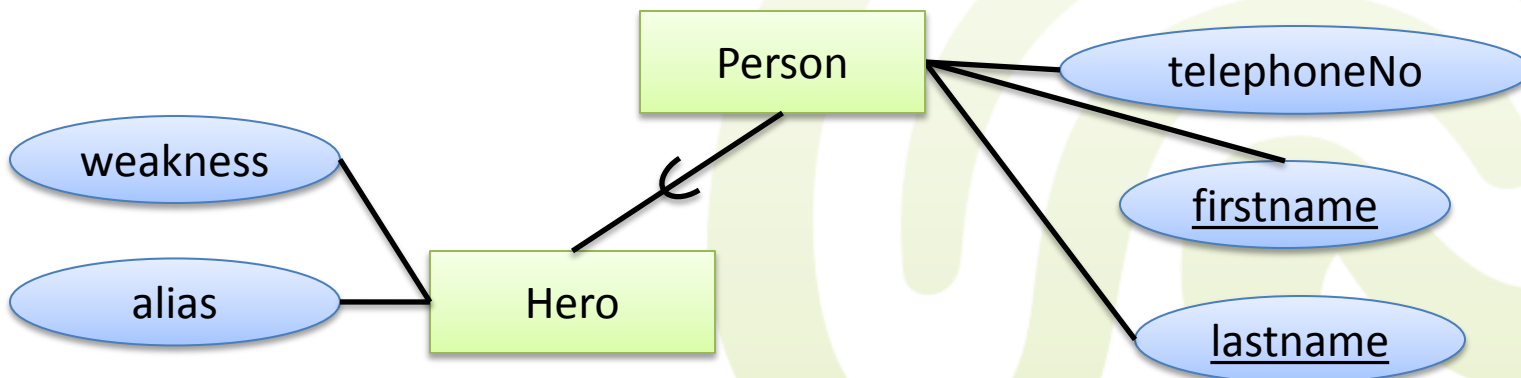
- After modeling a conceptual schema (e.g., using an **ER diagram**), the schema can be **automatically** transformed into a **relational schema**
- Remember:





## 5.3 Conversion from ER *Detour*


- Each **entity type**  $E$  with attributes  $A_1, \dots, A_n$  from domains  $D_1, \dots, D_n$  is converted into an  **$n$ -ary relation schema**  $E(A_1:D_1, \dots, A_n:D_n)$
- If there is a relationship type  $E$  is\_a  $F$  involved (**specialization**), the inheritance relationship can be expressed by copying all key attributes from  $F$





## 5.3 Conversion from ER *Detour*

- A **relationship type**  $R$  between entity types  $E_1, \dots, E_n$  is converted to a relation schema whose attributes are all the key attributes of  $E_i$ 
  - If keys share the same name, they have to be renamed
- If the relationship type has **own attributes** they are also copied to the relation schema

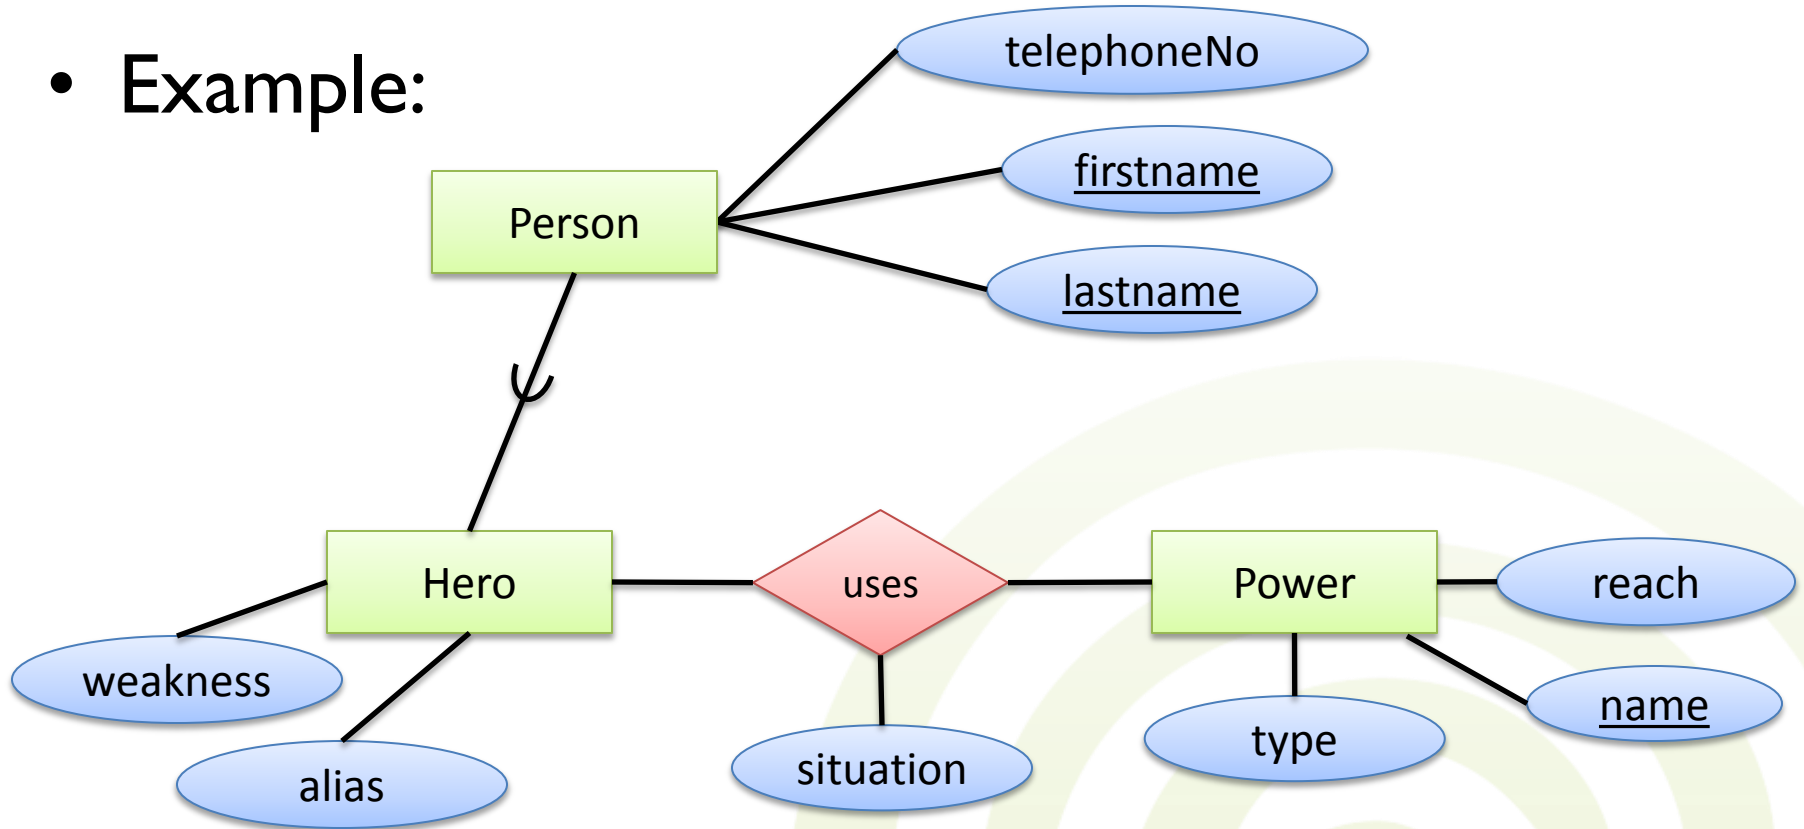






## 5.3 Conversion from ER *Detour*

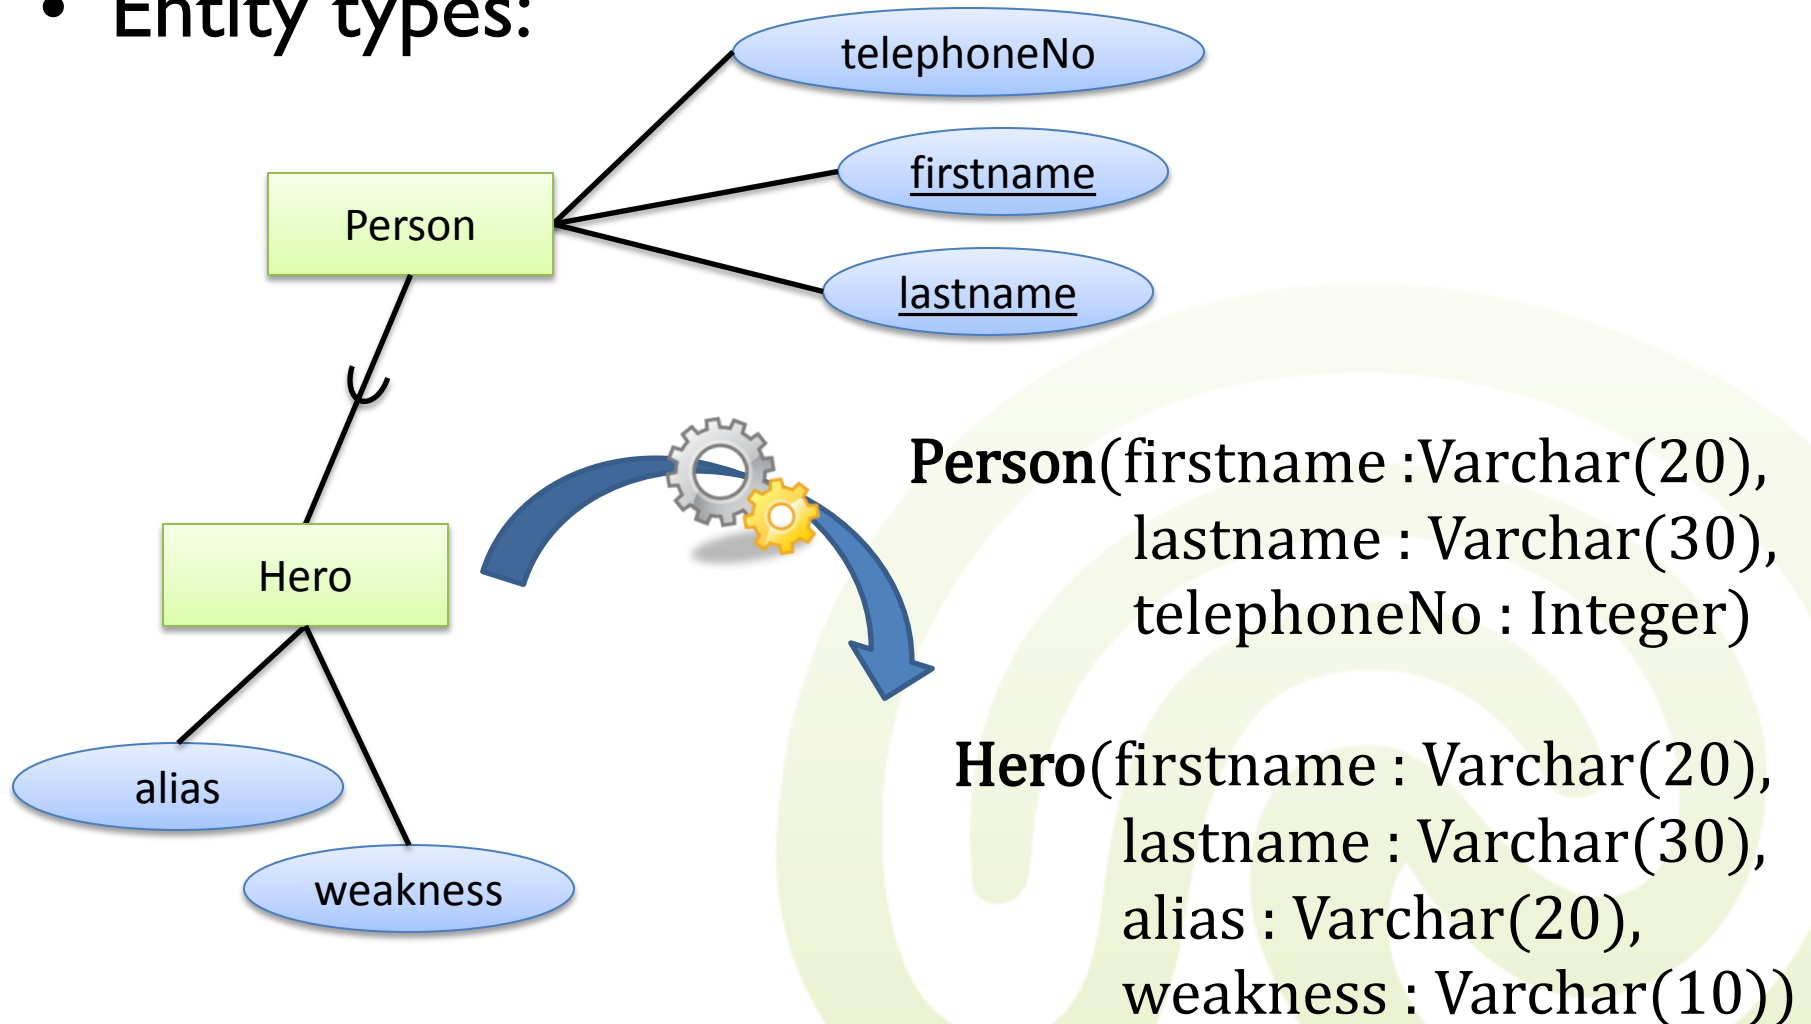
- Example:





## 5.3 Conversion from ER *Detour*


- Entity types:






## 5.3 Conversion from ER *Detour*

- Relationship types:



**Hero**(firstname : Varchar(20),  
lastname : Varchar(30),  
alias : Varchar(20),  
weakness : Varchar(10))

**Power**(name : Varchar(20),  
type : Varchar(20),  
reach : Integer)



**Uses**(firstname : Varchar(20),  
lastname : Varchar(30),  
name : Varchar(20),  
situation : Varchar(30))



## 5.3 Conversion from ER *Detour*

- Note: The ER diagram is **semantically richer** than the relational model
- Examples:
  - No key constraints and functionalities yet
  - Integrity constraints like disjoint/overlapping generalization cannot be expressed
  - ...
- Therefore, it usually is a really good idea to create an ER diagram before coding a logical schema



## 5.4 Integrity Constraints


- Integrity constraints are difficult to model in ER
  - Basically annotations to the diagram, especially for **behavioral constraints**
    - **Example:** The popularity rating of any assistant should always be less than the respective professor's
- But some **structural constraints** can directly be expressed
  - Key constraints
  - Functionalities





## 5.4 Inherent Constraints

- A relation is defined as a **set** of tuples
  - All tuples have to be **distinct**, i.e., no two tuples can have the same combinations of values for all attributes
  - So-called **uniqueness (unique key) constraint**
- Any subset of a relation type's attributes is called a **superkey**, if any two tuples can never share the same values with respect to this subset
  - The set of all attributes is always a superkey, but there may be smaller sets
  - Superkeys may have redundant attributes





## 5.4 Inherent Constraints

- A minimal superkey is called a **key**
  - “Minimal” means that no attribute can be removed without losing the superkey property
  - Of course, a relation can have several keys
  - The key property is determined from the **semantics** of the attributes, not from the current data instances
  - Example:
    - Relation `address(street, number, zip code, city)`
    - Keys:  
`{street, number, zip code}` and `{street, number, city}`



## 5.4 Inherent Constraints

- The set of keys of some relation  $R$  is called **candidate key set** or  $\text{cand}(R)$
- For each relation, a single candidate key has to be chosen to identify tuples in the relation, the so-called **primary key**
  - Analogously to ER diagrams, the chosen primary key is often underlined in the relation schema
    - Example: `address(street, number, zip-code, city)`
  - Though any candidate key can be chosen, it is usually better to choose a primary key with a **small** number of attributes






## 5.4 Inherent Constraints

- Assume that an ER diagram is converted into relation schemas, what are the **candidate keys** of the relation schemas?
- If the relation schema...
  - ... has been derived from some **entity type**  $E$  with key attributes  $K_E := \{A_1, \dots, A_n\}$ , then  $K_E \in \text{cand}(R)$
  - ... has been derived from an **N:M relationship type** between  $E$  and  $F$ , then  $K_E \cup K_F \in \text{cand}(R)$



## 5.4 Inherent Constraints

- If the relation schema...
  - ... has been derived from an **1:1 relationship type** between  $E$  and  $F$ , then  $\text{cand}(E) \cup \text{cand}(F) = \text{cand}(R)$
  - ... has been derived from an **N:1 relationship type** between  $E_1, \dots, E_n$  and  $F$ , then  $K_{E_1} \cup \dots \cup K_{E_n} \in \text{cand}(R)$ 
    - In this case, it might also be a good to add  $F$ 's key attributes





## 5.4 Inherent Constraints

- Another constraint from ER diagrams is whether a value has to be provided for some attribute
  - NULL values, allowed by default
  - Again, this is a semantic property
- A second inherent constraint for each relation is that primary keys must **never** be NULL
  - So-called **entity integrity constraint**
  - Example: address( street NOT NULL,  
number NOT NULL,  
zip-code,  
city NOT NULL )




## 5.4 Inherent Constraints

- A third integrity constraint applies to all relation schemas that have been derived from **relationship types**
  - Relationship types borrow their (primary) keys from the entities involved, so-called **foreign keys**
  - Relationships only exist, if the respective entities exist
  - So-called **referential integrity constraint**



## 5.4 Inherent Constraints

- Example:



**Uses**(firstname : Varchar(20),  
lastname : Varchar(30),  
name : Varchar(20),  
situation : Varchar(30))

Foreign key from Hero → firstname  
Foreign key from Hero → lastname  
Foreign key from Power → name

- If a tuple ('Clark', 'Kent', 'X-Ray vision', 'Bomb threat') exists in **Uses**, there has to be a tuple ('Clark', 'Kent', ...) in Hero and a tuple ('X-Ray vision', ...) in Power



## 5.4 Inherent Constraints

- All three **structural constraints** have to be checked by the database
  - Unique key constraint
  - Entity integrity constraint
  - Referential integrity constraint
- This is especially necessary when inserting, deleting, or updating tuples in relations



## 5.4 First Normal Form

- There is another major **constraint** on the attributes' data types in the relational model
  - The value of any attribute must be **atomic**, that is, it **cannot be composed** of several other attributes
    - If this property is met, the relation is often referred to as being in **first normal form** (1NF or minimal form)
    - In particular, set-valued and relation-valued attributes (tables within tables) are prohibited





## 5.4 First Normal Form

- Example of a **set-valued** column:
  - A person may own several telephones (home, office, cell, ...)

Person	firstName	lastName	telephoneNo
	Clark Joseph	Kent	5555678
	Louise	Lane	{3914533, 3556576, 5463456}
	Lex	Luthor	4543689
	Charles	Xavier	7658736
	Erik	Magnus	{1252345, 8766781}







## 5.4 First Normal Form

- Please note, it is possible to **model** composed attributes in ER models...
- To transform such a model into the relational model, a **normalization** step is needed
  - This is not always trivial, e.g., what happens to keys?

Person	firstName	lastName	telephoneNo
	Clark Joseph	Kent	555-5678
	Louise	Lane	391-4533
	Louise	Lane	355-6576
	Louise	Lane	546-3456
	Lex	Luthor	454-3689
	Charles	Xavier	765-8736
	Erik	Magnus	125-2345
	Erik	Magnus	876-6781



## 5.4 First Normal Form


- In a purely relational database, all relations are in first normal form
  - **Object-oriented** databases feature multi-valued attributes, thus closing the modeling gap
  - **Object-relational extensions** integrate user-defined types (UDTs) into relational databases
    - Oracle from version 9i, IBM DB2 from version 8.1, ...



## 5.5 From Theory to Practice

# Detour

- In the early 1970s, the **relational model** became a “hot topic” database research
  - Based on **set theory**
  - A relation is a subset of the **Cartesian** product over a list of **domains**
- Early “query interfaces” for the relational model:
  - **Relational algebra**
  - **Tuple relational calculus (SQUARE, SEQUEL)**
  - **Domain relational calculus (QBE)**
- Question: **How to build a working database management system using this theory?**





- 
- The 17-author paper
- "Just type UFI!"
- The shootout at the OK corral
- SQL 20th year reunion: 1975-1995



- The challenge of the **System R** project was to create a **working prototype system**
  - Theory is good
  - But developers were willing to sacrifice theoretical beauty and clarity for the sake of usability and performance
- **Vocabulary change**
  - Mathematical terms were too unfamiliar for most people
  - **Table** = relation
  - **Row** = tuple
  - **Column** = attribute
  - **Data type, domain** = domain





- **Design decisions:**  
During the development of System R, two major and very controversial decisions had been made
  - Allow **duplicate tuples**
  - Allow **NULL values**
- Those decisions are still subject to discussions...





- **Duplicates**

- In a relation, there **cannot** be any **duplicate tuples**
- Also, query results cannot contain duplicates
  - The relational algebra and relational calculi all have implicit **duplicate elimination**








- **Practical considerations**

- You want to query for **name** and **birth year** of all **students** of TU Braunschweig
- The result returns roughly **13,000 tuples**
- Probably there are **some** duplicates
- It's **1973**, and your computer has **16 kilobytes** of main memory and a very slow external storage device...
- To eliminate duplicates, you need to **store** the result, **sort** it, and **scan** for adjacent duplicate lines
  - System R engineers concluded that this effort is not worth the effect
  - Duplicate elimination in result sets happens only on-request







- **Decision:** Don't eliminate duplicates in results
- What about the tables?
  - Again: Ensuring that no duplicates end up in the tables requires some work
  - Engineers also concluded that there is actually no need in enforcing the no-duplicate policy
    - If the user wants duplicates and is willing to deal with all the arising problems – then that's fine
- **Decision:** Allow duplicates in tables
- As a result, the theory underlying relational databases shifted from **set theory** to **multi-set theory**
  - Straightforward, only notation is more complicated



## 5.5 From Theory to Practice

# Detour

- Sometimes, an attribute value is **not known** or an attribute does **not apply** for an entity
  - Example: What value should the attribute *universityDegree* take for the entity *Heinz Müller*, if Heinz Müller does not have any degree?
  - Example: You regularly observe the weather and store temperature, wind strength, and air pressure every hour – and then your barometer breaks... What now?





- Possible solution:  
For each domain, **define a value** indicating that data is not available, not known, not applicable, ...
  - For example, use *none* for Heinz Müller's degree, use *-1* for missing pressure data, ...
  - Problem:
    - You need such a special value for each domain or use case
    - You need special failure handling for queries, e.g.  
“compute average of all pressure values that are not *-1*”



## 5.5 From Theory to Practice

## Detour

- Again, system designers chose the simplest solution (regarding implementation): **NULL values**
  - **NULL** is a special value which is usable in any domain and represents that data is just there
    - There are many interpretations of what NULL actually means
  - Systems have some default rules how to deal with NULL values
    - Aggregation functions usually ignore rows with NULL values (which is good in most, but not all cases)
    - Three-valued logic
    - However, creates some strange anomalies




- Another tricky problem:  
How should users **query** the DB?
- Classical answer:
  - Relational algebra and relational calculi
  - **Problem:** More and more **non-expert users**
- More “natural” query interfaces:
  - **QBE** (query by example)
  - **SEQUEL** (structured English query language)
  - **SQL**: the current standard; derived from SEQUEL



# Preview – Relational Algebra

- How do you work with relations?
- **Relational algebra!**
  - Proposed by Edgar F. Codd: “A Relational Model for Large Shared Data Banks,” Communications of the ACM, 1970
- The theoretical foundation of all relational databases
  - Describes how to manipulate relations and retrieve interesting parts of available relations
  - Relational algebra is mandatory for advanced tasks like query optimization





# Preview – Relational Algebra

- Elementary operations:
  - **Set algebra operations**
    - Set Union  $\cup$
    - Set Intersection  $\cap$
    - Set Difference  $\setminus$
    - Cartesian Product  $\times$
  - **New relational algebra operations**
    - Selection  $\sigma$
    - Projection  $\pi$
    - Renaming  $\rho$
- Additional derived operations (for convenience)
  - All sorts of joins  $\bowtie, \ltimes, \ltimes, \dots$
  - Division  $\div$
  - ...



# Preview – Relational Calculi

- Beside the **relational algebra**, there are two other major **query** paradigms within the relational model
  - **Tuple relational calculus (TRC)**
  - **Domain relational calculus (DRC)**
- All three provide the **theoretical** foundation of the relational database model
- They are mandatory for certain DB features:
  - Relational algebra → **Query optimization**
  - TRC → **SQL** query language
  - DRC → **Query-by-example** paradigm








# Preview – Relational Calculi

- Relational algebra has some **procedural** aspects
  - You specify an **order of operations** describing how to retrieve data
- Relational calculi (TRC, DRC) are **declarative**
  - You just **specify** how the desired **tuples look like**
  - The query contains no information about how to create the result set
  - Provides an alternative approach to querying



# Next Lecture

- Relational Algebra
  - Basic relational algebra operations
  - Additional derived operations
- Query Optimization
- Advanced relational algebra
  - Outer Joins
  - Aggregation

 $\sigma$  $\pi$ 