



ifis

Institut für Informationssysteme
Technische Universität Braunschweig

Relational Database Systems I

Wolf-Tilo Balke

Simon Barthel

Institut für Informationssysteme

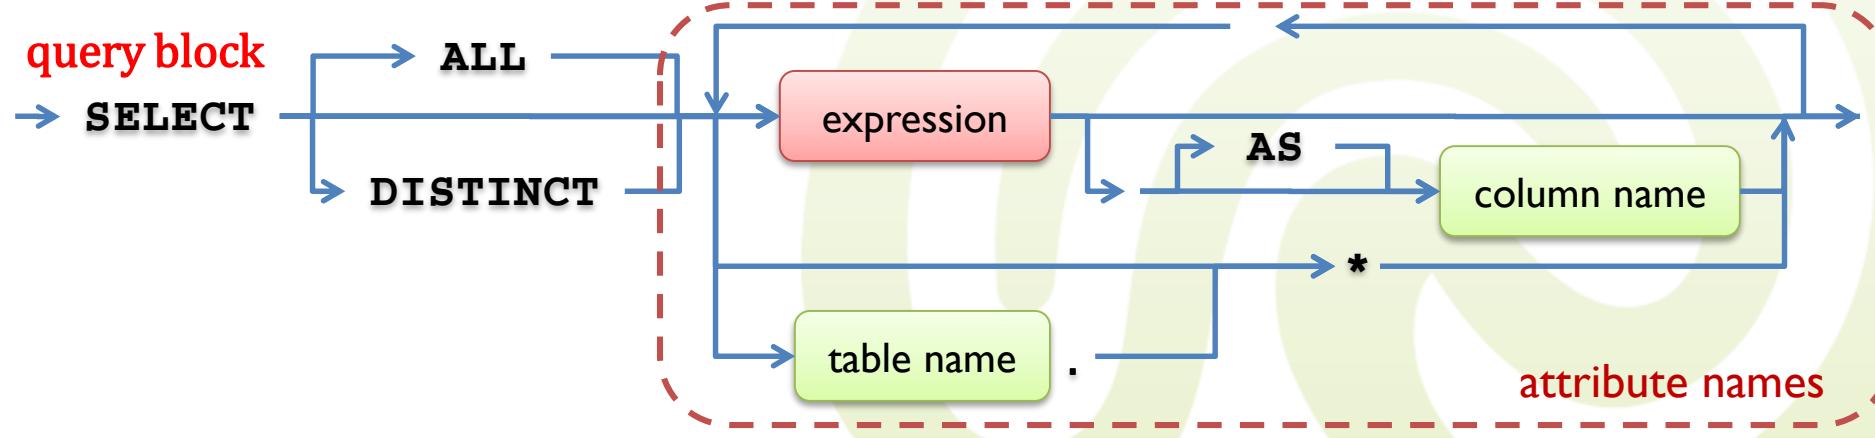
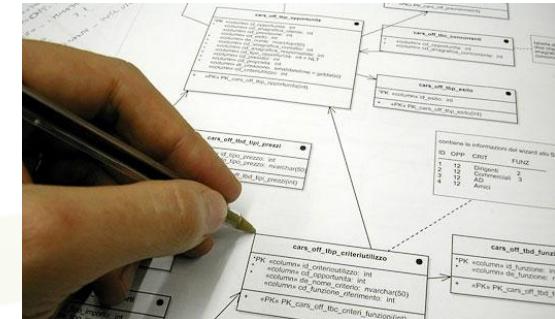
Technische Universität Braunschweig

www.ifis.cs.tu-bs.de



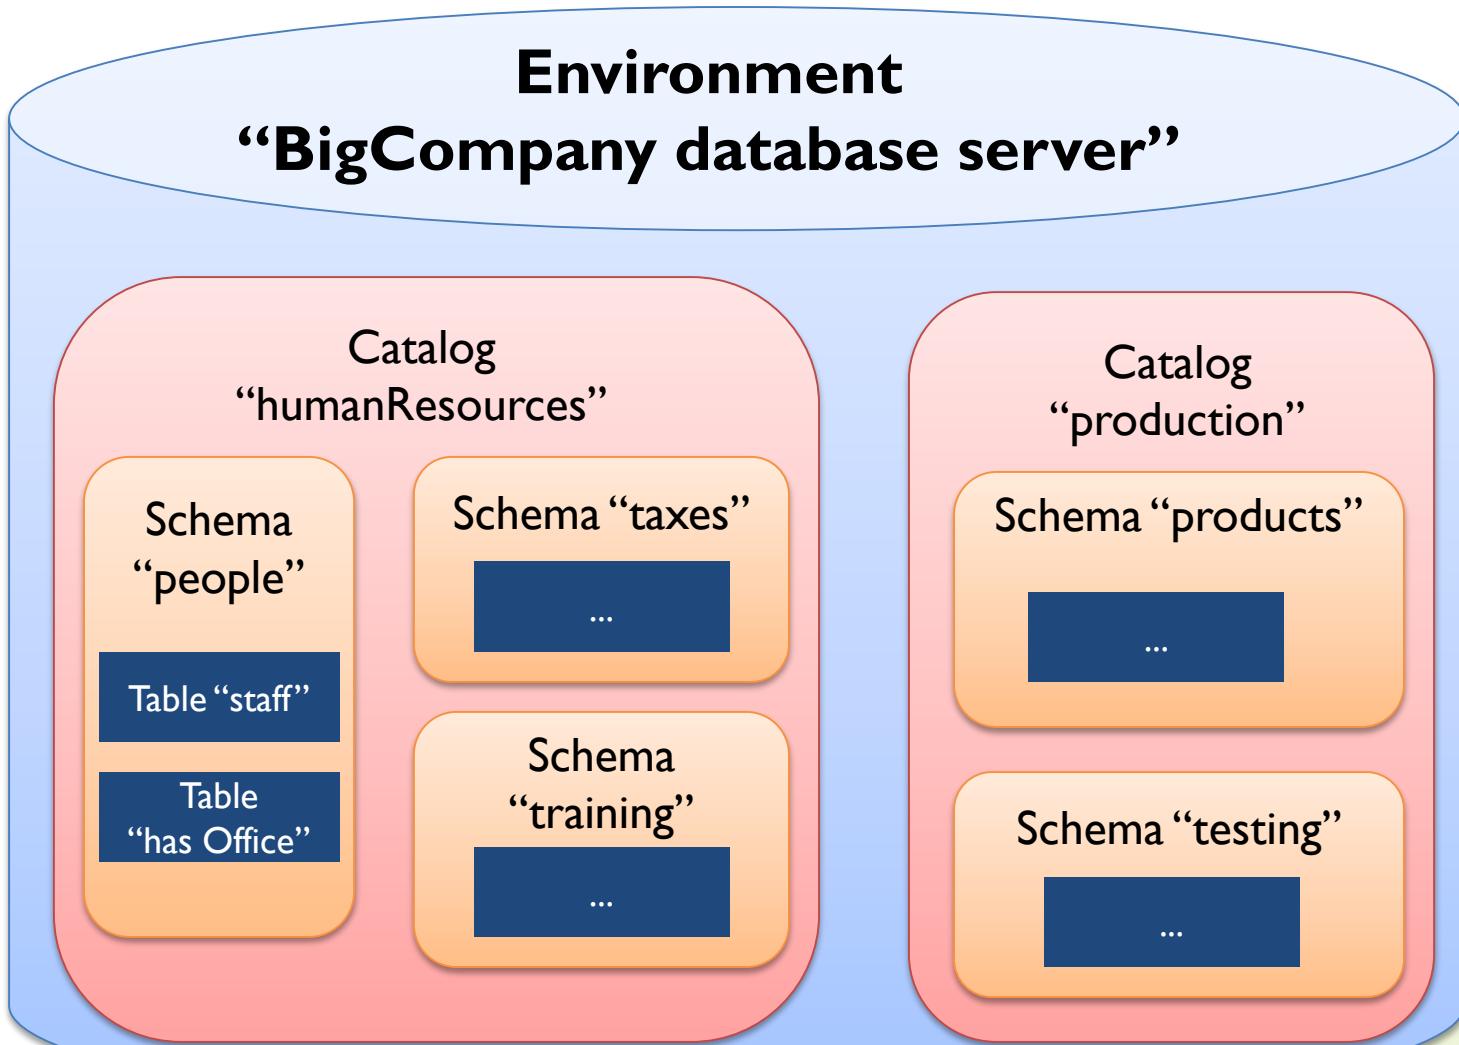
Overview

- SQL data definition language
- SQL data manipulation language
(apart from **SELECT**)
- $\text{SQL} \neq \text{SQL}$
- Some advanced SQL concepts



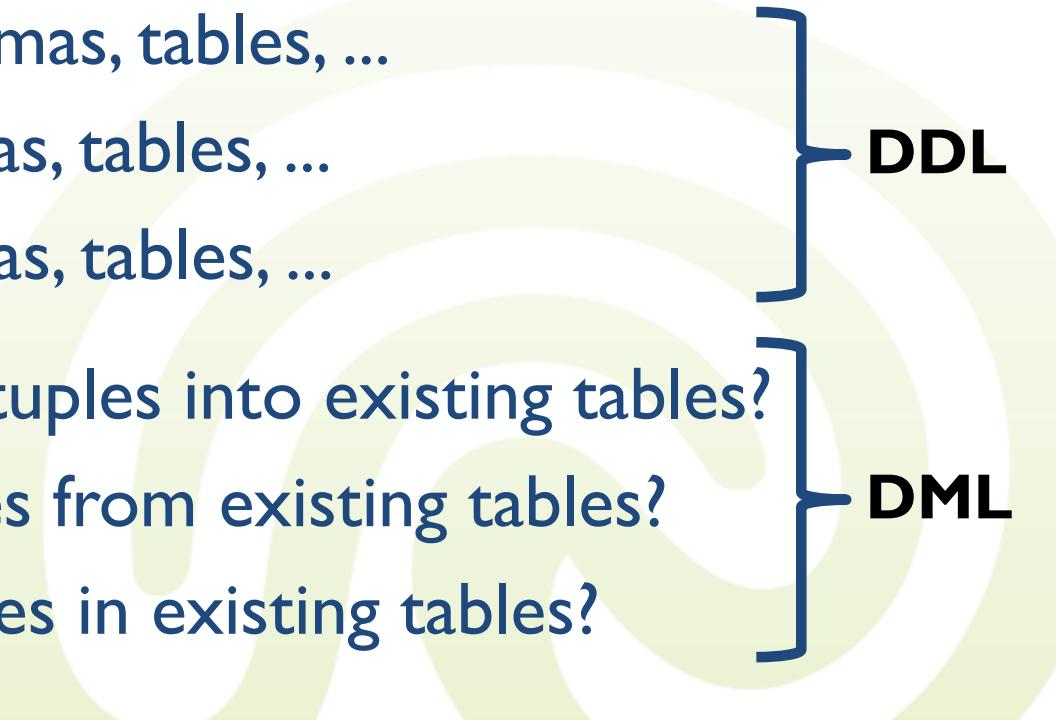


9. I Recap





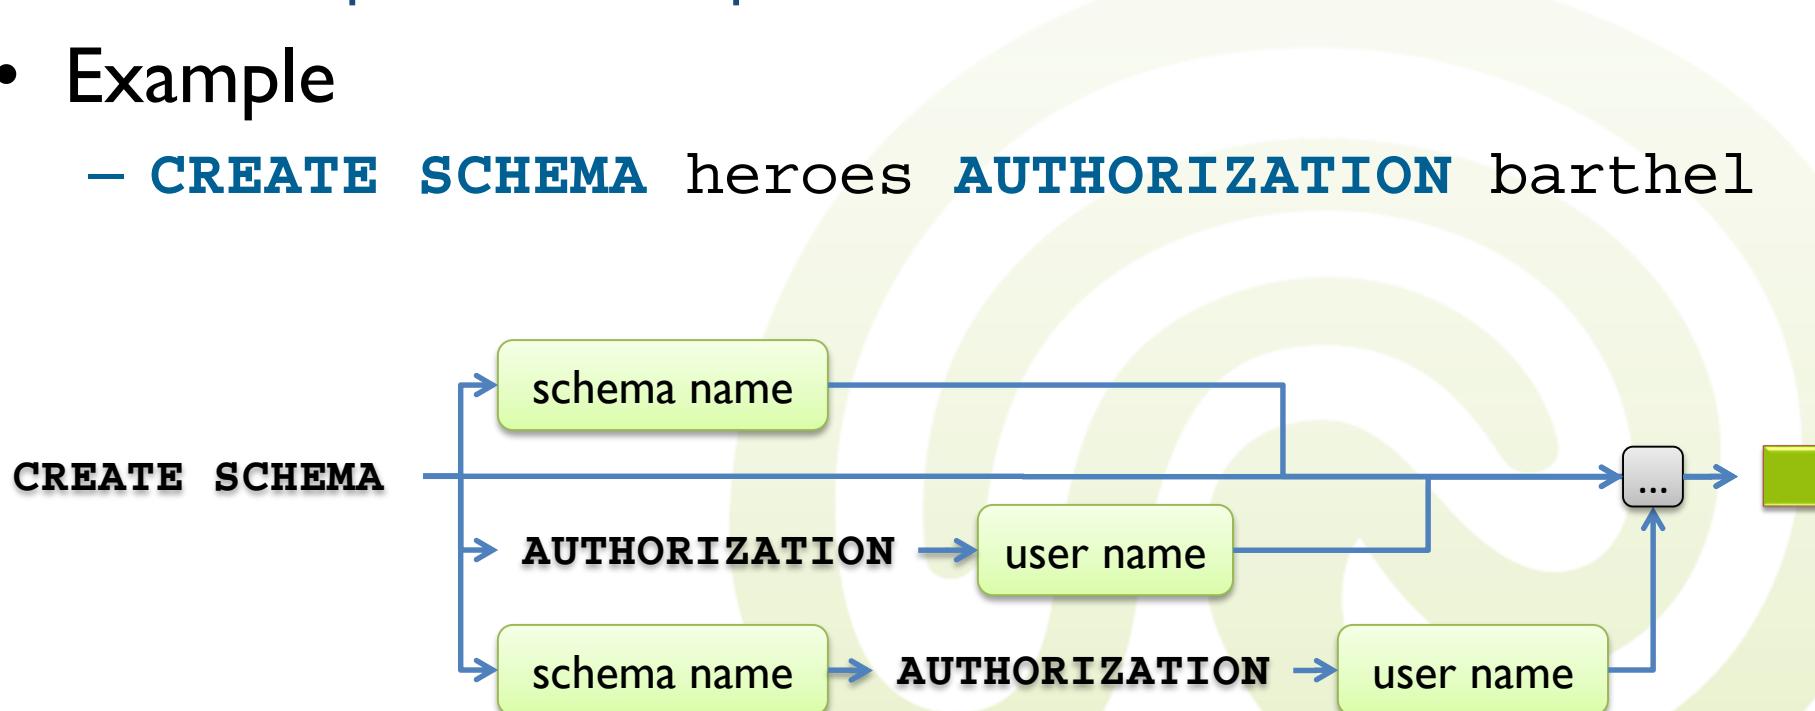
9.1 SQL DDL

- Last week, you learned how to query an existing relational database ...
 - What's missing?
 - How to **create** schemas, tables, ...
 - How to **drop** schemas, tables, ...
 - How to **alter** schemas, tables, ...
 - How to **insert** new tuples into existing tables?
 - How to **delete** tuples from existing tables?
 - How to **update** tuples in existing tables?
- 



9.1 SQL DDL: Schemas

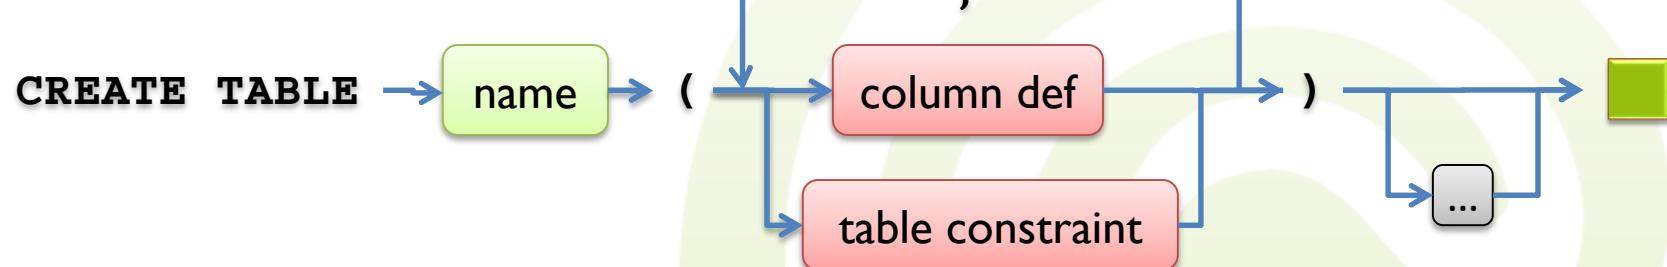
- **CREATE SCHEMA** creates a **new schema** with a given name for a given **owner**
 - If no schema name is provided, the **current username** is used
 - If no explicit owner is provided, also the **current user** is used
- Example
 - **CREATE SCHEMA heroes AUTHORIZATION barthel**





9.1 SQL DDL: Tables

- **CREATE TABLE** creates a **new table** with a given name
 - Contains column definition for **each column**
 - Contains additional table-specific **structural constraints**





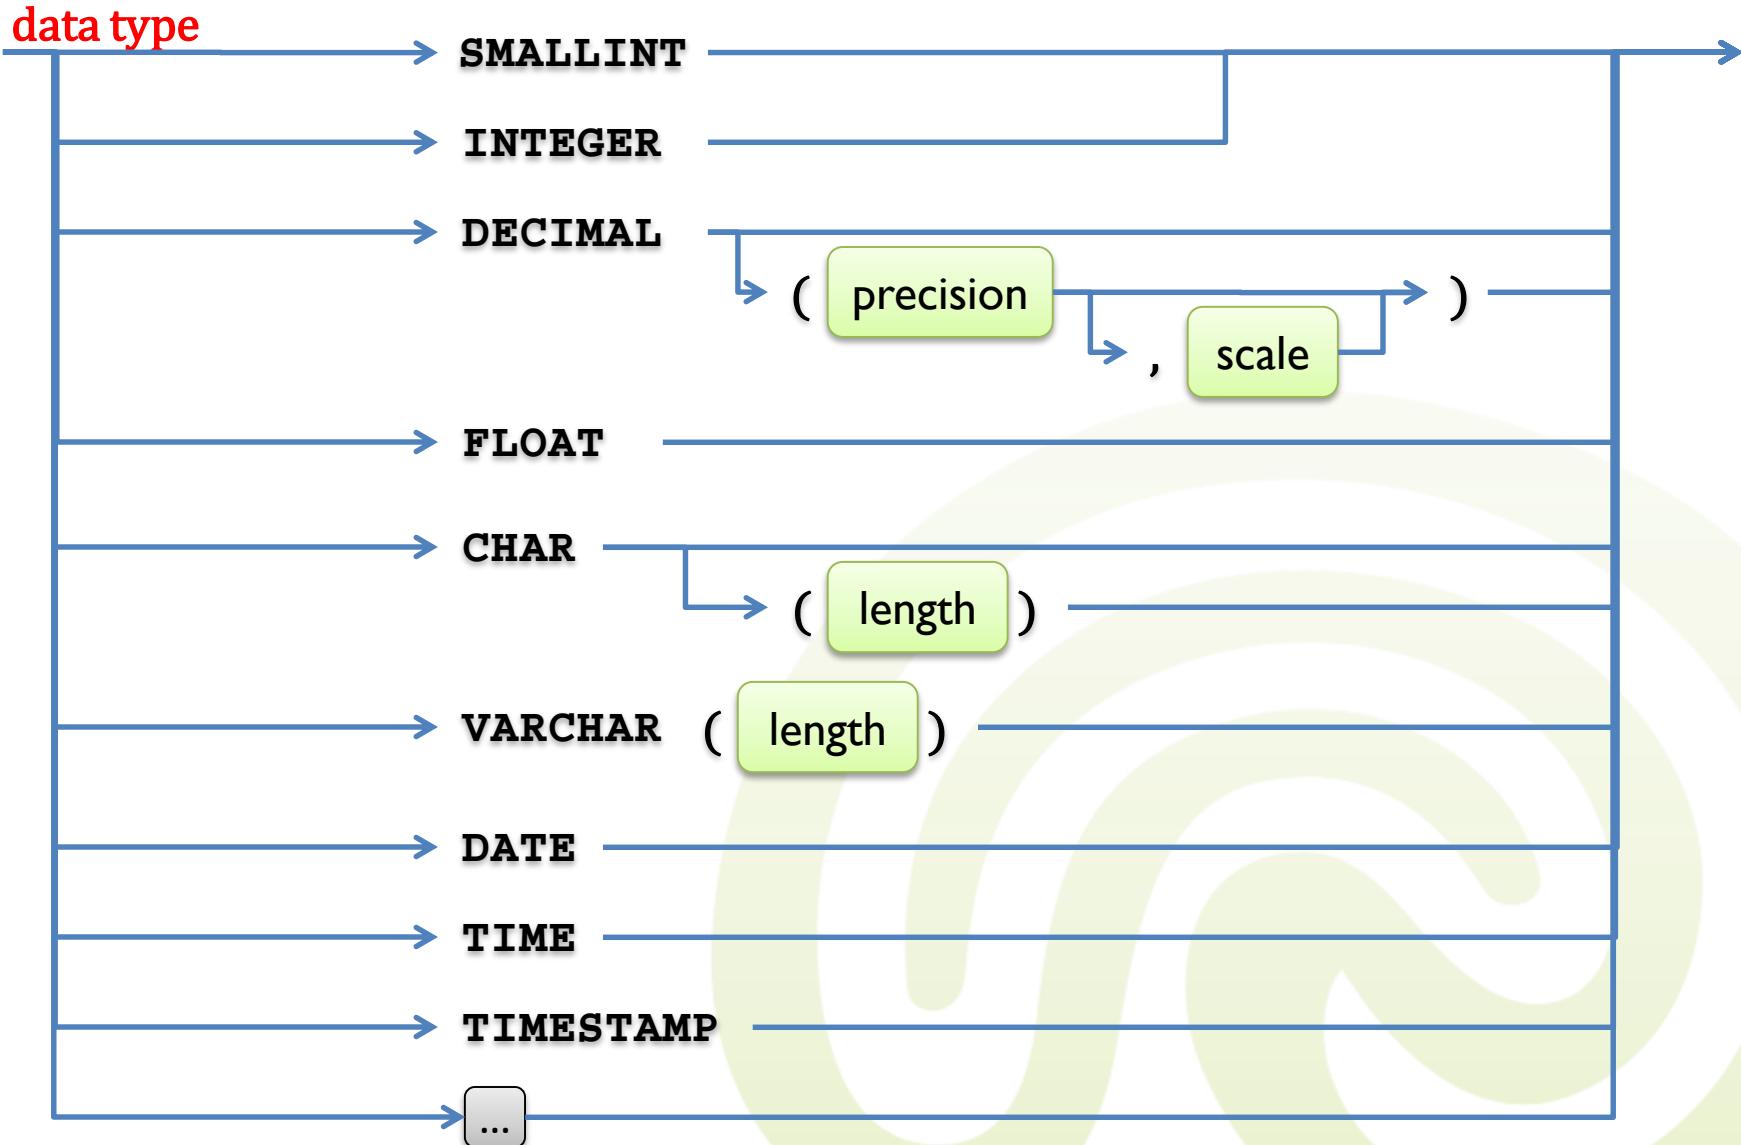
9.1 SQL DDL: Tables

- Each column has a **name** and a **data type**
- Each column may have multiple **column options**
- Example:
 - `CREATE TABLE person (
 name VARCHAR(200),
 age INTEGER
)`





9.1 SQL DDL: Tables



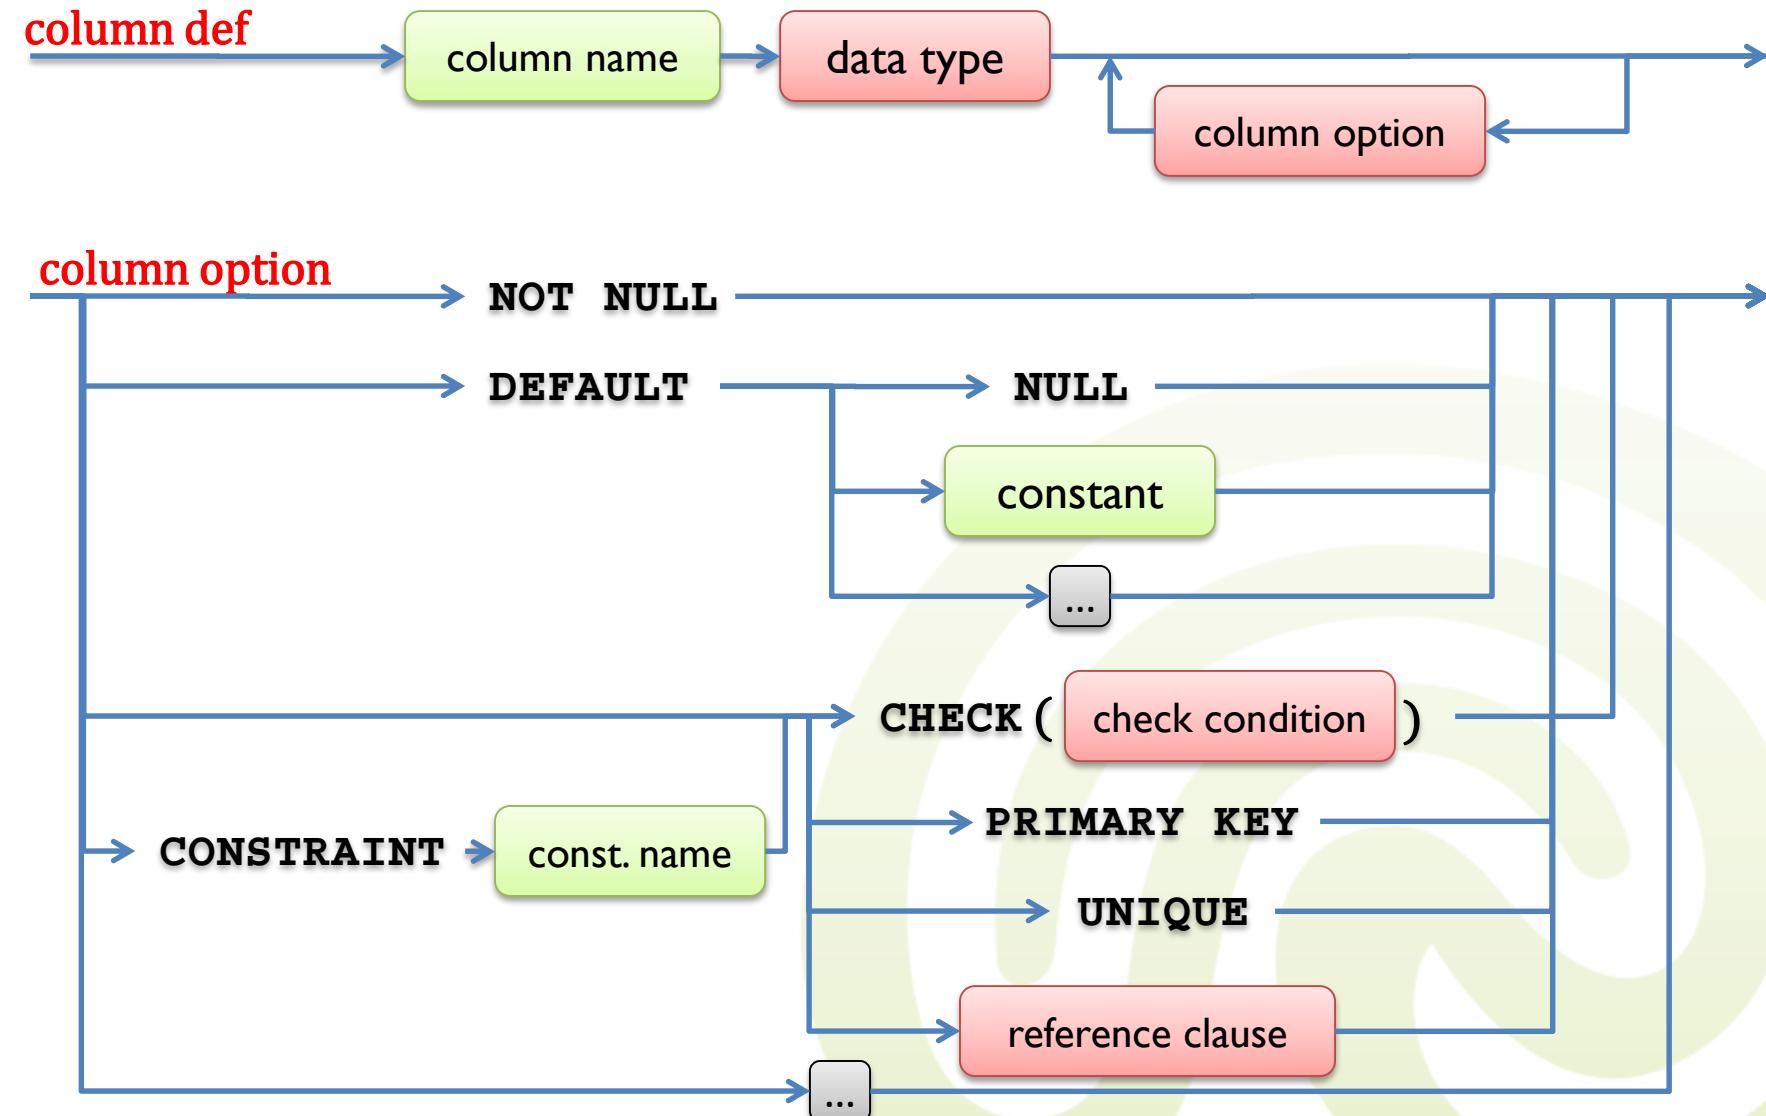


9.1 SQL DDL: Tables

Name	Syntax	description
Integer	INTEGER	Integer value
Float	FLOAT	Floating point number of approximate precision (the supported precision is implementation-dependent)
Numeric	NUMERIC (p , s)	A number with p digits before the decimal point and s digits after the decimal point
Character	CHAR (x)	A textual string of length x
Character varying	VARCHAR (x)	A textual string of length at most x
Date	DATE	Year, month, and day
Time	TIME	Hours, minutes, and seconds
Timestamp	TIMESTAMP	A date and a time



9.1 SQL DDL: Tables





9.1 SQL DDL: Tables

– NOT NULL:

The **NULL** value is not allowed for the column

– Example:

- `CREATE TABLE person (
 name VARCHAR(200) NOT NULL,
 age INTEGER NOT NULL
)`



– DEFAULT:

Defines the default value if a value is not explicitly set

- Usually a constant or **NULL**
- If omitted, **NULL** is the default



9.1 SQL DDL: Tables

- **Column constraints:**

- Restricts possible values for the current column
- May have a **unique name** indicated by
CONSTRAINT <name>
 - If name is omitted, system creates a default name
- **CHECK:** User-defined constraint.
To be valid, values have to satisfy the condition.
- **Example:**
 - `CREATE TABLE person (
 name VARCHAR(200),
 age INTEGER CONSTRAINT adult
 CHECK (age >= 18)
)`



9.1 SQL DDL: Tables

- **UNIQUE**: No duplicate values are allowed within this attribute
 - This option can only be used if the uniqueness constraints concerns only a single attribute
 - For multi-attribute uniqueness constraints, there is a different option (later)
 - Implies **NOT NULL**
 - Note: In DB2, **NOT NULL** has to be specified nevertheless ...

– Example:

- **CREATE TABLE** person (
 name **VARCHAR(200)** **NOT NULL UNIQUE**,
 age **INTEGER NOT NULL**
)



9.1 SQL DDL: Tables

- **PRIMARY KEY:** Each table may have a **primary key** (optionally, but recommended) made up of at least one column
 - This option can only be used if the primary key consists of only one column
 - For multi-column primary keys, you need a different option (later)
 - Implies **NOT NULL** and **UNIQUE**
 - Again, DB2 needs an explicit **NOT NULL**
- Additionally, a **referential clause** may be specified (see next slides)





9.1 SQL DDL: Referential Integrity

- Rows in tables may **refer** to rows in other tables to capture **relationships**
- Of course, you should not be allowed to refer to a non-existing row
 - **Referential integrity** between **primary keys** and **foreign keys** ensures that references are correct

Name	Reg. No.	Corres. No.	Case Book No.
Marks, Grace	1180		
County Kingston Penitentiary	City, Village or Township	Nativity	P.O. Address
May 4, 1852	England	Age	M. F. S. W.
Public Prison	Rate	Occupation	Religion
		Domestic	Epis.
Habits	Homicidal	Suicidal	Heredity
Cause	No. Attacks		Duration
Correspondent			
In Event of Death			
Recovered	Unconscious	Date Discharged	Date of Death
Aug. 18, 1853.			
Form 112-100-11-18 Hospital for Insane Series			
Library Bureau of Canada			

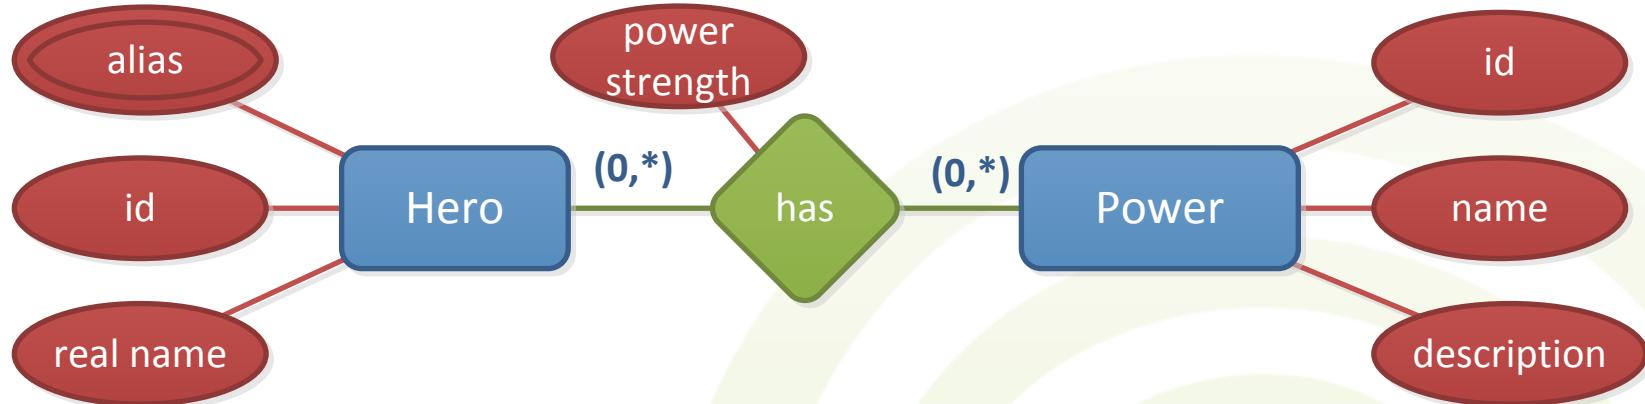




9.1 SQL DDL: Referential Integrity

Example:

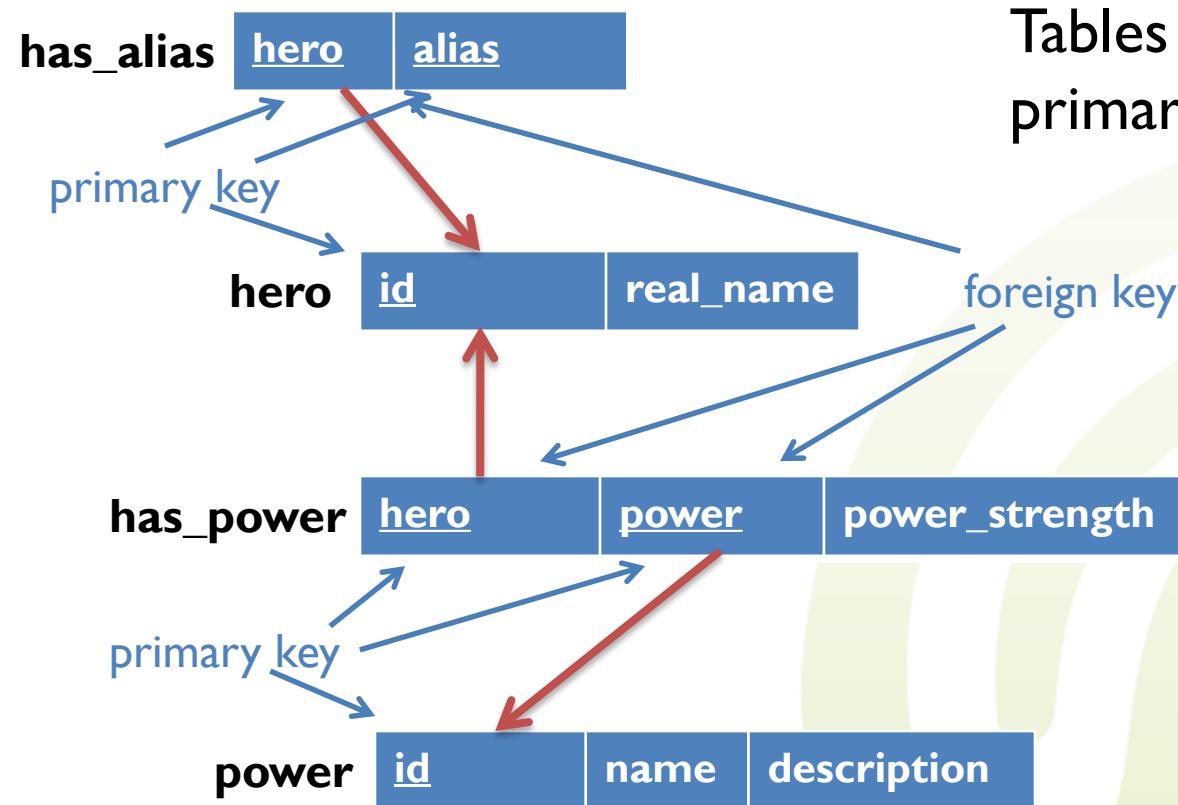
Conceptual ER schema:





9.1 SQL DDL: Referential Integrity

Resulting tables:



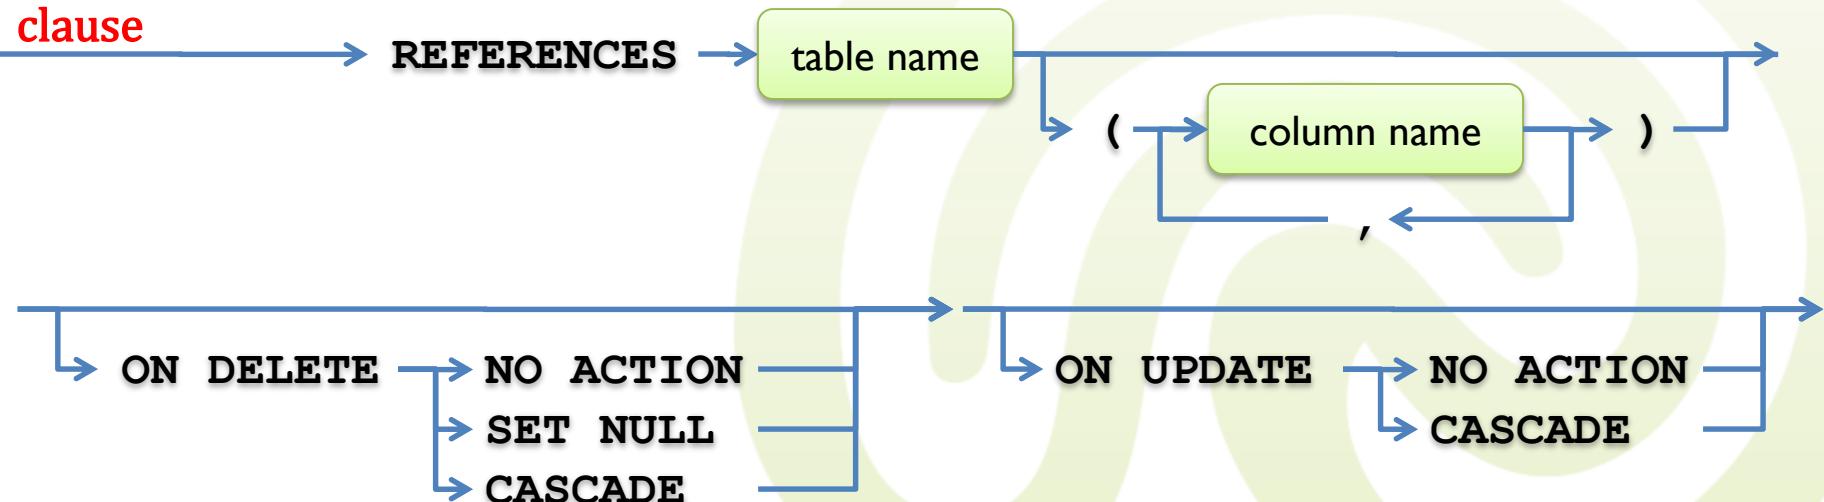
Tables refer to others by primary keys and foreign keys



9.1 SQL DDL: Referential Integrity

- Referential integrity can be defined using the **REFERENCES clause**
 - Either used by constraints in **column options** or within **table constraints**

**REFERENCES-
clause**





9.1 SQL DDL: Referential Integrity

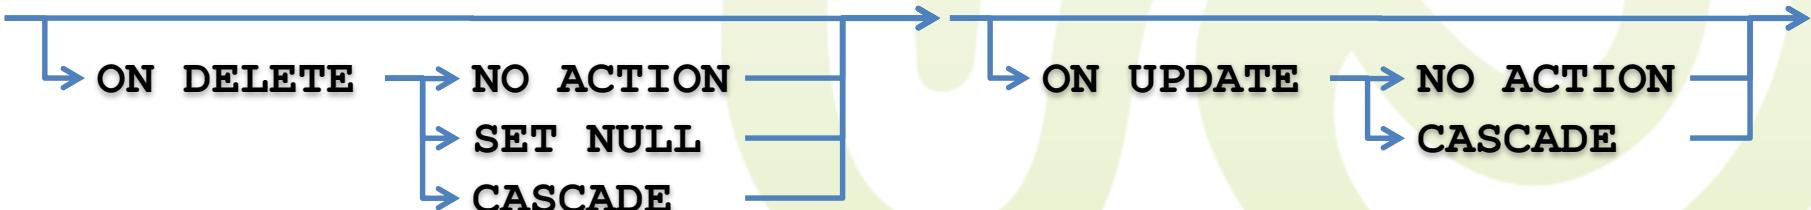
- **Example:**

- ```
CREATE TABLE employee(
 id INTEGER NOT NULL PRIMARY KEY,
 name VARCHAR(100) NOT NULL
)
```
- ```
CREATE TABLE managed_by(
    employee INTEGER NOT NULL
        REFERENCES employee,
    manager INTEGER NOT NULL
        REFERENCES employee
)
```



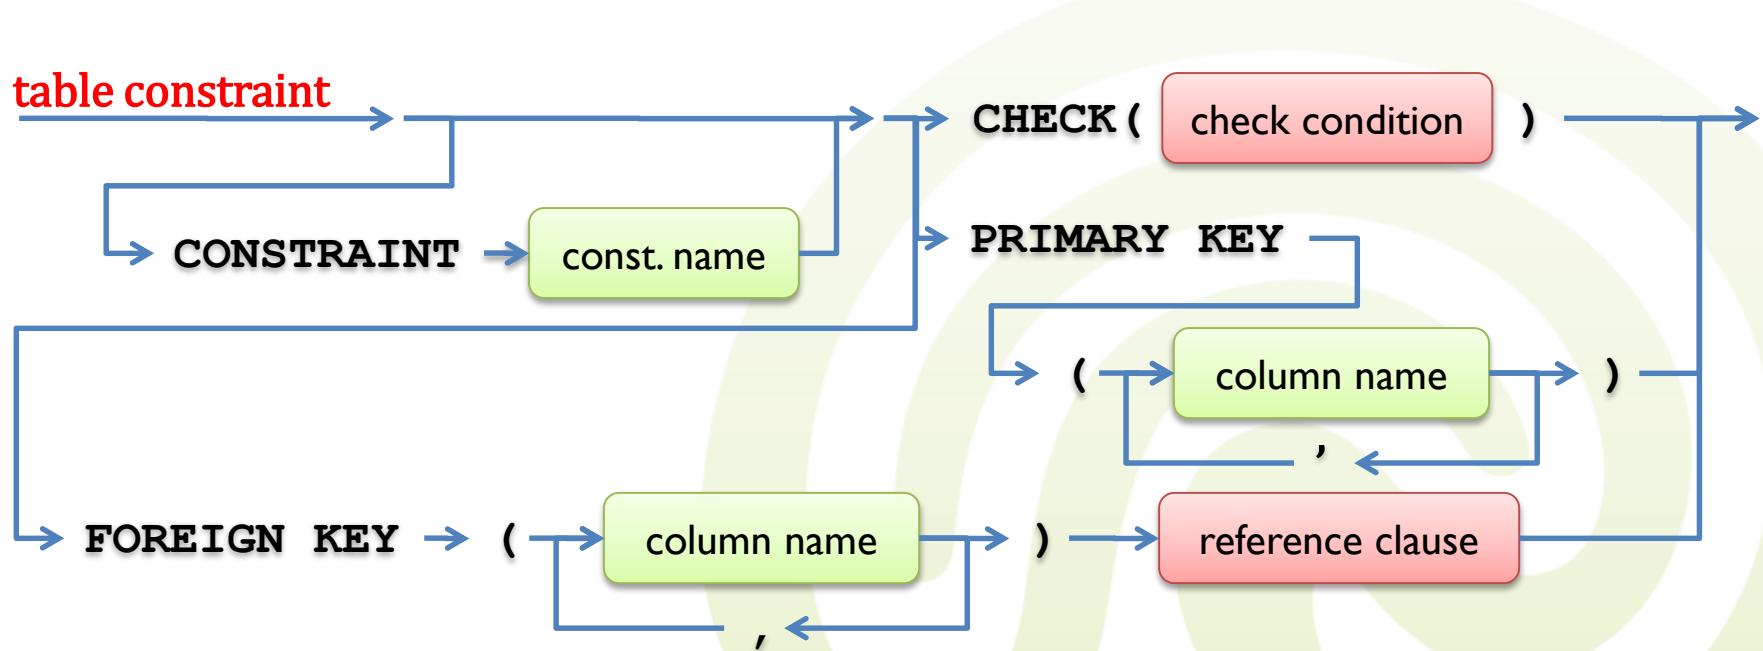
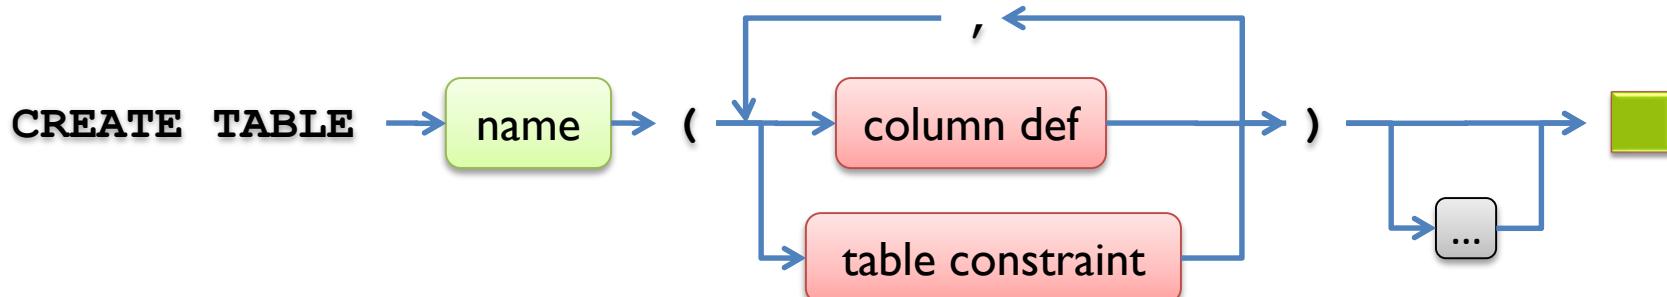
9.1 SQL DDL: Referential Integrity

- Optionally, you may specify what happens if a row that is referenced will be deleted or modified
 - **ON DELETE**: If a referenced row is deleted, ...
 - **NO ACTION**: ...reject the deletion (that is, it cannot be performed)
 - **SET NULL**: ...delete it and set all referencing foreign keys to **NULL**
 - **CASCADE**: ...delete it along with all rows referring to it
 - **ON UPDATE**: If the primary key of a referenced row is modified, ...
 - **NO ACTION**: ...reject the modification (that is, it cannot be performed)
 - **CASCADE**: ...change all values of referencing foreign keys
 - Default:
 - **ON DELETE NO ACTION ON UPDATE NO ACTION**





9.1 SQL DDL:Table Constraints



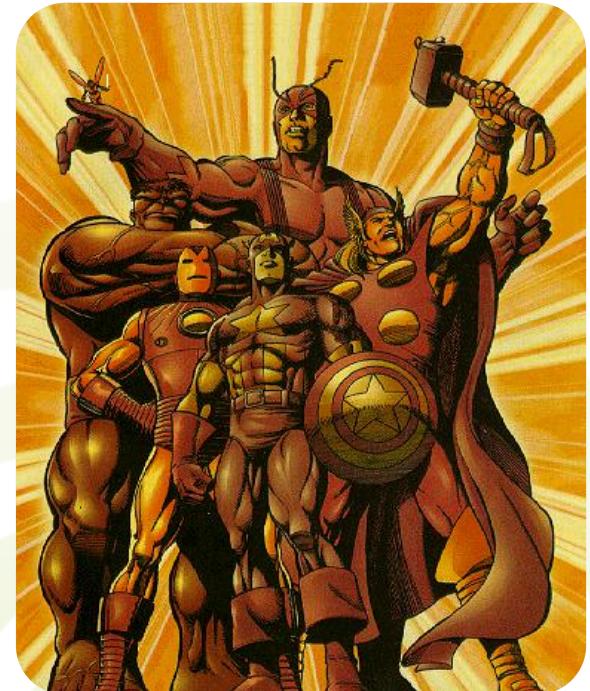
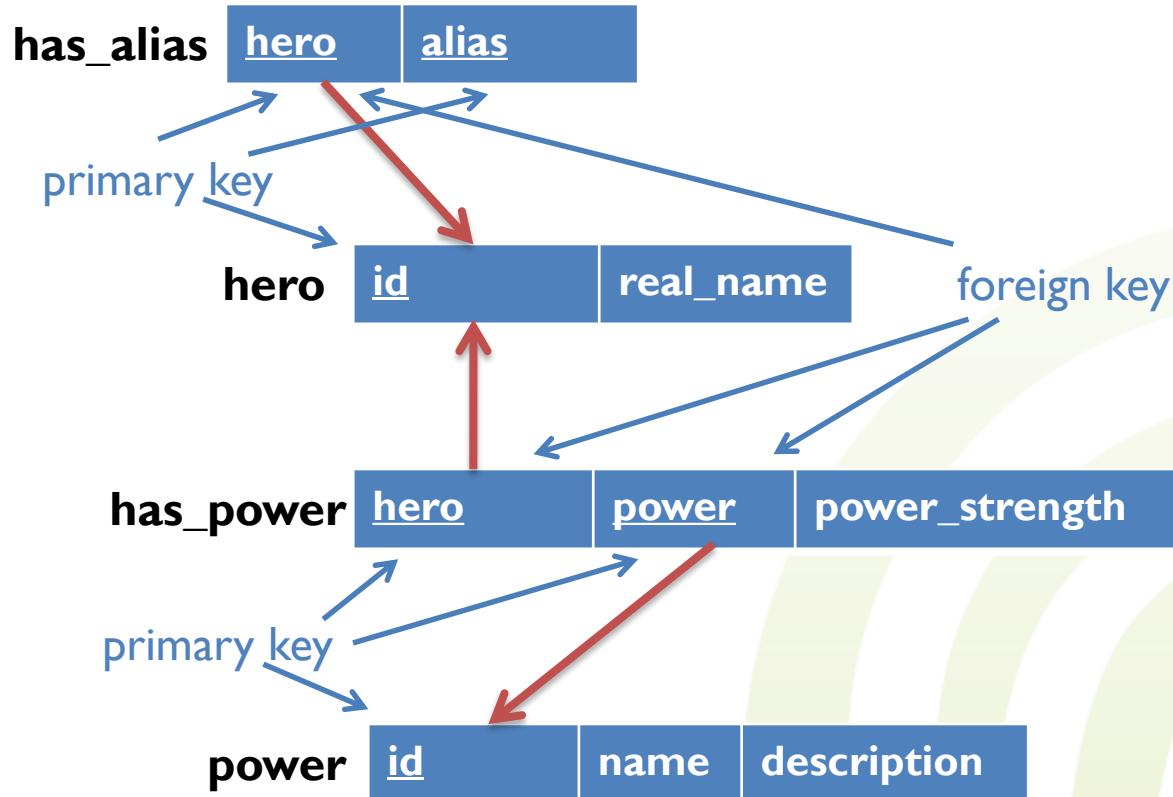


9.1 SQL DDL: Table Constraints

- Table constraints behave similar to constraints in column options
 - If no name is provided, a **name is automatically generated**
 - The **CHECK condition** may contain any Boolean predicate
 - In contrast to column options, table constraints may declare **primary keys** consisting of **multiple attributes**
 - **Foreign keys** declare references to primary keys of other tables
 - See referential integrity



9.1 SQL DDL: Table Example





9.1 SQL DDL:Table Example

```
CREATE TABLE hero(  
    id INTEGER NOT NULL PRIMARY KEY,  
    real_name VARCHAR(100)  
)
```

```
CREATE TABLE power(  
    id INTEGER NOT NULL PRIMARY KEY,  
    name VARCHAR(100),  
    description VARCHAR(255)  
)
```

```
CREATE TABLE has_alias (
```

```
    hero INTEGER REFERENCES hero  
        ON DELETE CASCADE  
        ON UPDATE CASCADE,  
    alias VARCHAR(100) NOT NULL,  
    PRIMARY KEY (hero, alias)  
)
```

link has_alias to hero

delete alias if hero is deleted

update alias if hero is updated

composed primary key



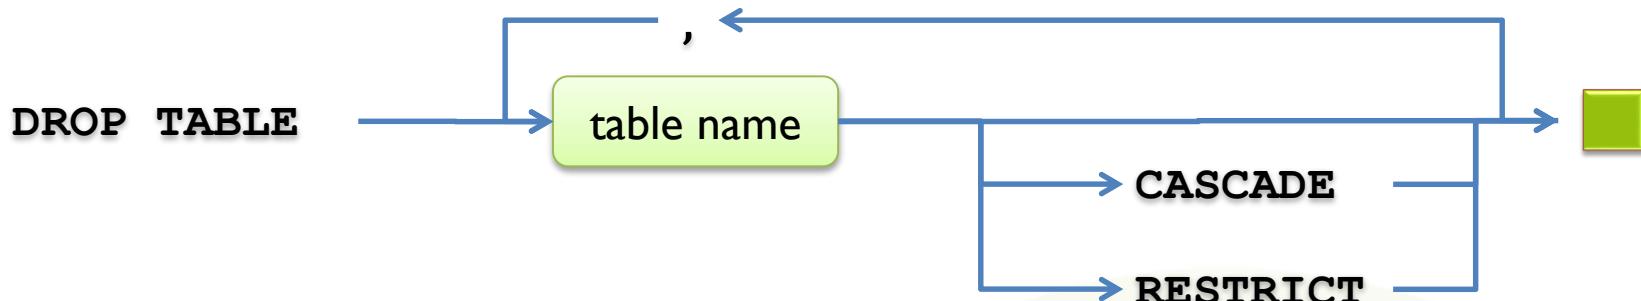
9.1 SQL DDL: Table Example

- `CREATE TABLE has_power(`
 `hero INTEGER NOT NULL,`
 `power INTEGER NOT NULL,`
 `power_strength INTEGER NOT NULL,`
 `PRIMARY KEY (hero, power),`
 `FOREIGN KEY (hero) REFERENCES hero`
 `ON DELETE CASCADE`
 `ON UPDATE CASCADE,`
 `FOREIGN KEY (power) REFERENCES power`
 `ON DELETE CASCADE`
 `ON UPDATE CASCADE`
)



9.1 SQL DDL: Drop Tables

- For **deleting** tables, there is the **DROP TABLE** command

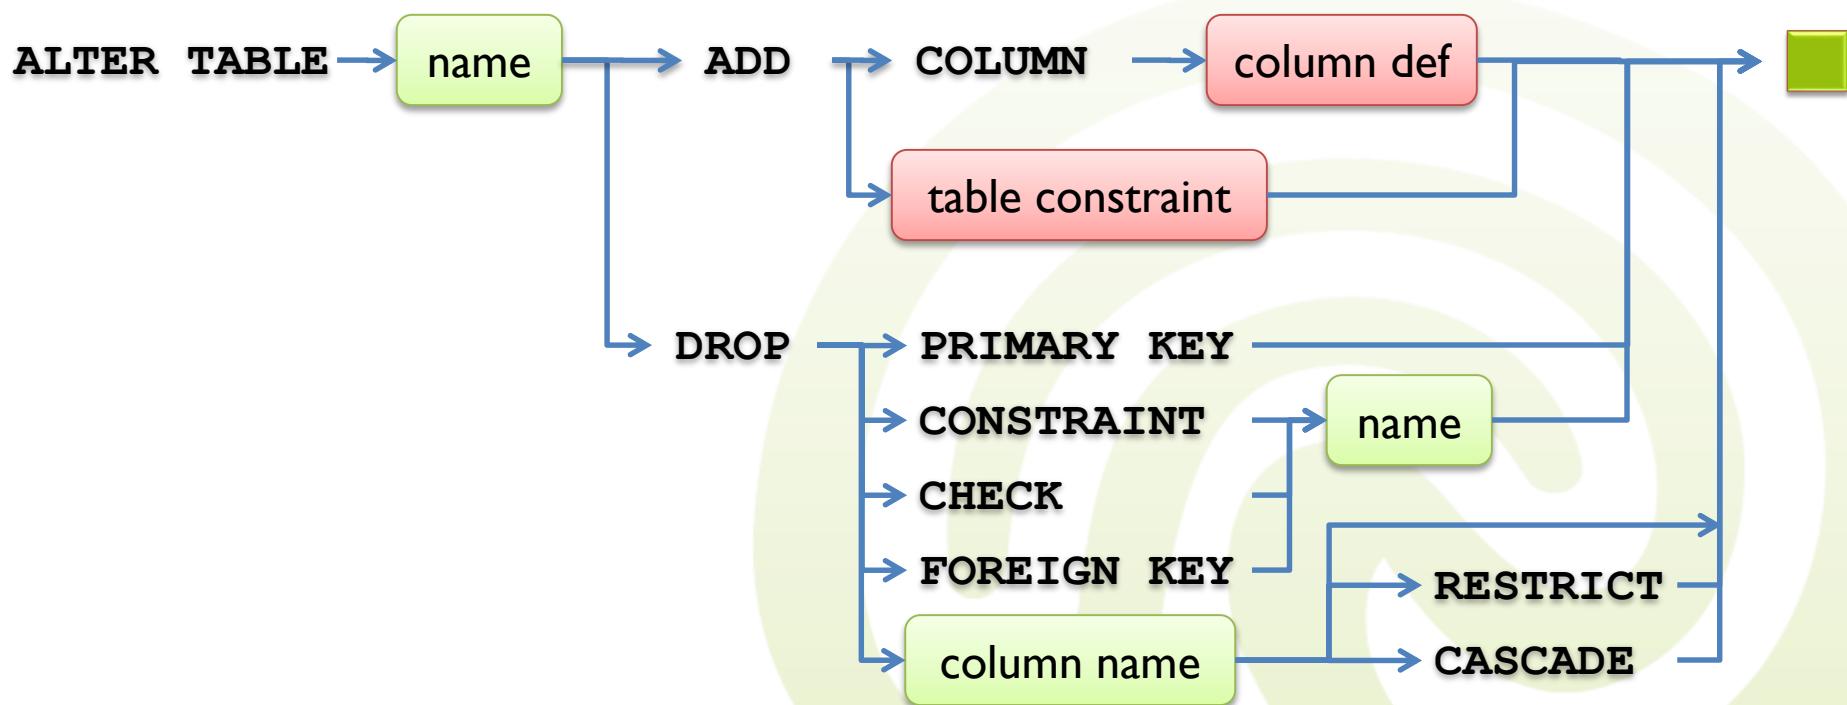


- If **RESTRICT** is used, you may only drop empty tables that are not referenced by any other table
- If **CASCADE** is used, all referencing tables are also deleted (including all stored rows)
 - DB2 does not support **CASCADE** ...
- If neither is used, the table does not have to be empty, but must not be referenced by another one
- Example
 - DROP TABLE hero CASCADE, power CASCADE**



9.1 SQL DDL: Alter Tables

- After a table has been created, you may **alter** it by adding/removing **columns** or **constraints**





9.1 SQL DDL: Alter Tables

- If you add a **new column** with a **NOT NULL** constraint, you also need to provide a **default value**
- When **dropping** a **column**, you must either choose
 - **CASCADE** to also delete any views, indexes, and constraints dependent on that column
 - **RESTRICT** to allow the drop only if there is no referring column (default)
- If the **name** of a constraint is **auto-generated**, you need to look it up in the **system catalog**
- **Example:**
 - **ALTER TABLE** has_power **DROP** **power_strength**
 - **ALTER TABLE** has_power
ADD COLUMN **since** **DATE**



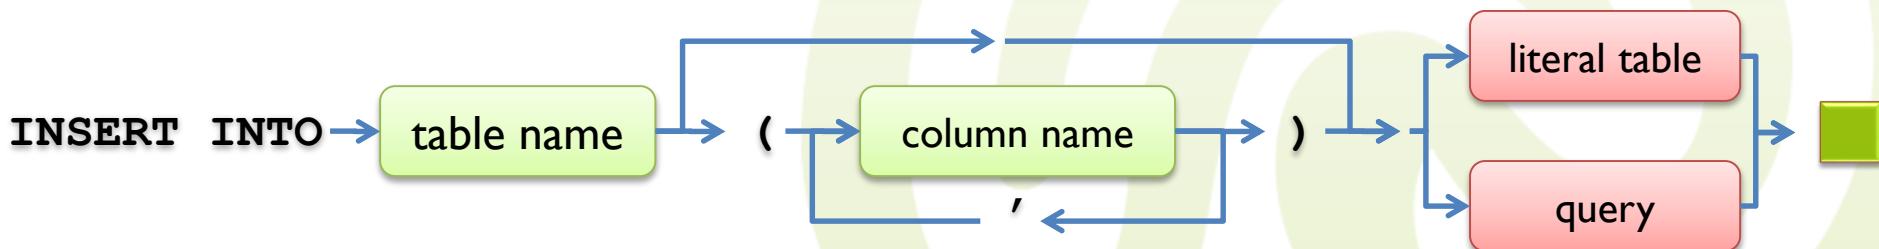
9.2 SQL DML

- **Data definition language (DDL)**
 - Creating, changing, altering schemas, tables, ...
 - CREATE SCHEMA
 - CREATE TABLE
 - ALTER TABLE
 - DROP TABLE
- **Data manipulation language (DML)**
 - Querying
 - SELECT
 - Adding and updating data
 - INSERT INTO
 - UPDATE
 - DELETE



9.2 SQL DML: Insert

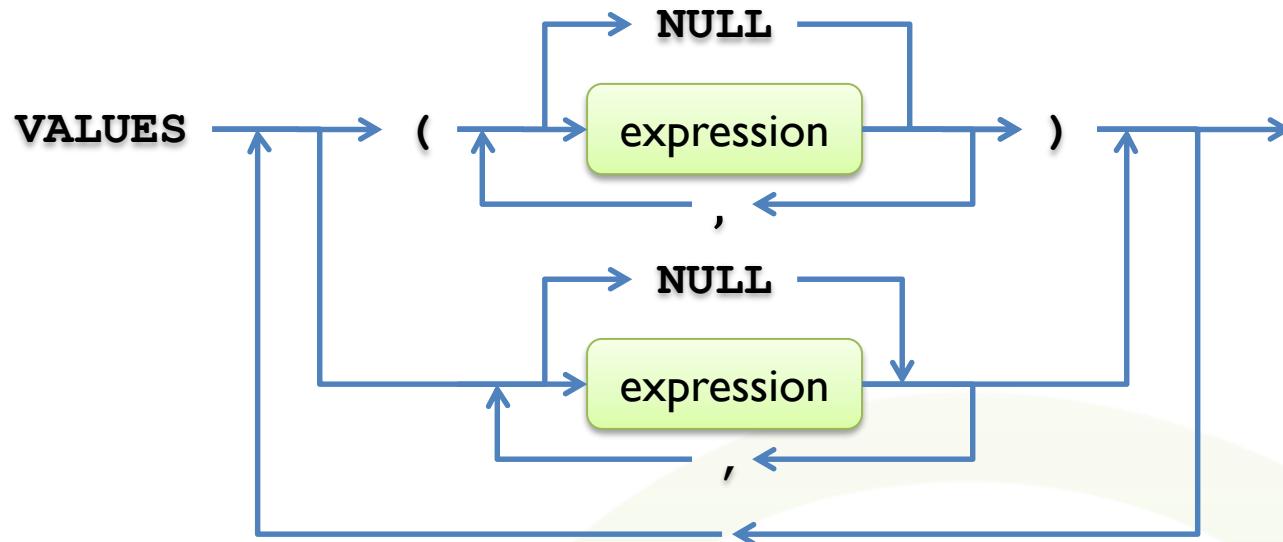
- Now we have wonderful, empty tables
- We need to put data into them!
 - **INSERT INTO** statement
 - You can specify into what **columns** you want to insert data
 - Default: All columns
 - New values are stated as a **literal table** or **inline view (query)**
 - Of course the **attribute domains** have to match





9.2 SQL DML: Insert

literal table



- A **literal table** is defined extensionally:

`VALUES ('James', 'Howlet')`

James	Howlet
-------	--------

`VALUES ('Charles', 'Xavier'), ('Jean', 'Grey')`

Charles	Xavier
Jean	Grey

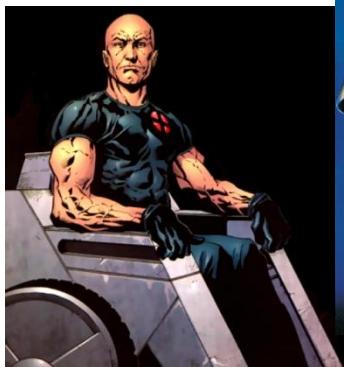
`VALUES 'Wolverine', ('Professor X'), 'Phoenix'`

Wolverine
Professor X
Phoenix



9.2 SQL DML: Insert

- `INSERT INTO hero(id, real_name) VALUES
(1, 'Charles F. Xavier'),
(2, 'Jean Grey')`
`INSERT INTO has_alias VALUES
(1, 'Professor X'),
(1, 'Onslaught'),
(2, 'Phoenix'),
(2, 'Marvel Girl')`





9.2 SQL DML: Insert

- Of course, subqueries may also be used in **INSERT** statements

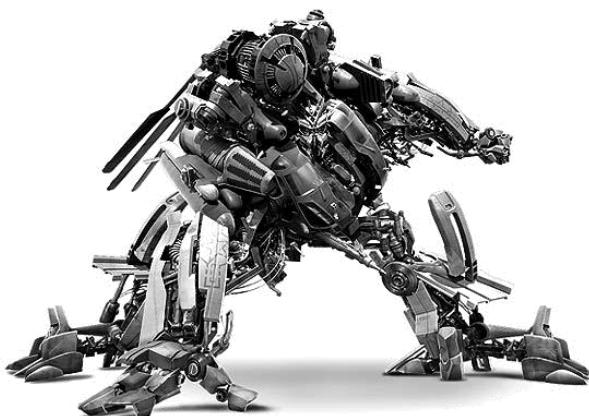
- `INSERT INTO heroes_starting_with_a(`
`SELECT * FROM hero`
`WHERE real_name LIKE 'A%`
`)`





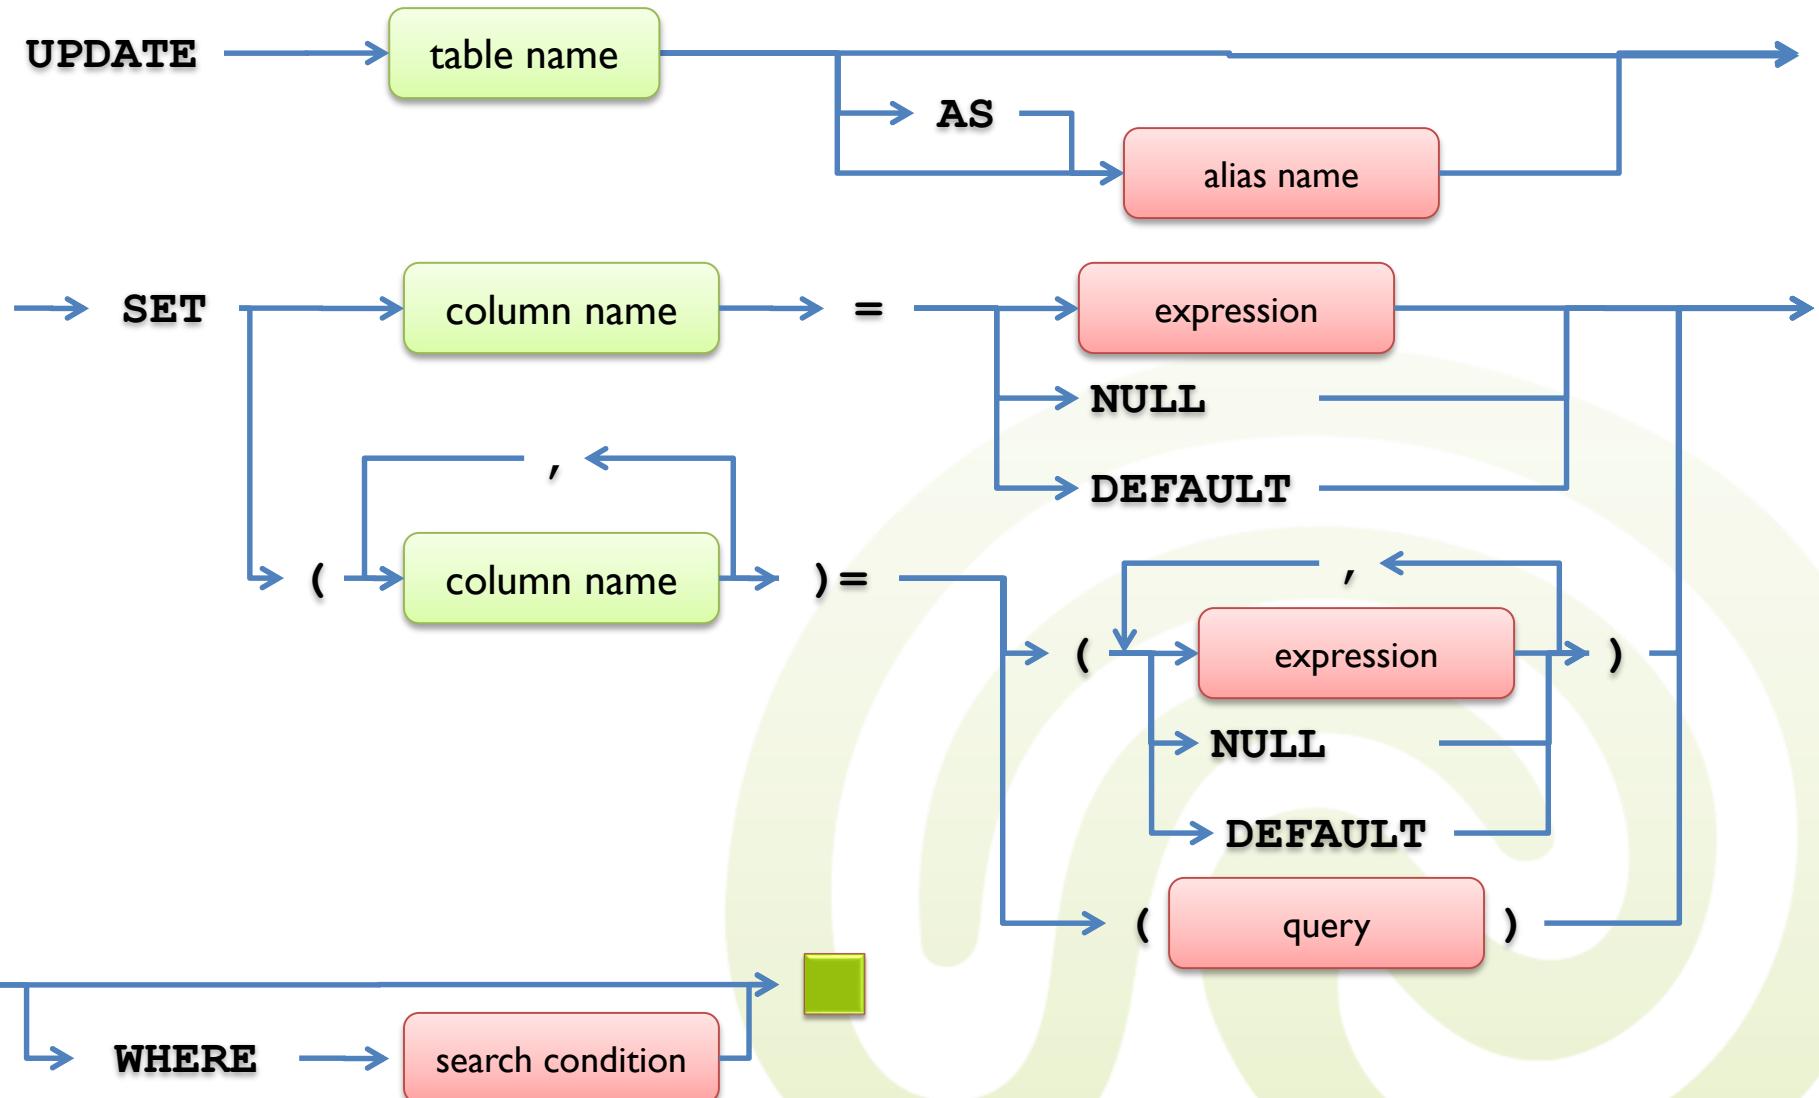
9.2 SQL DML: Update

- Existing rows can also be changed using the **UPDATE statement**
 - Very similar to the **SELECT** statement
 - Update **finds rows** fulfilling a **given condition** and **changes** some of its **rows** by assigning **new values**





9.2 SQL DML: Update





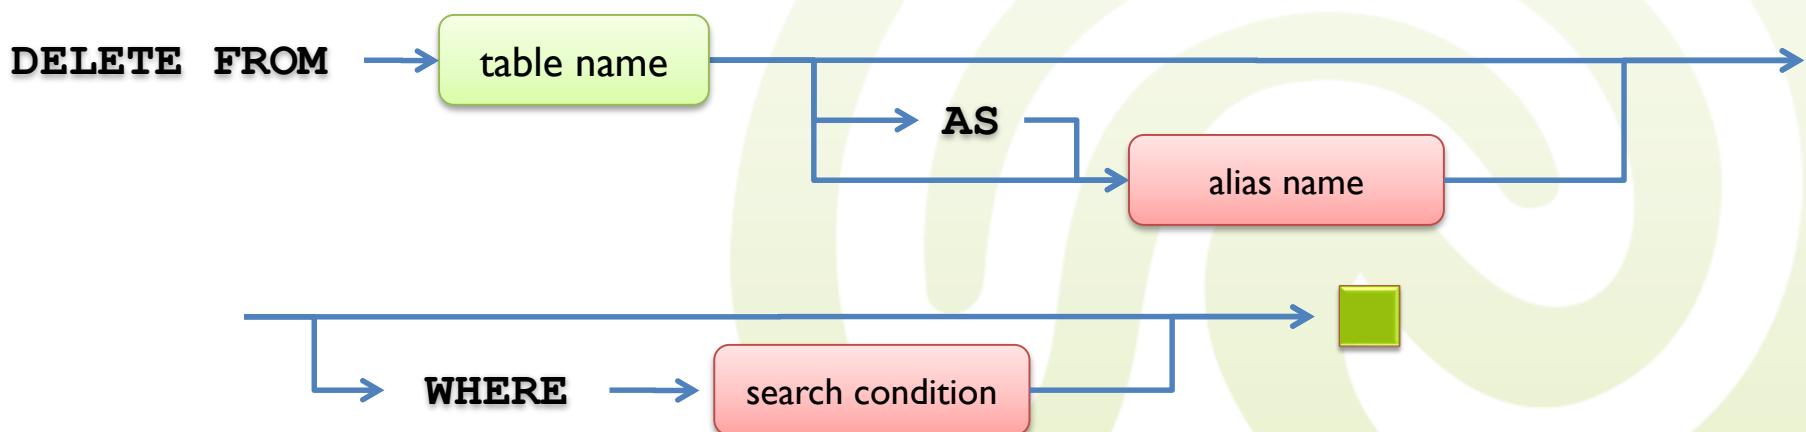
9.2 SQL DML: Update

- Hide the real name of each hero
 - `UPDATE hero SET real_name = NULL`
- Multiply all `power_strength` values by 10
 - `UPDATE has_power SET power_strength = power_strength * 10`
- Change the name of hero with id 1
 - `UPDATE hero SET name= 'Charles Francis Xavier' WHERE id = 1`
- Change name and id of Jean Grey
 - `UPDATE hero SET (id, name) = ('3', 'Jean Grey-Summers') WHERE name = 'Jean Grey'`
 - Change of id is propagated to other tables when `ON UPDATE CASCADE` is used in table definition
- Again, subqueries can be used in the `WHERE` clause



9.2 SQL DML: Delete

- The **DELETE statement** is used to delete rows from a table
 - Deletes all rows satisfying a certain search condition
 - **Example:**
 - Delete Jean Grey
 - `DELETE FROM hero WHERE name = 'Jean Grey'`
 - Delete all heroes
 - `DELETE FROM hero`

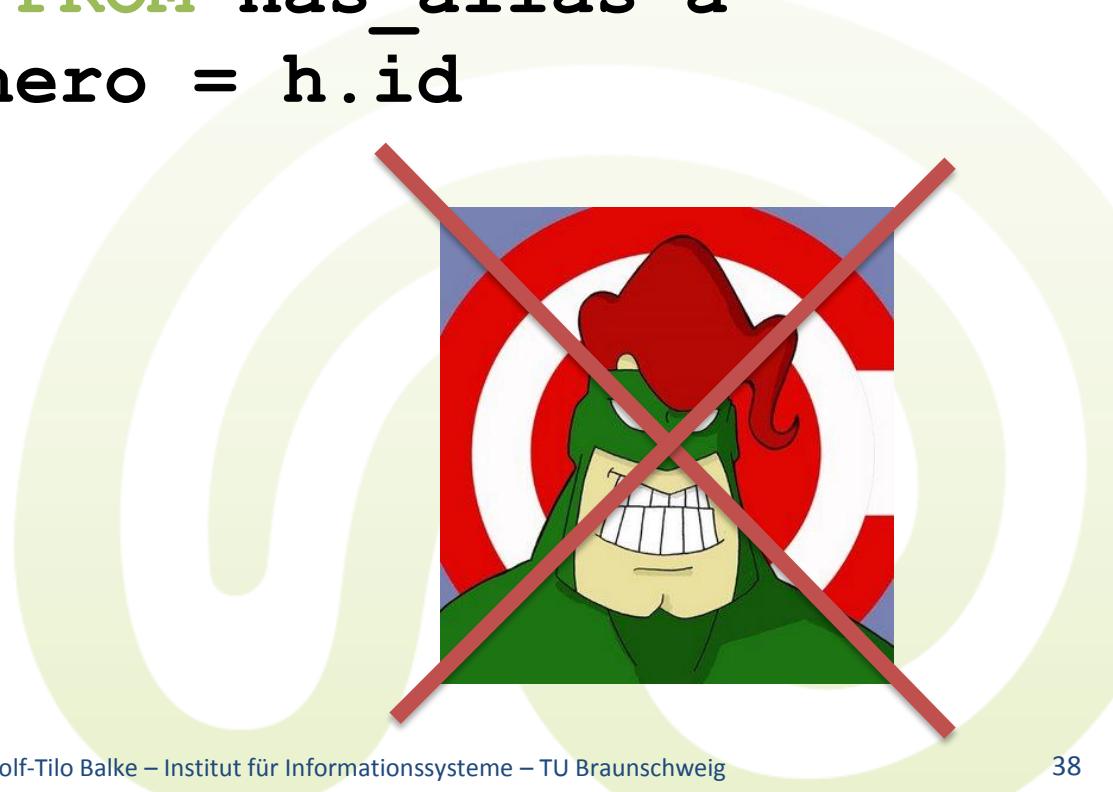




9.2 SQL DML: Delete

- Again, subqueries can be used here

- ```
– DELETE FROM hero h
 WHERE NOT EXISTS (
 SELECT * FROM has_alias a
 WHERE a.hero = h.id
)
```





## 9.3 SQL ≠ SQL

*Detour*

- First, the **good news**:
  - SQL has been standardized by the ISO in 1987
  - The standard is well-maintained and under active development
    - SQL-86, SQL-89, **SQL-92**, **SQL:1999**, **SQL:2003**, SQL:2006, SQL:2008
  - Many “big” database vendors participate in the standardization process:
    - IBM, Ingres, Microsoft, Oracle, Sybase, ...





## 9.3 SQL ≠ SQL

*Detour*

### A timeline of SQL standardization:

- 1986
  - ANSI SQL
  - Relations, attributes, views
  - **SELECT ... FROM ... WHERE ...**
- 1987
  - SQL-86 (ISO/IEC 9075:1986)
- 1989
  - SQL-89 (SQLI)
    - ≈ SQL-86 + restricted referential integrity





## 9.3 SQL ≠ SQL

*Detour*

- 1992
  - SQL-92 (SQL2)
  - 3 parts, 1120 pages
  - Entry Level
    - ≈ SQL-89 + **CHECK** (attribute)
  - Intermediate Level
    - ⊇ Entry Level + domains, **CHECK** (relation),  
**CASE**, **CAST**, **JOIN**, **EXCEPT**, **INTERSECT**
  - Full Level
    - ⊇ Intermediate Level + assertions, nested select,  
nested from



# 9.3 SQL ≠ SQL

*Detour*

- 1999/2000
  - SQL:1999 (SQL3)
  - 5 parts, 2084 pages
  - ≈ SQL-92 + object-orientation, **recursive queries**, triggers, OLAP, user-defined types, regular expressions
  - Computationally complete, object-oriented database programming language, descriptive and procedural
  - Core (about 180 features)
    - ≈ SQL92 Entry Level + parts of Intermediate and Full Level
  - 9 Packages (about 240 features)
    - enhanced date/time, enhanced integrity, OLAP, PSM, CLI, basic object support, enhanced object support, trigger, SQL/MM



## 9.3 SQL ≠ SQL

*Detour*

- Recursive queries (SQL:1999):
  - How to find all rivers that flow into the North Sea?
  - Multiple joins?

| flows_into | river  | mouth     |
|------------|--------|-----------|
|            | Oker   | Aller     |
|            | Aller  | Weser     |
|            | Weser  | North Sea |
|            | Elbe   | North Sea |
|            | Edder  | Wietze    |
|            | Flöth  | Wietze    |
|            | Wietze | Aller     |
|            | Isar   | Danube    |
|            | Inn    | Danube    |



# 9.3 SQL ≠ SQL

*Detour*

- Solution:

```
WITH RECURSIVE flows_into_ns(river) AS (
 SELECT river
 FROM flows_into
 WHERE mouth = 'North Sea'
 UNION
 SELECT fi.river
 FROM flows_into AS fi
 JOIN flows_into_ns AS fins
 ON fi.mouth = fins.river
)
SELECT river FROM flows_into_ns
```

- Note: DB2 uses a different syntax
  - No **RECURSIVE**, **UNION ALL** instead of **UNION**, no explicit **JOIN**



# 9.3 SQL ≠ SQL

*Detour*

- 2003
  - SQL:2003
  - 14 parts, 3606 pages
  - **MULTISET** as an explicit construct (with numerous operations, such as: **MULTISET UNION**, **MULTISET EXCEPT**, **MULTISET INTERSECT**, **CARDINALITY**)
  - Sequence generators
    - **CREATE SEQUENCE <sequence name> AS <type name> [START WITH <value>] [INCREMENT BY <value>] [NO MINVALUE | MINVALUE <value>] [NO MAXVALUE | MAXVALUE <value>] [NO CYCLE | CYCLE]**
  - Base type **XML** for mappings between SQL and XML



## 9.3 SQL ≠ SQL

*Detour*

- 2006
  - SQL:2006
  - Successor of SQL:2003
  - A new extension for XML handling
    - Importing, storing, querying, and manipulating XML data
    - Support for XQuery
    - Concurrently access (object-)relational data and XML data
- 2008
  - SQL:2008
  - Successor of SQL:2006
  - “Maintenance release” (some new but minor features)



## 9.3 SQL ≠ SQL

*Detour*

- Well, here are the **bad news**:
  - There are still **too many variants** of SQL  
**(both syntactic and semantic differences)**
    - True application portability remains a challenge
  - The standard has been used to introduce  
**two kinds of features**:
    1. Features that are well-understood and widely implemented
    2. New and largely untried technologies, hoping that vendors follow the lead and deliver new functionalities
  - **Vendors don't care** too much about the standard



## 9.3 SQL ≠ SQL

*Detour*

- A **common myth** among software developers:



If your application uses  
only standard SQL,  
then it is portable.

- If you don't believe me, here are some examples ...



## 9.3 SQL ≠ SQL

*Detour*

- ```
CREATE TABLE name (
    first VARCHAR(100),
    middle VARCHAR(100),
    last VARCHAR(100)
)
```

```
INSERT INTO name VALUES ('George', 'Walker', 'Bush')
INSERT INTO name VALUES ('Horst', '', 'Kr')
INSERT INTO name VALUES ('Angela', NULL, 'Merkel')
```

- ' '(empty string) means that we know that there is no middle name
- **NULL** means that we don't know whether there is a middle name
- Sounds like a good design? What do you think?
 - According to the SQL standard, this approach is fine ...
 - ... unless your RDBMS is Oracle (' ' is the same as **NULL**)



9.3 SQL ≠ SQL

Detour

- What about terminology?
 - The SQL standard defines the following notions:
 - Environment
 - Cluster
 - Catalog
 - Schema
 - The reality:
 - Database server
 - (unsupported)
 - Database
 - Schema
 - But attention:
 - In MySQL, there are no catalogs, “schema” and “database” are synonyms
 - In Oracle, there is exactly one schema per user;
CREATE/ALTER SCHEMA x <command> executes **<command>** on all objects located in schema x



9.3 SQL ≠ SQL

Detour

- The **statement terminator** “;”
 - According to the SQL standard, (almost) every SQL statement has to be terminated by a semicolon
 - What's happening in practice?
 - Many RDBMS treat the terminator as being optional (which is fine, but may cause some problems)
 - Some RDBMS either strictly require a terminator or complain if it is present
 - In some RDBMS, this behavior can be configured ...
 - Summary:

No matter what you do, it causes problems!



9.3 SQL ≠ SQL

Detour

- The **BOOLEAN** data type

- ```
CREATE TABLE customers (
 id INTEGER PRIMARY KEY,
 name VARCHAR(100),
 is_vip BOOLEAN,
 is_blacklisted BOOLEAN
)
```

```
SELECT id, name FROM customers
WHERE is_vip AND NOT is_blacklisted
```

- Practice?

- Not supported by Oracle, DB2, and MS SQL Server
  - Official workarounds: Use CHAR or INTEGER ...
- Supported by MySQL and PostgreSQL
  - Where in MySQL **BOOLEAN** is just a short hand for **TINYINT(1)** ...



## 9.3 SQL ≠ SQL

*Detour*

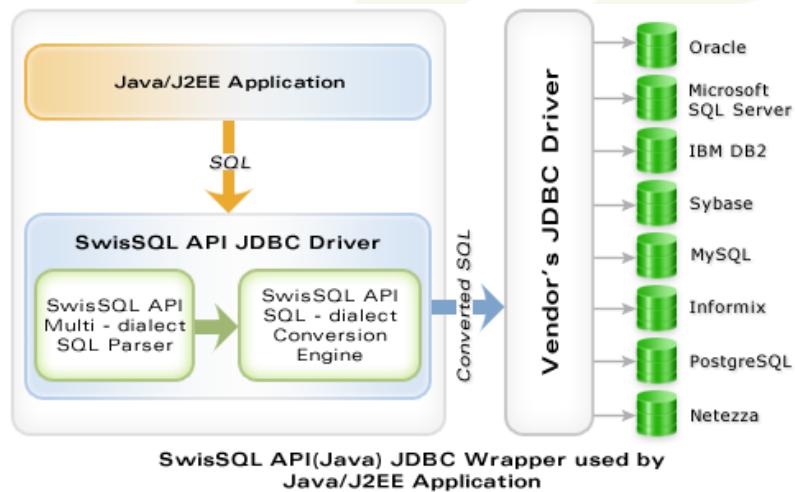
- **Summary**
  - **SQL is not SQL**
  - In some cases, even **identical SQL code works differently** on different RDBMS
- **Current trends?**
  - Open-source RDBMS (PostgreSQL, MySQL, Derby, ...) typically try to adhere to the standard
    - However, many advanced features are not supported yet
  - Recently, DB2 added support for Oracle's SQL (in DB2 9.7)



# 9.3 SQL ≠ SQL

*Detour*

- Helpful tools:
  - **SQL::Translator (Perl module)**  
<http://search.cpan.org/dist/SQL-Translator>
  - **SwisSQL Console (GUI tool, discontinued)**  
<http://www.swissql.com/products/sql-translator/sql-converter.html>
  - **SwisSQL API (.NET/Java SQL wrapper, discontinued)**  
<http://www.swissql.com/products/sqlone-apijava/sqlone-apijava.html>  
<http://www.swissql.com/products/sqlone-apidotnet/sqlone-apidotnet.html>





## 9.4 Advanced Concepts

- Type casting
- Ranking functions
- CASE expressions

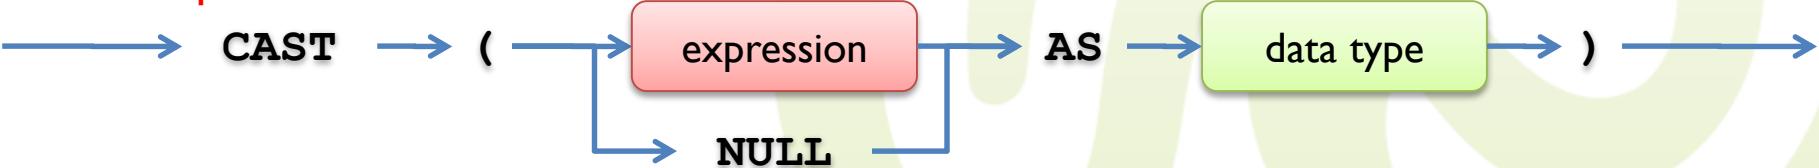




## 9.4 Type Casting

- SQL is a **strongly typed** language
  - Basically, this means that e.g. **INTEGER** is different from **VARCHAR(100)**
- If data types are incompatible, **type casting** may be used to make them compatible
  - **CAST** expression
  - During casting, precision may be lost (e.g. **FLOAT** → **INTEGER**)
  - Example:
    - **CAST (power\_strength AS NUMERIC(3, 2))**
    - **CAST (alias|| real\_name AS CHAR(255))**

**CAST expression**





## 9.4 Type Casting

- In DB2, possible castings are:

| Source                                             | Target                                                            |
|----------------------------------------------------|-------------------------------------------------------------------|
| SMALLINT, INTEGER,<br>DECIMAL, FLOAT               | SMALLINT, INTEGER, DECIMAL, FLOAT                                 |
| CHAR, VARCHAR,<br>LONG VARCHAR, CLOB               | CHAR, VARCHAR, LONG VARCHAR, CLOB,<br>BLOB                        |
| CHAR, VARCHAR                                      | SMALLINT, INTEGER, DECIMAL, DATE,<br>TIME, TIMESTAMP, VARGRAPHICS |
| SMALLINT, INTEGER,<br>DECIMAL                      | CHAR                                                              |
| DATE, TIME, TIMESTAMP                              | CHAR, VARCHAR                                                     |
| DATE                                               | DATE                                                              |
| TIME                                               | TIME                                                              |
| TIMESTAMP                                          | TIMESTAMP                                                         |
| GRAPHICS, VARGRAPHICS,<br>LONG VARGRAPHICS, DBCLOB | GRAPHICS, VARGRAPHICS,<br>LONG VARGRAPHICS, DBCLOB                |



## 9.4 Ranking Functions

- Since SQL:2003, there are special functions for working with **result lists**
- Examples:
  - Output only every other row of the list
  - Create a ranking with explicit ranks (1, 2, 3, ...)
  - On what rank position is some given row?

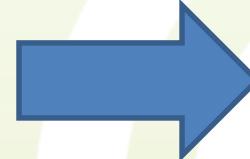


## 9.4 Ranking Functions

- **ROW\_NUMBER()** returns the **position** of each row in the result list
- **Example:**

```
SELECT name, salary,
 ROW_NUMBER() OVER (
 ORDER BY salary DESC
) AS pos
FROM salary
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Christoph | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Joachim   | 45000     |



| name      | salary    | pos |
|-----------|-----------|-----|
| Larry     | 200000000 | 1   |
| Wolf-Tilo | 75000     | 2   |
| Joachim   | 45000     | 3   |
| Christoph | 45000     | 4   |

Depending on the implementation, the last two rows may switch positions

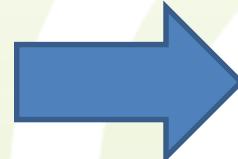


# 9.4 Ranking Functions

– Example: At which position is Wolf-Tilo?

- `SELECT name, salary,  
ROW_NUMBER() OVER (  
 ORDER BY salary DESC  
) AS pos  
FROM salary  
WHERE name = 'Wolf-Tilo'`

| person | name      | salary    |
|--------|-----------|-----------|
|        | Christoph | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Joachim   | 45000     |



| name      | salary | pos |
|-----------|--------|-----|
| Wolf-Tilo | 75000  | 2   |



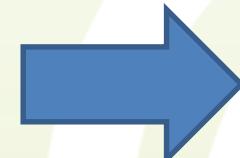
## 9.4 Ranking Functions

– Example: Show only rows at even positions:

```
• SELECT name, salary, ROW_NUMBER() OVER (
 ORDER BY salary DESC
) AS pos
FROM salary
WHERE (pos % 2) = 0
```

modulo

| person | name      | salary    |
|--------|-----------|-----------|
|        | Christoph | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Joachim   | 45000     |



| name      | salary | pos |
|-----------|--------|-----|
| Wolf-Tilo | 75000  | 2   |
| Christoph | 45000  | 4   |



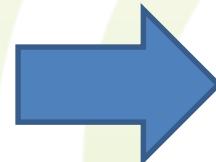
## 9.4 Ranking Functions

- **RANK()** returns the **rank** of each row in the result list

- **Example:**

```
SELECT name, salary, RANK() OVER (
 ORDER BY salary DESC
) AS rank
FROM salary
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Christoph | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Joachim   | 45000     |



| name      | salary    | rank |
|-----------|-----------|------|
| Larry     | 200000000 | 1    |
| Wolf-Tilo | 75000     | 2    |
| Joachim   | 45000     | 3    |
| Christoph | 45000     | 3    |

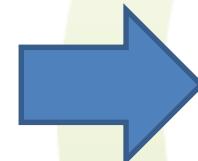


## 9.4 Ranking Functions

- **DENSE\_RANK()** works like **RANK()** but does not skip ranks on ties (as it is usually done)
- **Example:**

```
SELECT name, salary, RANK() OVER (
 ORDER BY salary ASC
) AS rank, DENSE_RANK() OVER (
 ORDER BY salary ASC
) AS drank
FROM salary
```

| person | name      | salary    |
|--------|-----------|-----------|
|        | Christoph | 45000     |
|        | Wolf-Tilo | 75000     |
|        | Larry     | 200000000 |
|        | Joachim   | 45000     |



|  | name      | salary    | rank | drank |
|--|-----------|-----------|------|-------|
|  | Christoph | 45000     | 1    | 1     |
|  | Joachim   | 45000     | 1    | 1     |
|  | Wolf-Tilo | 75000     | 3    | 2     |
|  | Larry     | 200000000 | 4    | 3     |



## 9.4 CASE Expressions

- Very often **codes** are used for storing more complex information
  - Retrieving the account information for owner “Clark” with appropriate account descriptions needs a join
  - Indicate all customers with a negative balance with the string “not creditworthy” in the query result

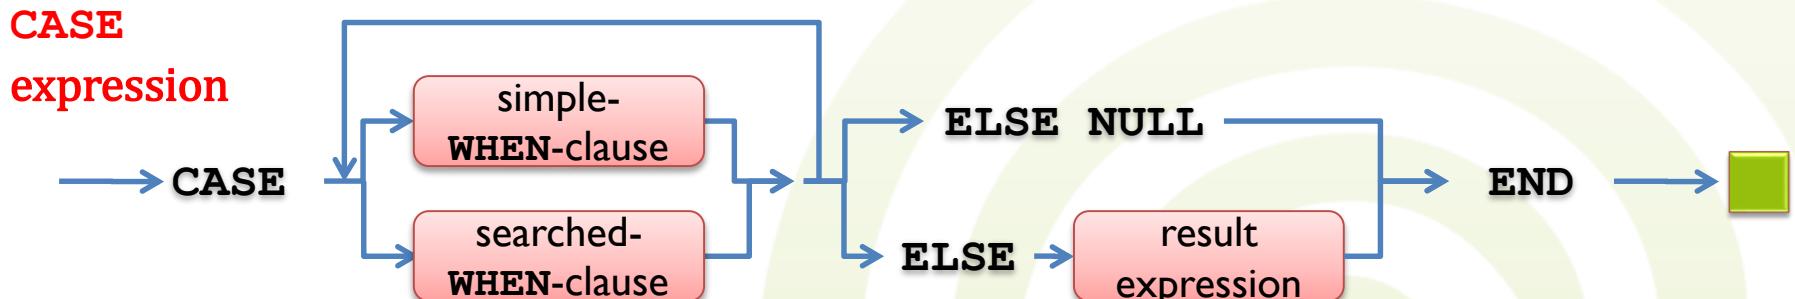
| account | owner | balance  | type |
|---------|-------|----------|------|
|         | Clark | 367,00   | 0    |
|         | Louis | -675,00  | 0    |
|         | Clark | 54987,00 | I    |

| acc_type | type | description         |
|----------|------|---------------------|
|          | 0    | checking account    |
|          | I    | savings account     |
|          | 2    | credit card account |



## 9.4 CASE Expressions

- The **CASE expression** allows a value to be selected based on the evaluation of one or more conditions (similar to “if-then-else”)
  - Comes in two flavors:



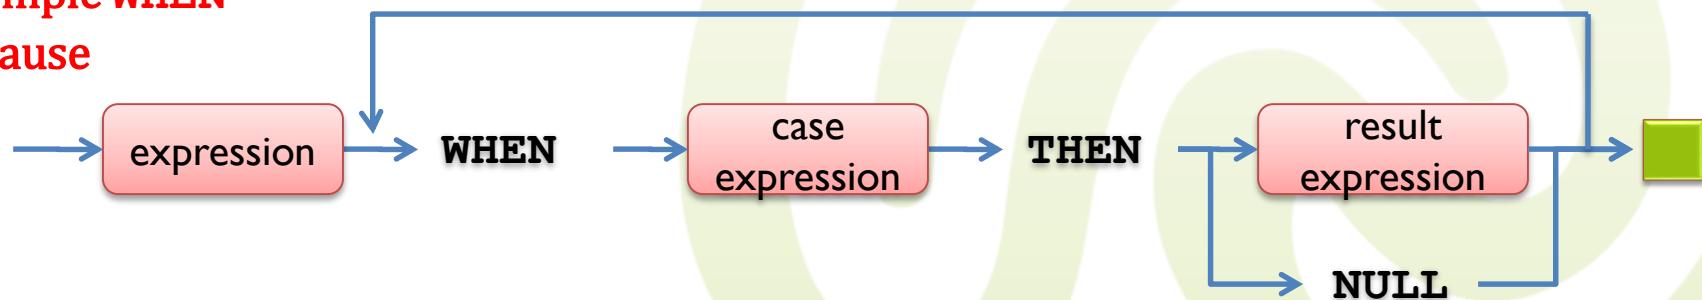


## 9.4 CASE Expressions

- The **simple WHEN clause**:

- **Compares an expression to each case expression one by one**
  - If expression is equal to search value, the corresponding result expression is returned
- **If no match is found,** then some default (**ELSE clause**) is returned
  - If **ELSE** is omitted, then **NULL** is returned

Simple WHEN  
clause





## 9.4 CASE Expressions

- **Example:** simple WHEN clause
  - Directly decode the account type
    - **SELECT** owner,  
**CASE** type  
  **WHEN** 0 **THEN** 'checking account'  
  **WHEN** 1 **THEN** 'savings account'  
  **WHEN** 2 **THEN** 'credit card account'  
**END**  
**FROM** account

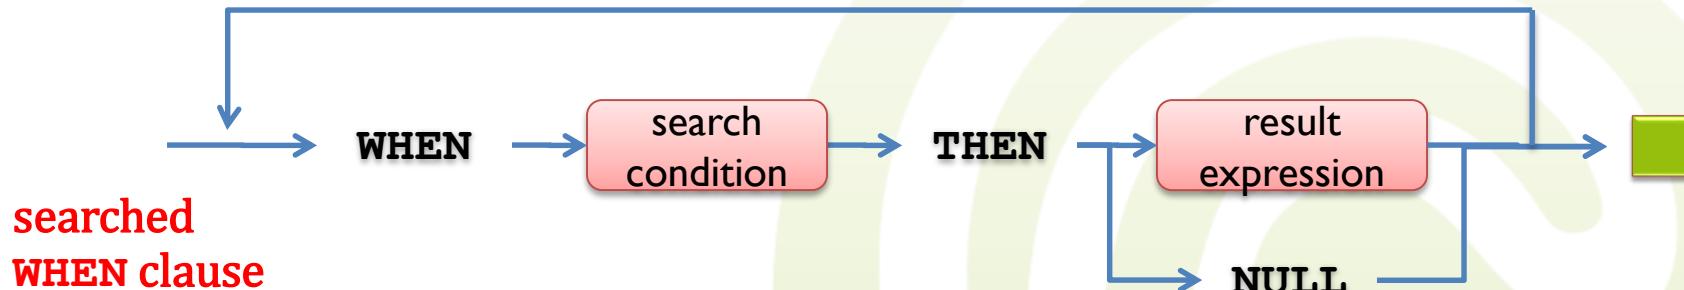
| account | owner | balance  | type |
|---------|-------|----------|------|
|         | Clark | 367,00   | 0    |
|         | Louis | -675,00  | 0    |
|         | Clark | 54987,00 | 1    |



## 9.4 CASE Expressions

- The **searched WHEN clause**:

- Checks search conditions from left to right
- **Stops as soon as a search condition evaluates to true**
  - Returns the corresponding result then
- **If no condition is true, the value given by the ELSE clause is returned (or NULL, if there is no ELSE clause)**





## 9.4 CASE Expressions

- **Example:** searched WHEN clause
  - Retrieve credit rating of customers based on their checking accounts
  - `SELECT owner,  
CASE  
 WHEN balance < 0 THEN 'not credit-worthy'  
 WHEN balance = 0 THEN 'questionable'  
 ELSE 'credit-worthy'  
END  
FROM account  
WHERE type = 0`

| account | owner    | balance | type |
|---------|----------|---------|------|
| Clark   | 367,00   | 0       |      |
| Louis   | -675,00  | 0       |      |
| Clark   | 54987,00 | 1       |      |



## 9.5 More on SQL

- There are **many more SQL** statements than we are covering in our lectures
- Moreover, there are many different SQL dialects
- If you don't want to get mad, do the following:
  - Don't care too much about the **SQL standard** (unless you are actually implementing an RDBMS)
  - Read the **SQL manuals of your RDBMS!**





# 9.5 More on SQL

- Example: DB2's SQL reference (version 9.7)
  - <http://publib.boulder.ibm.com/infocenter/db2luw/v9r7>

The screenshot shows the IBM DB2 9.7 for Linux, UNIX and Windows Infocenter. The left sidebar contains a navigation tree with categories like Messages, Multicultural support, Security, Searching text, SQL (with subtopics like How to read the syntax diagrams, Conventions used for the SQL topics, Language elements, Functions, Procedures, and Queries), and pureXML. The main content area displays a detailed syntax diagram for the `select-statement`. The diagram is represented by a series of nested arrows and text labels indicating the structure of the statement. Below the diagram, a descriptive text explains what a `select-statement` is and how it can be used. Further down, there is information about authorization and a section on `common-table-expression` with its own syntax diagram.

IBM DB2 9.7 for Linux, UNIX and Windows

Country/region [select]

Search scope: All topics

Contents

- Messages
- Multicultural support
- Security
- Searching text
- SQL
  - How to read the syntax diagrams
  - Conventions used for the SQL topics
  - Language elements
  - Functions
  - Procedures
  - Queries
    - Queries and table expressions
    - subselect
    - fullselect
    - select-statement**
  - Statements
  - Catalog views
  - SQL and XML limits
  - Reserved schema names and reserved words
  - Statements allowed in routines
  - Communications areas, descriptor areas, and
  - Explain tables
  - Explain register values- pureXML

Database fundamentals > SQL > Queries

DB2 Version 9.7 for Linux, UNIX, and Windows

### select-statement

```
>>-+-----+---+-----+--fullselect--●----->
| | .,- ----- . | |
| | V | | |
'-WITH---common-table-expression+-'

>>-+-----+●+-----+-----+-----+-->
+read-only-clause+ 'optimize-for-clause'
'-update-clause---'

>>-+-----+-----+-----+-----+--<
-isolation-clause-'
```

The `select-statement` is the form of a query that can be directly specified in a `DECLARE CURSOR` statement, or prepared and then referenced in a `DECLARE CURSOR` statement. It can also be issued through the use of dynamic SQL statements using the command line processor (or similar tools), causing a result table to be displayed on the user's screen. In either case, the table specified by a `select-statement` is the result of the fullselect.

The authorization for a `select-statement` is described in the Authorization section in "SQL queries".

### common-table-expression

```
>>-table-name---+-----+-----+-----+-----+-->
| | .,- ----- . | |
| | V | | | (1) | |
'- (----column-name---)-----'
```

<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.sql.ref.doc/doc/r0000879.html>

FoxyProxy: Patterns