



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# Relational Database Systems I

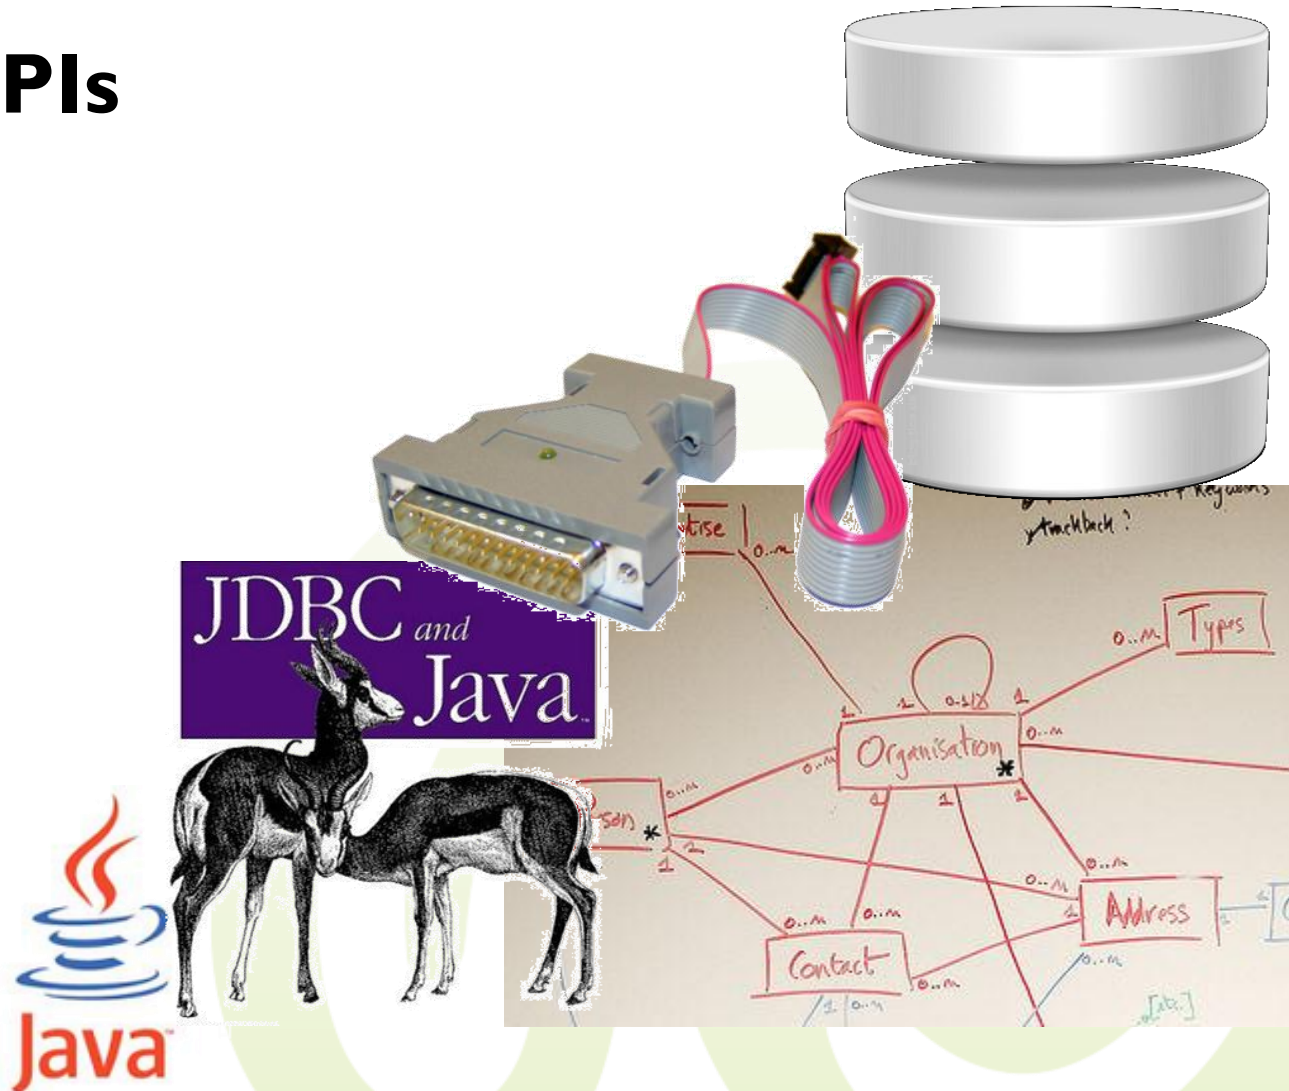
**Wolf-Tilo Balke**  
**Simon Barthel**

Institut für Informationssysteme  
Technische Universität Braunschweig  
[www.ifis.cs.tu-bs.de](http://www.ifis.cs.tu-bs.de)



# Overview

- **Database APIs**
  - CLI
  - ODBC
  - JDBC





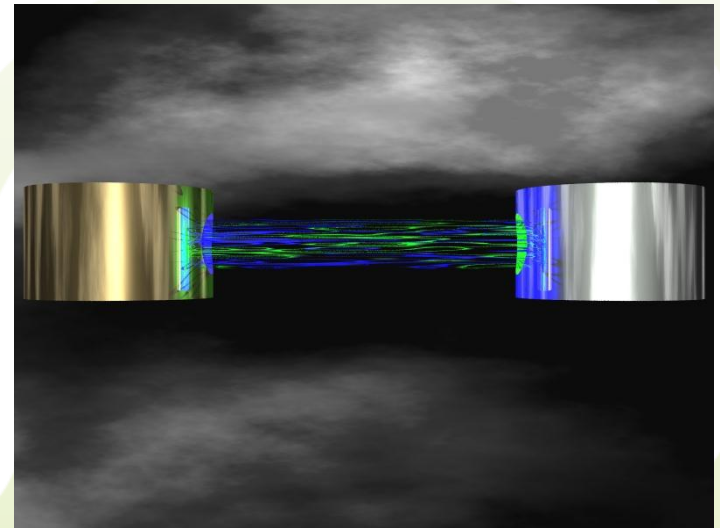
# 12 Accessing Databases

- Database access using a **library**  
(application programming interface, **API**)
  - Most popular approach
  - Prominent examples
    - **CLI** (Call level interface)
    - **ODBC** (Open Database Connectivity)
    - **JDBC** (Java Database Connectivity)



# 12 Accessing Databases

- General steps in using database APIs
  - Set up the **environment**
  - **Define** and **establish** connections to the DBMS
  - **Create** and **execute** statements (strings)
  - **Process** the results (cursor concept)
  - **Close** the connections





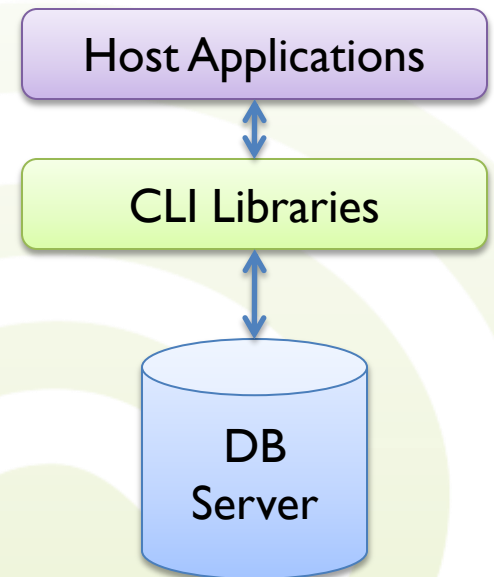
## 12.1 CLI

- The **Call Level Interface (CLI)** is an ISO software standard developed in the early 1990s
  - Defines how programs send queries to DBMS and how result sets are returned
  - Was originally targeted for **C** and **Cobol**
- Vision: **Common Application Environment**
  - Set of standards and tools to develop open applications
  - Allows to integrate different programming teams and DB vendors



## 12.1 CLI

- CLI libraries are provided by the **DB vendors**
  - Each library is specific for the respective DBMS and follows the individual DBMS's syntax
  - However, vendor libraries all follow the CLI standard and can be used interchangeably by all applications





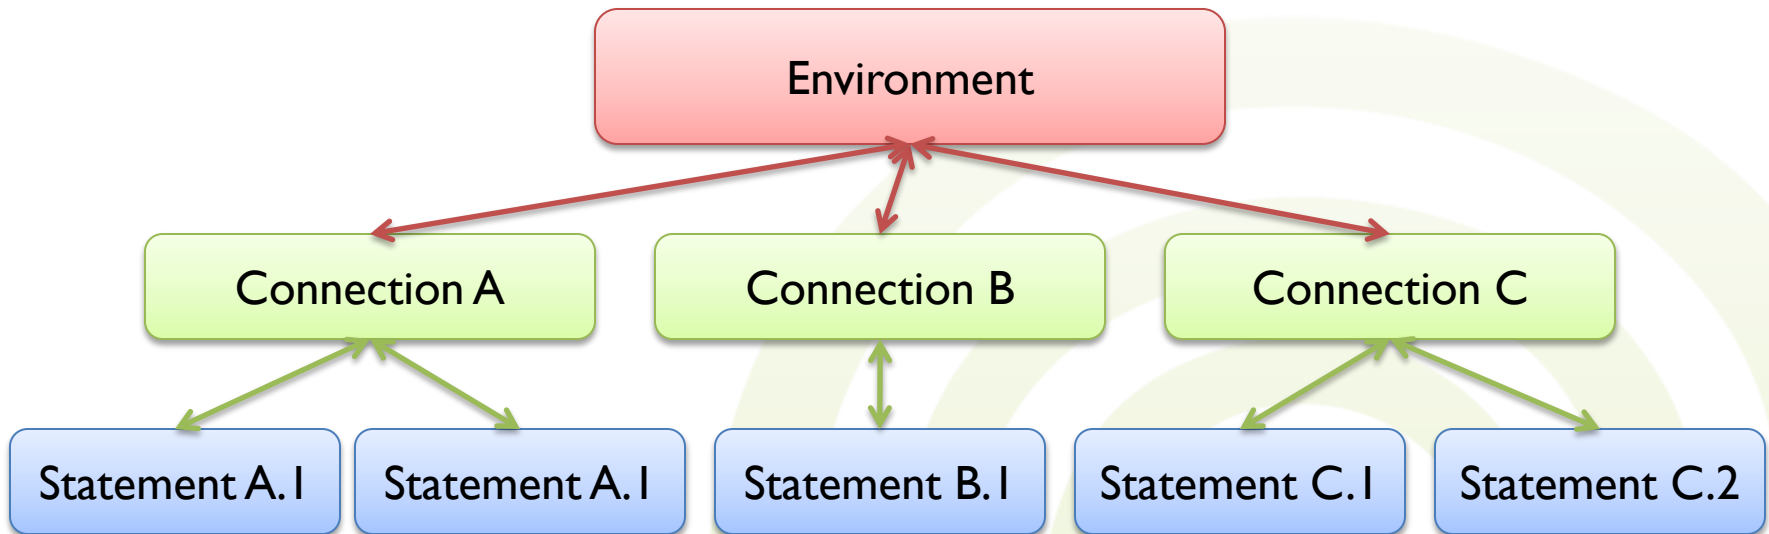
# 12.1 CLI

- Host language connects and accesses DB using following concepts
  - **Environments:** Represent the DBMS installation
    - Properties and specifics
  - **Connections:** A current session with the DBMS
    - URL, username, password, session context
  - **Statements:** SQL statements to be passed to DBMS via a connection
  - **Descriptions:** Records about tuples from a query or parameters of a statement
    - Number of attributes and respective types
    - Parameters in function calls



## 12.1 CLI

- An **environment** can host several **connections**, while a **connection** can host several **statements**







## 12.1 CLI

- When working with **CLI**, **following steps** have to be performed
  - Include the CLI function libraries and open database connections
  - Metadata about the database, tables, and columns can be retrieved
  - Define variables to contain SQL statement information
  - Execute the query and manipulate the result set in a (implicitly declared) cursor
  - Terminate statements, connections and the environment



# 12.1 CLI: Handle Concept

- Function `SQLAllocHandle(T, I, O)` is used to create data structures (variables), which are called environment, connection, and statement handles
  - **T: Handle type**, e.g., an environment, a connection, or a statement
  - **I: Input handle**, container structure at next higher level (statement < connection < environment)
  - **O: Output handle** (pointer to new handle)
- Example for handling statements:
  - `SQLAllocHandle(SQL_HANDLE_STMT, myCon, myStat);`
    - `myCon` is a previously created connection handle.
    - `myStat` is the name of the statement handle that will be created.



# 12.1 CLI: Handle Concept

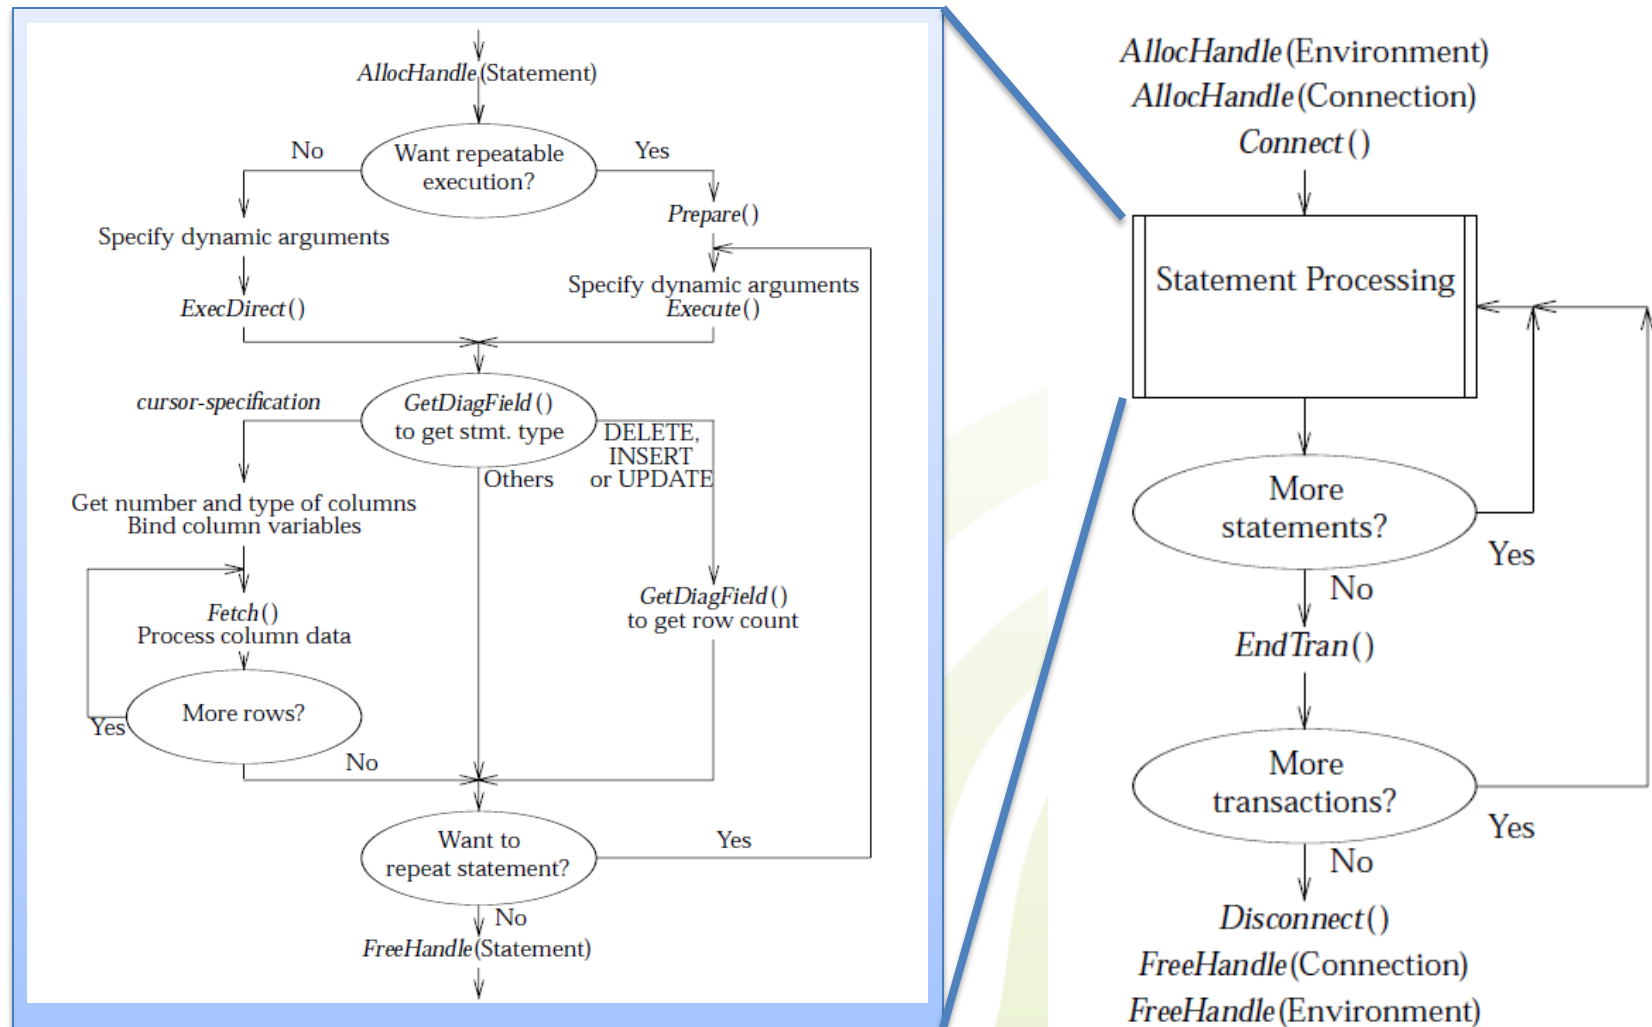
- For details please read the manual... 😊
  - **Example in C:**

```
#include <sqlcli.h>
SQLRETURN ReturnCodeEnv;
SQLHENV EnvironmentHandle;

...
ReturnCodeEnv =
SQLAllocHandle(SQL_HANDLE_ENV,
SQL_NULL_HANDLE, &EnvironmentHandle);
```



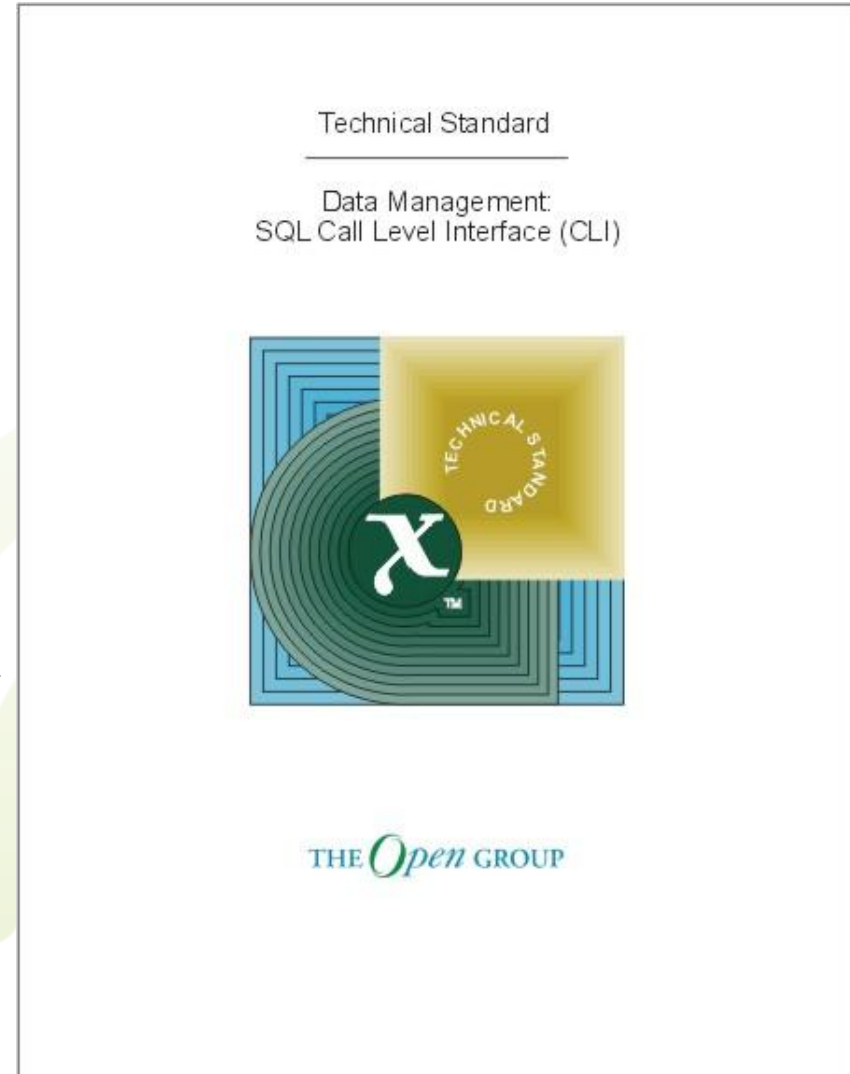
# 12.1 CLI Basic Control Flow





# 12.1 CLI

- The complete technical standard is available freely from the Open Group
- Specification C451
- Over 300 pages...
- <http://www.opengroup.org/products/publications/catalog/c451.htm>





## 12.2 ODBC

- The **Open Database Connectivity (ODBC)** provides a standardized application programming interface to DBMS using the CLI standard
  - Development driven by **Microsoft** in 1992, later versions aligned with X/Open and ISO/IEC
  - Builds on **several CLI specifications**, but does not implement the full SQL features
  - Central for the design was **independence** of programming language, operating system, and DBMS
  - Implements the standardized **middleware** concept



## 12.2 ODBC

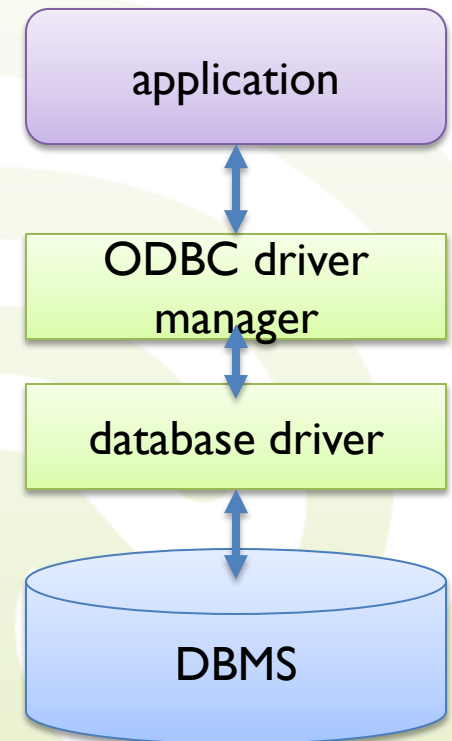
- Basic idea: **The DBMS is virtualized**
  - The person with **specialized knowledge** to make the application logic work with the database is the driver developer and not the application programmer
  - Application developers write to a generic DBMS interface and **loadable drivers** map that logic to vendor-specific commands





## 12.2 ODBC

- Being a **middleware solution** a basic implementation of ODBC always contains...
  - A generic **ODBC “driver manager” library** to interpret the applications’ commands
    - Defines standard types and features
  - And a set of **database drivers** to provide the DBMS-specific details
    - Each database vendors can write an individual driver to map ODBC commands







## 12.2 ODBC

- ODBC supports different **numbers of tiers** that have to be passed to access the databases
  - **Tier 1**
    - **Direct access** to database files by the database drivers (usually only used for desktop databases)
  - **Tier 2**
    - The database drivers prepare the requests and **pass them on to the DBMS** for execution (which is the normal case)
  - **Tier 3**
    - The database drivers prepare the requests and pass them to a specific **ODBC gateway** that manages the communication to the DBMS (e.g., via a low level interface) for execution



## 12.2 ODBC

- ODBC development has driven by the need of **easy application programming**
  - Originally in Microsoft's Visual Basic
  - But quickly adapted for use in C, C++ and other languages
- The ODBC architecture also has certain **drawbacks**
  - **Large client networks** may need a variety of drivers increasing the system-administration overhead
  - **Multi-tier ODBC drivers** can ease this problem



## 12.2 ODBC

- ODBC uses **standard CLI calls...**
- The concept of **handles** is used to set up the environment and connections
  - First, applications have to allocate an **environment handle** by calling `SQLAllocEnv`
  - Then, a handle for a **database connection** (`SQLAllocConnect`) has to be allocated before calling **connection functions** like `SQLConnect`



## 12.2 ODBC

- To process SQL statements...
  - an application must first acquire a **statement handle** by calling SQLAllocStmt
  - There is a function for **direct execution** of a SQL statement (SQLExecDirect) and functions to **prepare** and **execute statements** (SQLPrepare and SQLExecute)
  - An application can use **named cursors** by getting and setting the cursor name for a statement (SQLGetCursorName and SQLSetCursorName)
  - An application **retrieves a row** in a result set by calling SQLFetch
  - ...



## 12.2 ODBC

- As part of ODBC's termination logic...
  - Every application should free statement handles using `SQLFreeStmtclose`
  - Close the connection and free the connection and environment handles by calling `SQLDisconnect`, `SQLFreeConnect`, and `SQLFreeEnv`
  - We won't go into more implementation details here, but consider the exact use for the case of JDBC



## 12.3 JDBC



- **JDBC** provides a standard **Java** library for **accessing tabular data**
  - Tabular data usually means a **relational DBMS**
  - API provides standard way to **connect** to a DB
  - API allows to perform **dynamic queries**
  - Method to create stored (**parameterized**) **queries**
  - Provides data types for Java/DB **impedance mismatch**
    - **Result sets** with rows and columns
    - Methods for accessing table **meta data**
  - Provides functionality **independent** of chosen DBMS



## 12.3 JDBC



- JDBC does **not standardize SQL**
  - SQL statements are treated as **Java strings**
  - In case of full dynamic SQL, sometimes excessive **string manipulation** is necessary
  - If DBMS uses different/extended SQL syntax, this has to be considered by the programmer
- JDBC is not an acronym, but a registered product **trademark** by Sun Microsystems
  - However, usually, it is assumed that it stands for **Java Database Connectivity**



## 12.3 JDBC



- Why not just use ODBC?
  - ODBC is based on **binary libraries** (usually written in C)
    - Native calls necessary
    - Not platform-independent which is one of Java's goals
  - 1:1 translation from ODBC to Java does not work as ODBC heavily **relies on pointers**
  - ODBC API is more **complex** and **littered** (and thus harder to learn and use)
    - For example, programmer needs to worry about byte alignment and advanced connection properties explicitly
  - Intention was to create a “**pure**” **Java** and fully portable API
    - No installation required, JDBC can easily be bundled into the application archive





## 12.3 JDBC

- JDBC is composed of **two primary components**
- **JDBC API:** An **programming interface** for database connectivity.
  - Written in **100% pure Java**
  - Completely **independent** of platform, vendor, and DBMS
  - Provided by the Sun in its **Java SDK by default**
    - Usually to be found in `java.sql` and `javax.sql`



## 12.3 JDBC

- **JDBC driver**

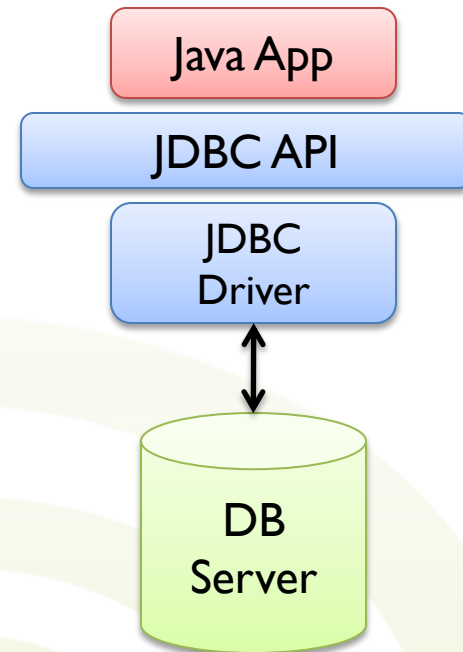
- Implementation of the respective API interface, responsible for **communication** with the **database**
- Interface implementation in Java, but may depend on any amount of binary libraries, middleware, or other tools
- Heavily dependent on the used DBMS
- Usually provided by the **DB vendor**





## 12.3 JDBC

- General Architecture
  - Java application uses API
  - API uses driver
  - Driver communicates with DB
- If you change the DBMS, you need to
  - Provide a new driver
  - Change configuration of driver
  - Assuming the SQL syntax is compatible, you are done
    - If not, you are in trouble...





## 12.3 JDBC: Versions

- There are several versions of JDBC, each with improved functionality

Version	Year	Java Version
JDBC 4.0	2006	Java 6
<b>JDBC 3.0</b>	<b>2001</b>	<b>Java 1.4 &amp; Java 5</b>
JDBC 2.1	1999	Java 1.3
JDBC 1.2	1997	Java 1.1

- JDBC drivers are written for a specific JDBC version
  - Driver should match the JDBC version
  - However, most features also work with outdated drivers
- JDBC 3.0 is most common
- However, JDBC 4 is gaining momentum as it contains many useful features and improvements



## 12.3 JDBC: Levels

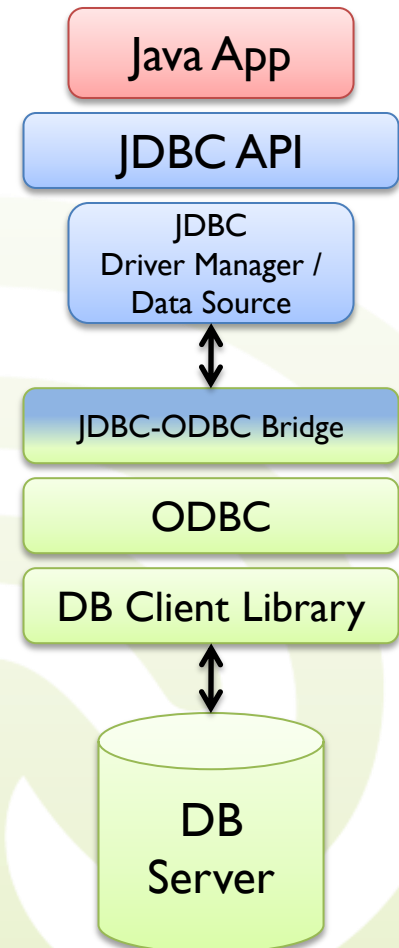
- Beside versions, there are JDBC levels
  - Comparable to ODBC tiers
  - For each level, there are different drivers
    - Be careful when picking your driver! You need the right version and correct level!
  - All levels offer the same functionality (i.e., API is the same), but use different means of driver implementation and communication with the DBMS
    - Different performance and portability properties



## 12.3 JDBC: Levels

- **Level I: JDBC/ODBC bridge**
  - JDBC driver just translates requests to ODBC calls
    - **Performance overhead** due to translation
  - Needs correctly installed ODBC drivers on every client machine
    - Distribution difficult
    - ODBC drivers are not platform-independent

### JDBC Level I (ODBC Bridge)



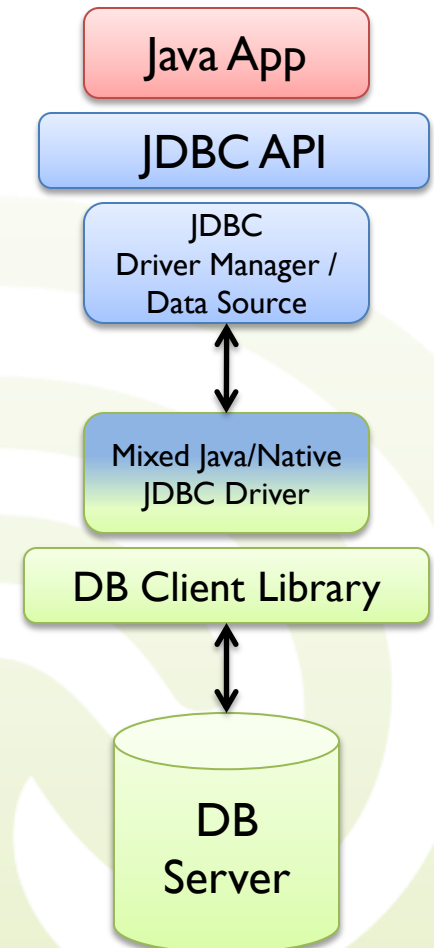


## 12.3 JDBC: Levels

- **Level 2: Native API**

- JDBC driver uses native calls to connect to a proprietary client software which handles DB connections
  - e.g. ORACLE client (which is a 1.7 GB installation)
- Difficult to port and with deployment problems
- Often used as cheap-and-dirty solution for older systems
  - Also, may be a good idea if application is running on the same machine as the DBMS

### JDBC Level 2 (Native API)



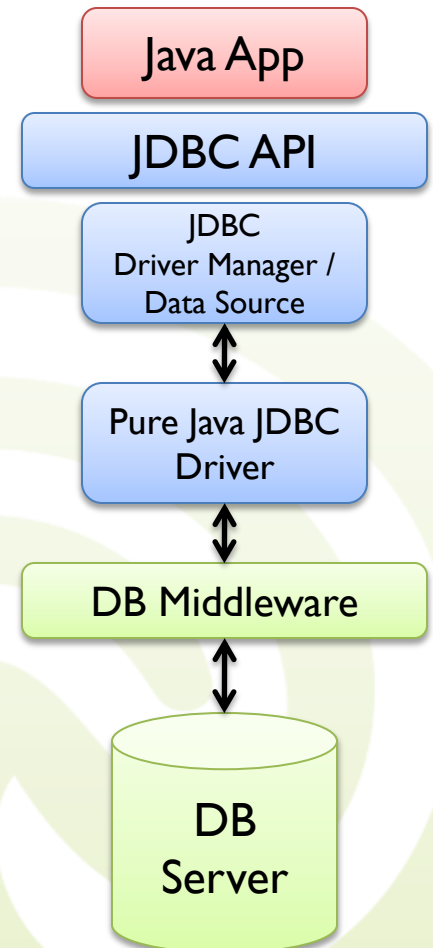


## 12.3 JDBC: Levels

- **Level 3: Middleware**

- JDBC driver communicates with a **middleware** software instead of the DBMS
- Often used for **large-scale enterprise applications** in a multi-tier-environment
- Vendor specific translation may happen at the middleware
  - Just one client driver for any used DBMS
- Middleware encapsulates the actual DBMS
  - Useful for advanced clustering, extended security, logging, caching, pooling, etc..

### JDBC Level 3 (Middleware)



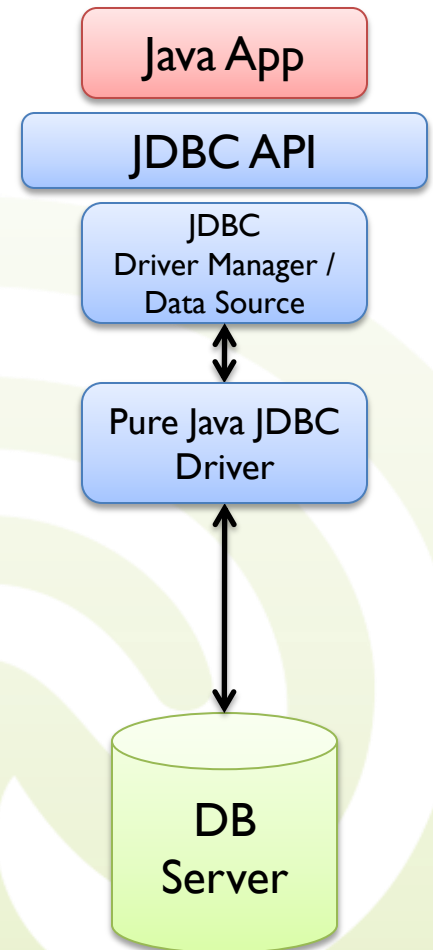




## 12.3 JDBC - Levels

- **Level 4: Direct pure Java**
  - Driver purely written in Java
    - No call translation
    - No installation, no deployment problems
    - Full portability due to platform-independence
  - Driver connects directly to the DBMS
    - You need a different driver for each different DBMS
    - Superior performance in remote scenarios
    - For access to a local DBMS, Level 1 might be better

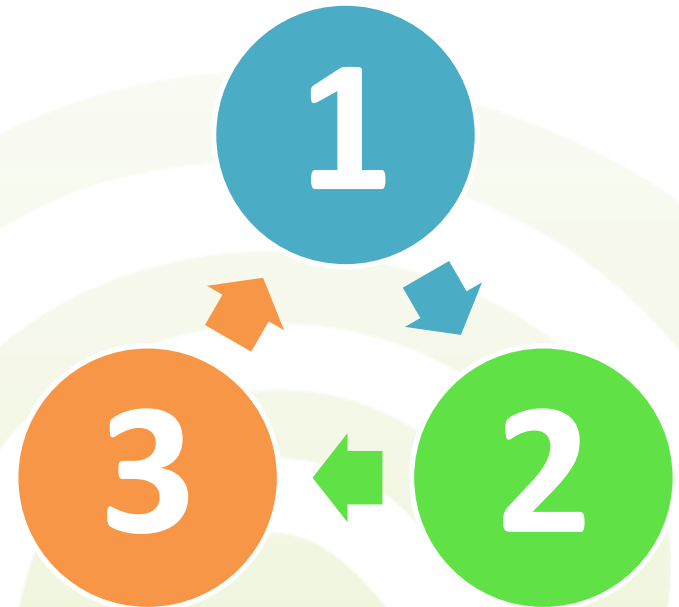
### JDBC Level 4 (Pure Java)





## 12.3 JDBC

- Basic steps when working with JDBC
  1. **Load** the driver
  2. **Define** a connection URL
  3. **Establish** a connection
  4. **Create** a statement(s)
  5. **Execute** a statement(s)
  6. **Process** the result(s)
  7. **Close** the connection





- A connection creates a query session for a given **user** within a specific **DBMS**
  - Each connection may have specific **properties**
  - The DBMS server is specified using an **URL**
  - Common URL format is
    - jdbc:[driverAlias]:[driverParameters]
    - DB2 Level 4:
      - jdbc:db2://[server][:port]/[db-name]
      - Example: jdbc:db2://dblab.ifis.cs.tu-bs.de:50000/DBLAB
    - MySQL Level 4:
      - jdbc:mysql://[server][:port]/[database]
      - Example: jdbc:mysql://dblab.ifis.cs.tu-bs.de:3306/rdb1
    - Further URL formats for most DBMS:
      - <http://www.redmountainsw.com/wordpress/archives/jdbc-connection-urls>



- To create a connection, the **DriverManager** is used
  - Returns a **Connection** object

```
public Connection getConnection(Properties connectionProps)
    throws SQLException {
    return DriverManager.getConnection (
        "jdbc:db2://" +
        connectionProps.getProperty("server") + ":" +
        connectionProps.getProperty("port") + "/" +
        connectionProps.getProperty("database"),
        connectionProps.getProperty("user"),
        connectionProps.getProperty("password")
    );
}
```



- The Connection object coordinates the whole database-application interaction
  - Contains DBMS instance metadata
  - Responsible for providing query statements
- Driver matching the driver alias in the URL must be contained in the classpath

```
Exception in thread "main" java.sql.SQLException: No  
suitable driver found for jdbc:db2://dblab.ifis.cs.tu-  
bs.de:50000/dblab
```



- Using the connection object, the DB **metadata** can be accessed
  - DBMS **name**, DBMS **version**, DBMS installation **properties**, available **schemas**, available **tables**, **columns** for a table, **primary keys** for a given table, and many **many more**

```
DatabaseMetaData metaData =  
    connection.getMetaData() ;  
String dbmsName =  
    metaData.getDatabaseProductName() ;  
ResultSet schemas =  
    metaData.getSchemas() ;
```



## 12.3 JDBC: Statements

*Detour*

- To actually execute a SQL statement, JDBC provides Statement objects
  - Any kind of statement is created by the **Connection** object
  - When a query is executed, a **ResultSet** is returned
    - ResultSet encapsulates SQL result tables
    - Result is stored on the server and transferred to the client **row by row**





## 12.3 JDBC: Statements

*Detour*

- There are three Statements facilities provided by JDBC
  - **Statement**
    - Used for simple or very rare statements
    - A real **dynamic fashion** of executing SQL
    - The whole SQL statements is provided as a **String**
  - **PreparedStatement**
    - Used for frequent statements
    - **Semi-dynamic** statements
    - Statement is provided as **parameterized String**
    - For each execution, **parameters are replaced by values**
    - Usually, **performance is much higher** due to lower overhead and better query plan optimization
  - **CallableStatement**
    - Used to execute server-side stored procedure (UDF)





## 12.3 JDBC: Statements

*Detour*

- Using **simple statements**
  - Create a **statement object** with the connection
  - Call one of the three execution methods
  - **executeQuery(String query)**:
    - Use for **SELECT** queries
    - Returns a **ResultSet**
  - **executeUpdate(String query)**:
    - Use for any DDL or data changing statements (**INSERT, UPDATE, DELETE**)
    - Returns an integer with number of affected rows
  - **execute(String query)**:
    - Advanced feature for multi-**ResultSet** queries



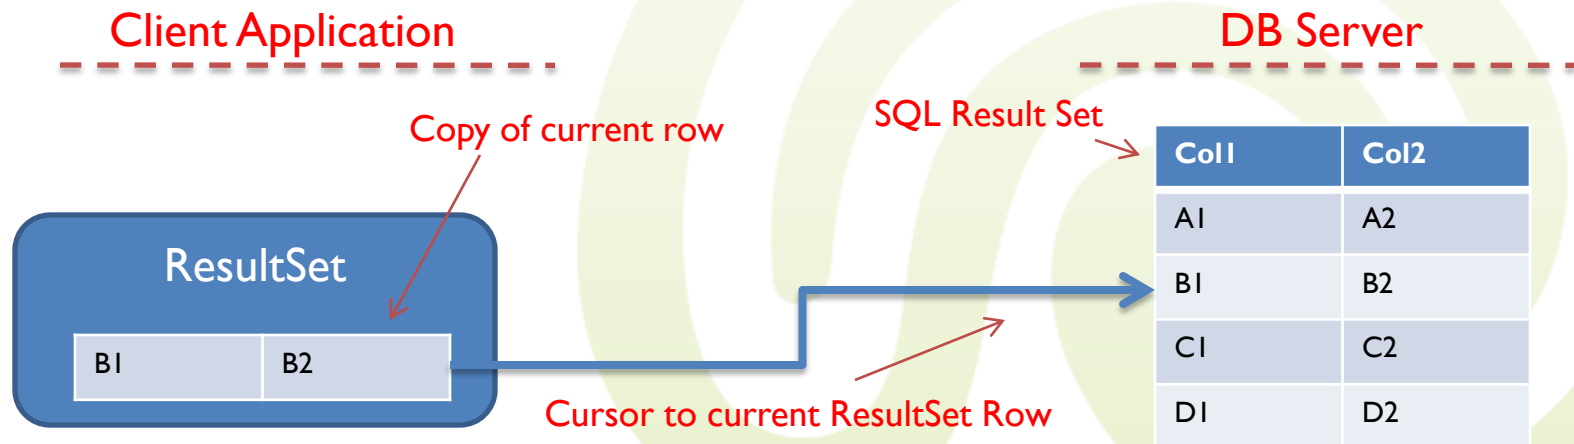
```
Statement stmt =  
    connection.createStatement() ;  
ResultSet rs = stmt.executeQuery (  
    "SELECT count(*) FROM IMDB.title") ;
```



## 12.3 JDBC: Result Sets

*Detour*

- **ResultSet** encapsulate a query result table
  - Rows are retrieved one after another from the server
    - Inspired by (but not compatible to) the Iterator-Interface
    - A **cursor** is held pointing the actual row in the server-side result set
      - like row cursor EmbeddedSQL
  - At first, the result set points **before the first row** (i.e. to no row at all)
  - For each row, the column values may be retrieved individually by **column getter methods**





## 12.3 JDBC: Statements

*Detour*

- After all result set rows have been read, the statement is marked complete
  - Statements and Results are usually **garbage collected** by Java
  - However, it is highly recommended to **manually close** (using the `close()`-method) statements (and thus result sets) to save system resources in the mean time



## 12.3 JDBC: Result Sets

*Detour*

- Default ResultSet only allows moving the cursor forward and is read only
  - `TYPE_FORWARD_ONLY` and `CONCUR_READ_ONLY`
- Movement operations for `TYPE_FORWARD_ONLY`
  - `next()`: Moves the cursor to the next row
    - Returns true if successful, false if there is no next row
  - `afterLast()`: Move the cursor after the last row

```
// iterate over all result set rows
while (rs.next()) {
    // do something with the current row
}
```



## 12.3 JDBC: Result Sets

*Detour*

- There are multiple getter methods for the column values
  - Columns may be either accessed by **name** or by **number (starting at 1)**
  - There are getters for each **variable type**
    - Each SQL data type is mapped to a Java data type
    - All getters are named `getTYPE()` (e.g. `getInt()`, `getDouble()`, ...)
- Example: Return all IDs and realNames of heroes

```
ResultSet rs = stmt.executeQuery(  
    "SELECT id, realName FROM heroes");  
while(rs.next()) {  
    int id = rs.getInt(1); // or: rs.getInt("id")  
    String realName = rs.getString("realName");  
    System.out.println(id + ":" + realName);  
}
```



## 12.3 JDBC: Data Types

*Detour*

**Example: Direct JDBC data types for DB2**

Java Data Type	SQL Data Type
short, boolean, byte	SMALLINT
int, java.lang.Integer	INTEGER
long, java.lang.Long	DECIMAL(19,0)
long, java.lang.Long	BIGINT
float, java.lang.Float	REAL
double, java.lang.Double	DOUBLE
java.math.BigDecimal	DECIMAL(p,s)
java.lang.String	CHAR(n)
java.lang.String	GRAPHIC(m)
java.lang.String	VARCHAR(n)
java.lang.String	VARGRAPHIC(m)
java.lang.String	CLOB(n)



## 12.3 JDBC: Data Types

*Detour*

Java Data Type	SQL Data Type
byte[]	CHAR(n) FOR BIT DATA
byte[]	VARCHAR(n) FOR BIT DATA
byte[]	BLOB(n)
byte[]	ROWID
java.sql.Blob	BLOB(n)
java.sql.Clob	CLOB(n)
java.sql.Clob	DBCLOB(m)
<b>java.sql.Date</b>	<b>DATE</b>
<b>java.sql.Time</b>	<b>TIME</b>
<b>java.sql.Timestamp</b>	<b>TIMESTAMP</b>
java.io.ByteArrayInputStream	BLOB(n)
java.io.StringReader	CLOB(n)
java.io.ByteArrayInputStream	CLOB(n)
com.ibm.db2.jcc.DB2RowID	ROWID
java.net.URL	DATALINK



## 12.3 JDBC: Result Sets

*Detour*

- Different ResultSet types may be used changing **navigation** and **update options**
  - It depends on the DBMS and the JDBC drivers which options / combinations are possible
- **Navigation** options
  - **TYPE\_FORWARD\_ONLY**
    - Only forward movement is possible
    - DBMS may materialize result set incrementally, it is not exactly clear how up-to-date the row is
  - **TYPE\_SCROLL\_INSENSITIVE**
    - Usually only works if DBMS supports rownums
    - Cursor may forward, backward, or directly to a particular row
    - While the result set is open, changes in the underlying data are not visible







## 12.3 JDBC: Result Sets

*Detour*

- Some additional navigation methods for **TYPE\_SCROLL\_INSENSITIVE**
  - **Previous()** Moves cursor to the previous row
  - **beforeFirst()**: Moves cursor before the first row
  - **relative(int x)**: Moves cursor x rows forward (or backward if x is negative)
  - **absolute(int x)**: Moves cursor to the given absolute row number

### – **TYPE\_SCROLL\_SENSITIVE**

- Same as **TYPE\_SCROLL\_INSENSITIVE**
- BUT changes to the underlying data are reflected in the result (i.e. rows are always up-to-date)
- May have bad performance





## 12.3 JDBC: Result Sets

*Detour*

- **Update options**

- **CONCUR\_READ\_ONLY**

- Results can only be read
    - Unlimited concurrency



- **CONCUR\_UPDATABLE**

- Result set may be **updated** (i.e. row/column values changed, rows inserted, etc.; similar to updateable view)
    - May **degenerate performance** in massively concurrent applications due to lock contention thrashing (see RDB2)





- Different result set types are declared during statement initialization
  - `Connection.createStatement`  
(`int resultSetTyp`, `int resultSetConcurrency`)
    - int values declared as constants in `ResultSet`

```
Statement stmt =  
    connection.createStatement (  
        ResultSet.TYPE_SCROLL_INSENSITIVE,  
        ResultSet.CONCUR_READ_ONLY) ;
```



## 12.3 JDBC: Result Sets

*Detour*

- If a result set is updatable, you may use a set of **update *TYPE*(column, value)** methods
  - Similar to get methods
  - Column may either be references by **name or number**
  - All update commands relate to the current row
  - Updates are only performed after you call **updateRow()**
    - There may be multiple update commands followed by one **updateRow()**
    - If you want to cancel the updates, call **cancelRowUpdates()**
    - If you move to another row without **updateRow()**, nothing happens

```
Statement stmt = connection. createStatement(ResultSet.TYPE_FORWARD_ONLY,  
    ResultSet.CONCUR_UPDATABLE);  
ResultSet rs= stmt.executeQuery  
    ("SELECT id, realName FROM heroes");  
while (rs.next()) {  
    rs.updateString("realName", rs.getString("realName").toUpperCase());  
    rs.updateRow();  
}
```



- Furthermore, new rows may be **inserted**
  - There is a special row for inserting (“**insert row**”)
  - To start an insertion, move cursor to the insert row by using `moveToInsertRow()`
  - While being in the insert row, you may only use get and update methods
  - You must update all columns
  - When all columns are updated, call `insertRow()` to commit the insert
    - Cursor will move to old position
    - There is no way to know where the row is inserted nor if it will appear in the current result set
  - Of course, this only works with updatable result sets



## 12.3 JDBC: Result Sets

*Detour*

- To delete the current row, call **deleteRow()**
  - Cannot be performed when being in the insert row

```
Statement stmt = connection.createStatement
    (ResultSet.TYPE_FORWARD_ONLY, ResultSet.CONCUR_UPDATABLE);
ResultSet rs= stmt.executeQuery
    ("SELECT id, realName FROM heroes");
//
rs.moveToInsertRow();
rs.updateInt(1, 999);
rs.updateString(2, "Peter Parker")
rs.insertRow();
//
while (rs.next()) {
    rs.deleteRow();
}
```





## 12.3 JDBC: Prepared Statements

- When performing a simple statement, roughly the following happens
  - The statement is composed in your app using **String manipulation**
  - The SQL String is **wrapped** and send to the **database** via the JDBC driver
  - The DBMS **parses** and **checks** the statement
  - The DBMS **compiles** the statement
  - The DBMS **optimizes** the statement and tries to find the best **access path**
  - The statement is **executed**
- When you execute the same/similar statement multiple times, all those steps are performed for each single statement



## 12.3 JDBC: Prepared Statements

- To avoid unnecessary overhead, **prepared statements** may be used
- Prepared statements use **parameterized SQL**
  - Use ? as markers for parameters
  - Example: “**SELECT \* FROM heroes WHERE id = ?**”
    - Generic SQL query for retrieving an hero by it's ID
  - Prepared Statements may either be used for queries or for updates / DDL operations







## 12.3 JDBC: Prepared Statements

- Prepared statements use the following workflow
  - When creating a (parameterized) prepared statement, it is wrapped, sent to the DBMS, parsed, checked, and optimized
    - **Only once for any number of execution**
  - Each time it is executed, the values for the parameters are transferred to the DBMS and the statement is executed
  - **Performance may be significantly higher** than when using simple statements





## 12.3 JDBC: Prepared Statements

- To supply values for the placeholders, use **set***TYPE*(number, value) methods
  - Like for the get and update methods, there are set methods for any data type
    - Placeholders are referenced by the position in the SQL string starting with 1

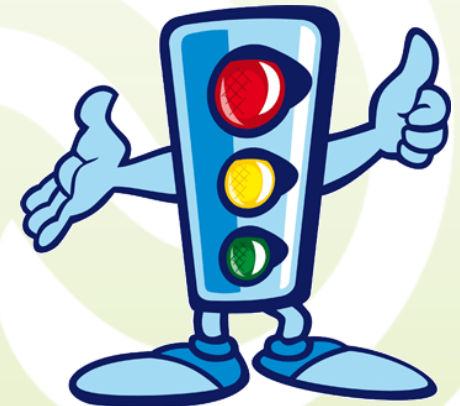
```
PreparedStatement moviesInYears = connection. prepareStatement
    ("SELECT * FROM movies WHERE releaseDate=>? AND releaseDate=<?");
for (int i=0; i<10; i++) {
    moviesInYears.setInt(1, 1990+i*2);
    moviesInYears.setInt(2, 1991+i*2);
    ResultSet rs = moviesInYears.executeQuery();
    // ... do something
}
```



## 12.3 JDBC: Transactions

*Detour*

- Of course, you may use **transactions** within JDBC
  - Depending on the DBMS, transactions are either enabled or disabled by default
    - Transactions disabled means that auto-commit mode is enabled and vice versa
  - Use **setAutoCommit**(boolean switch) to change transactional behavior
    - **true**: Use no transactions (every statement is an auto transaction)
    - **false**: Use transactions





## 12.3 JDBC: Transactions

*Detour*

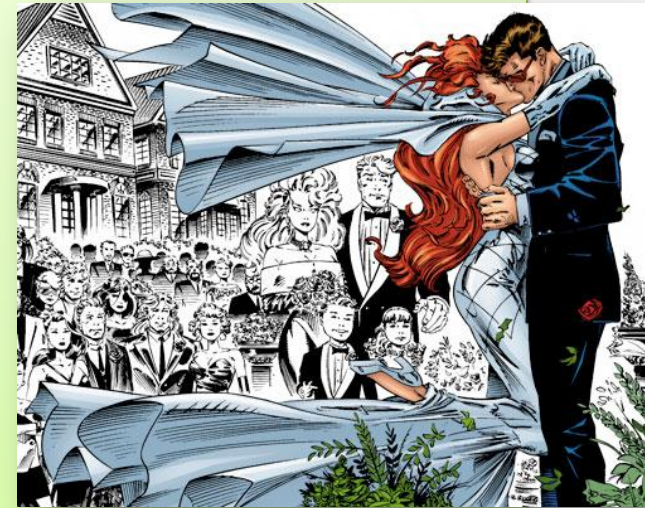
- When transactions are enabled, any number of statements is considered as one transaction until it is **committed** or **canceled**
  - Use `Connection.commit()` to commit a transaction
  - You may also create **save points** using `Connection.setSavepoint(String savepointName)`
  - Use `Connection.rollback()` to roll it back
    - `Connection.rollback(String savepointName)` returns to a given safe point



## 12.3 JDBC: Transactions

*Detour*

```
connection.setAutoCommit(false);
PreparedStatement changeName= connection. prepareStatement
    ("UPDATE hero SET name=? WHERE name=?");
//
changeName.setString(1, "Jean Grey-Summers");
changeName.setString (2, "Jean Grey");
changeName.executeUpdate();
//
changeName.setString(1, "Scott Grey-Summers");
changeName.setString (2, "Scott Summers");
changeName.executeUpdate();
//
connection.commit();
```







# Next Lecture

- General Problem of Object Persistence
- Manual persistence
  - Generating UUIDs
- Persistence frameworks
  - JPA
- Object databases
  - db4o

