



ifis

Institut für Informationssysteme
Technische Universität Braunschweig

Relational Database Systems I

**Wolf-Tilo Balke
Simon Barthel**

Institut für Informationssysteme
Technische Universität Braunschweig
www.ifis.cs.tu-bs.de



I3 Object Persistence

- General problem
- Manual persistence
 - Generating UIDs
- Persistence frameworks
 - JPA
- Object databases
 - db4o





Exam Stuff

- Written exam on March 1, 2013 at 11:30
 - Duration: 90 minutes
 - BI83, SN22.1, SN23.2
- Remember the 50% homework rule!
 - Who already took the exam doesn't need 50% (at least this time)
 - New Bachelor Prüfungs Ordnung (BPO)
 - Exercises = Studienleistung
 - Old Bachelor Prüfungs Ordnung (BPO)
 - Exercises are required for exam



Exam Stuff

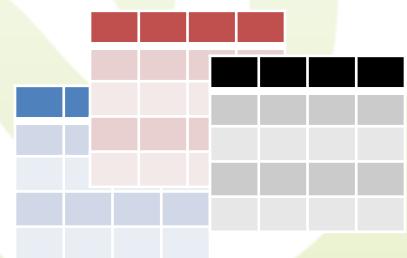
- You may bring **two sheets** of A4 paper with you
 - You decide what is written on these sheets
 - You may write on the **front and back** of these sheets
- **Important:**
The sheets must be **written on by hand**
 - This means: **No printouts, no copies!**
 - Otherwise, we will confiscate the sheets
- The exam will be about:
 - **All topics of the lecture** (excluding detours)
 - **All homework exercises**
 - **Detours covered in homework exercises** are **relevant** for the exam!





I3.1 Object Persistence

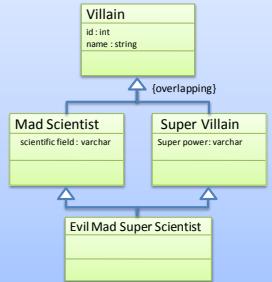
- Within typical software engineering cycles, usually **application data models** are developed
 - Data models are used **internally** in application, following its design and programming paradigms
 - Nowadays, most application data models are **object-oriented**
 - Also often called **domain model**
- When developing an application using a database there always is one huge problem
 - **How do you map your domain data model to your database data model?**
 - **Impedance mismatch!**



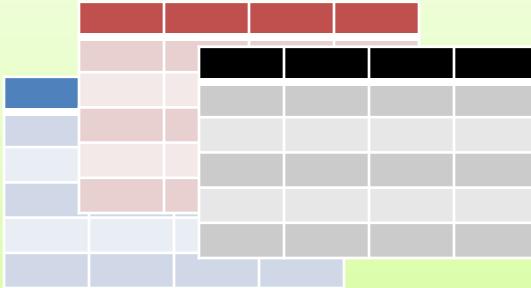


I3.1 Object Persistence

Application Data Model



Database Data Model



Application RDBMS RDBMS





I3.1 Object Persistence

- Model mapping is hard for object-oriented programming languages!
 - **Object model differs significantly from the relational model**
- In most cases, this leads to the fact that **developers adapt their domain model to the used database**
 - But good software engineering demands that your domain model follows your **business needs**, not the needs of the underlying storage technology!



I3.1 Object Persistence

- The object model:
 - Objects **represent objects in the real world**
 - Thus, objects have a **state** and a **behavior**
 - Example: This car is **blue** and has **33 PS**, it can **drive** and **honk the horn**
- Besides state and behavior, objects may have **complex relationships** to other objects
 - Usually, all relationships we already know from UML are possible
 - **Generalization, specialization, aggregation, composition, association, etc.**
 - **Example:**
 - A car may have 4 wheels and 2 doors
 - When the car drives, also the wheels are moving and rotating





I3. I Object Persistence

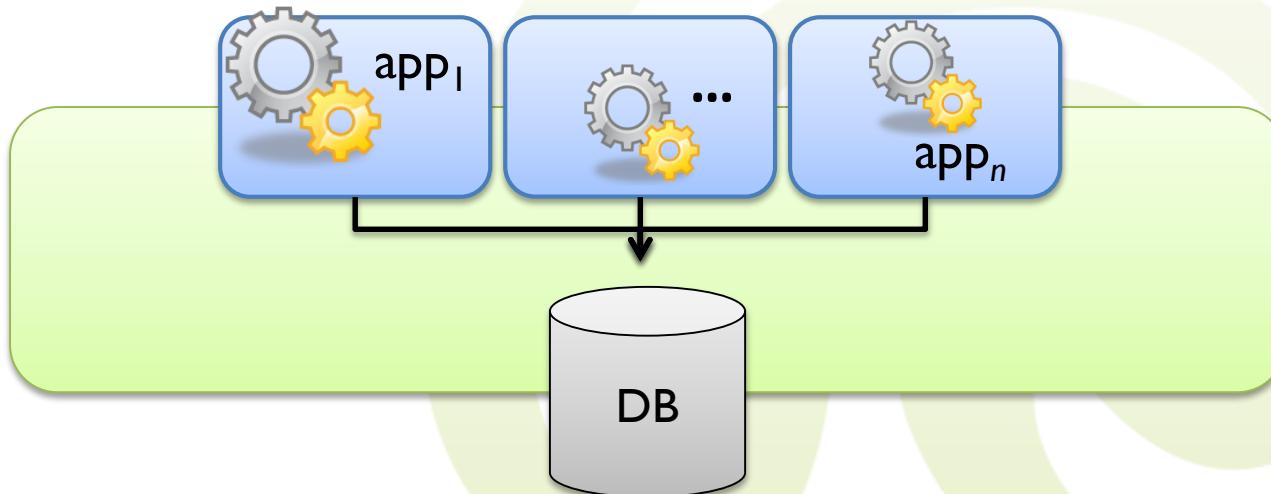
- Application objects in an application only “live” throughout the lifetime of the application
 - They are **transient**
- However, you may want to use a database to **store the state of an object**
 - Thus, the object may **persist** beyond the application’s termination
 - Object may thus be **retrieved** later or shared among **programs**
- Permanently storing the state of an object is called **object persistence**
 - Which also includes **restoring** the object





I3.1 Object Persistence

- When your object data is persistent, you may
 - Exchange it with other applications
 - Inspect it manually
 - Continue using it when you restart your application
 - Replicate and aggregate it
 - ...





13.1 Object Persistence

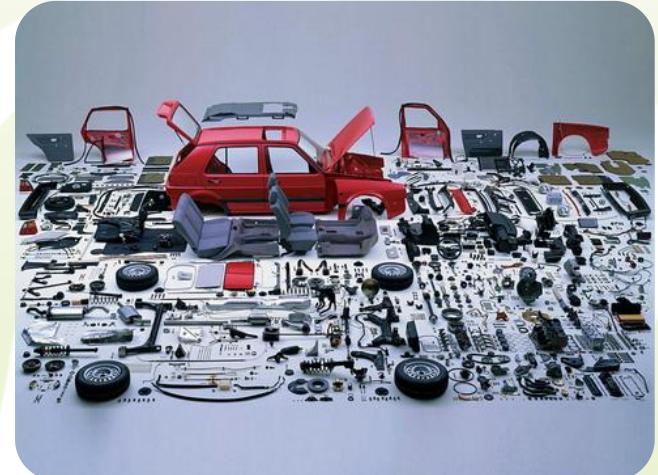
- There is a problem with **object persistence** in relational DBMS
- Each object has a **unique identity** that must be preserved
 - Objects have a unique **implicit immutable identity** independent of their state (values)
 - However, the identity cannot be accessed. It just “is.”
 - In the relational model, rows are **explicitly identified** by their **values**
 - **No duplicates**
 - The only safe assumption: All columns taken together form a **key**
 - To make objects persistent in a RDBMS, an **explicit identity (key) needs to be generated**
 - So called **unique identifier (UID)**





I3. I Object Persistence

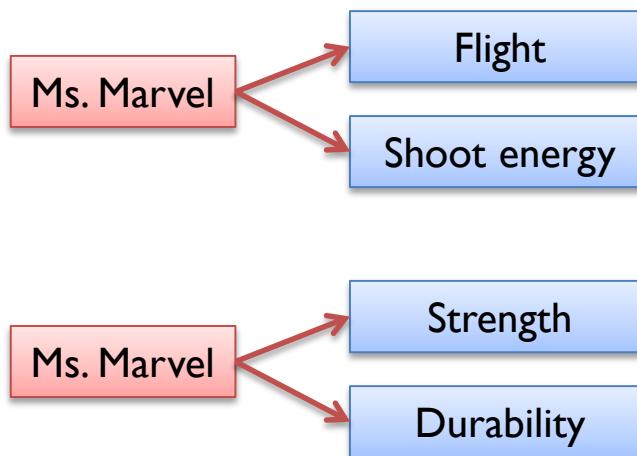
- Objects may have complex relationships (**object structure**)
 - As a whole, this structure is **hard to capture** in a **RDBMS**
 - Remember: You only have tables, columns and rows
 - Objects needs to be **disassembled** and stored in **multiple tables** linked by foreign key relationships
 - Think of this:
„Every time you want to park your car in the garage, you have to disassemble it and assemble it when you want to drive to work...“





I3.1 Object Persistence

- **Example:** Store following information in a RDB
 - A hero called **Ms. Marvel** has the powers to **fly** and **shoot energy bursts** out of her hands
 - Another hero also called **Ms. Marvel** has the powers of super human **strength** and **durability**
 - **Problem:** How to store this information in tables (in an extensible and flexible way)?



The diagram shows three tables designed to store superhero data, but it highlights the challenge of mapping multiple objects with the same name to different sets of attributes. A large red question mark is centered between the tables, symbolizing the problem.

Power
Flight
Strength
Durability
Shoot energy

Name
Ms. Marvel
Ms. Marvel

Name	Power 1	Power 2
Ms. Marvel	Flight	Shoot energy
Ms. Marvel	Strength	Durability



I3. I Object Persistence

- This problem is commonly known as the **multi-valued data types problem**
 - A more popular version of this is the **bill-of-material problem**
 - Most solutions rely on introducing new surrogate keys

id	Name
1	Ms. Marvel
2	Ms. Marvel

herold	Power
1	Flight
2	Strength
2	Durability
1	Shoot energy





I3. I Object Persistence

- Providing object persistence is a complex task
 - Provide means to **create, read, update, and delete** persistent objects
 - Called the **CRUD** operations
 - For each of these operations, respect the object's **identity and structure**
 - Create explicit identities if necessary
 - Break object structure into relationships among entities

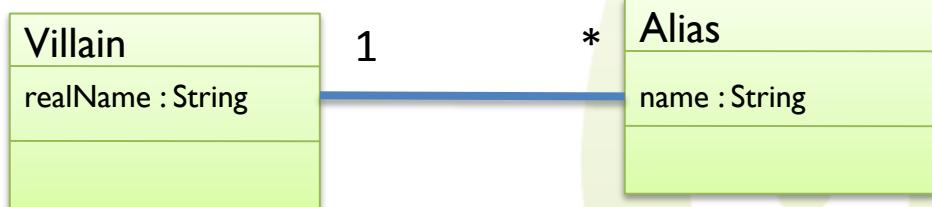




I3.2 Manual Persistence

Detour

- Of course, persistence can be achieved **manually** by using tools we already know
 - Example: **SQL** and **Java/JDBC**
 - However, this approach often is cumbersome and takes a lot of effort
- **Example:** Let's make a simple JavaBean persistent





I3.2 Manual Persistence

Detour

- **Beans**

Villain.java

```
public class Villain {  
    String realName;  
    Set<Alias> aliases;  
    // some getter/setter methods  
}
```

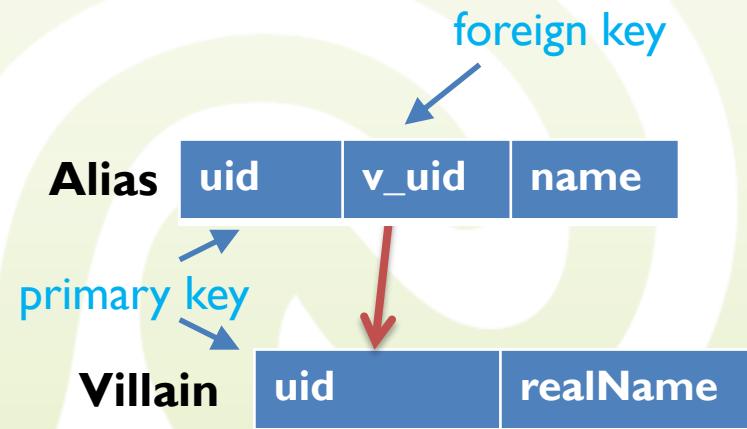
This is all the application actually needs!

Alias.java

```
public class Alias {  
    String name;  
    // some getter/setter  
}
```

- **Relational tables**

- Beans have **no primary keys** (names are not unique)
 - Generate a key somehow...





I3.2 Manual Persistence

Detour

CRUD create operation

Villain.java

```
int uid;  
public void createPersistent(Connection conn) throws SQLException {  
    Statement stmt = conn.createStatement();  
    uid = magicallyCreateAnUID(this);  
    stmt.executeUpdate("INSERT INTO villains (" + uid + ","  
        + realName + ")");  
    for (Alias alias : aliases) {  
        alias.createPersistent(this, conn); } }
```

Take care of identity

Take care of structure

Alias.java

```
int uid;  
public void createPersistent(Villain villain, Connection conn) throws  
    SQLException {  
    Statement stmt = conn.createStatement();  
    uid = magicallyCreateAnUID(this);  
    stmt.executeUpdate("INSERT INTO alias (" + uid + "," + villain.uid + ","  
        + name + ")"); }}
```



I3.2 Manual Persistence

Detour

- **Note:** The generated **UID** is usually not related to the “real” object identity (i.e. internal Java ID)
 - However, it serves the same purpose of **identifying** and connecting related **rows** in the RDBMS
 - Although the **UIDs** do not convey information that is needed within the application’s data model, it usually **must be carried along** to make the object persistence work





I3.2 Manual Persistence

Detour

- **CRUD read operation**
 - For reading an object, you need to **know its UID**
 - For accessing the full **object structure**, you have two options
 - **Eager loading:**
Loads the whole object structure as soon as the base object is requested
 - Good performance if you usually need the whole structure
 - **Lazy loading:**
Loads only the base object when requested, loads remaining parts of the object's structure when they are needed
 - Good performance if you usually only need parts of the structure



I3.2 Manual Persistence

Detour

CRUD read operation (lazy)

Villain.java

```
public static Villain readPersistent(int uid, Connection conn) throws  
SQLException {  
    Statement stmt = conn.createStatement();  
    Villain villain = new Villain();  
    ResultSet rs = stmt.executeQuery("SELECT * FROM villains WHERE  
        uid=" + uid);  
    rs.next();  
    villain.uid = uid;  
    villain.realName = rs.getString("realName");  
    villainAliases = null;  
    return villain;  
}
```

You need the UID, somehow...

Note:

Additional method besides the “normal” JavaBean methods



I3.2 Manual Persistence

Detour

CRUD read operation (lazy)

Alias.java

Note: This one is embedded in
a “normal” JavaBean accessor method!



```
public Set<Alias> getAliases() throws SQLException {  
    if (aliases == null) {  
        aliases = new HashSet<Alias>();  
        Statement stmt = ... // get statement from somewhere  
        ResultSet rs = stmt.executeQuery("SELECT * FROM alias WHERE  
            v_uid=" + uid);  
        while (rs.next()) {  
            Alias alias = new Alias();  
            alias.uid = rs.getInt("uid");  
            alias.name = rs.getString("name");  
            aliases.add(alias);  
        }  
    }  
    return aliases;  
}
```



I3.2 Manual Persistence

Detour

- **CRUD update:**
 - For updates, there are also various implementation approaches with different performance and safety properties
 - **Immediate updates:**
 - Directly persist updates after they occur
 - Augment each Bean's set method with additional JDBC calls that directly write the new value into the DBMS
 - Database immediately updated
 - **Performance might be very low**
 - Only few problems with transactional consistency



I3.2 Manual Persistence

Detour

– Explicit updates:

- **Persist updates on manual command**
- Just leave set methods as they are, but mark each object as “**dirty**” as long as it is used
- Provide an additional update method writing the whole Bean to the DB if it is dirty
 - **Traverse** the whole **object structure recursively** and also write all dirty, related objects
- Worst case:
Severe problems with transactional consistency



I3.2 Manual Persistence

Detour

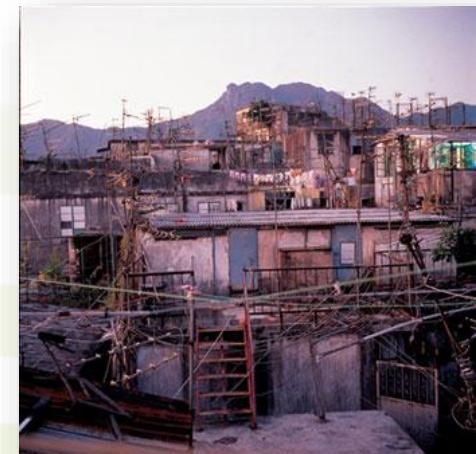
- **Summary:**
 - A considerable amount of **work is necessary**
 - There are **many fetching and update schemes** with different properties and performance impacts
 - **Lazy load vs. eager load**
 - **Immediate update vs. explicit update**
 - ...
 - There are several problems involved
 - **Generating IDs** and propagating them
 - Dealing with the **object structure**
 - Keeping **transactional consistency** in multi-user scenarios
 - Providing **sufficient performance**
 - How to perform **complex querying?**



I3.2 Manual Persistence

Detour

- Another major drawback is **code perturbation**
 - The code of the application data model is littered with code for dealing with persistence
 - Application models are “**dirty**” and bound to the chosen implementation
 - In this example, the initial Villain.java model grew from **26 lines** of code (class definition, attributes, getters & setters) to over **180 lines** of code (with simple (!) methods for persistence added)!





I 3.3 Persistence Frameworks

- **Persistence frameworks** come to rescue...
- Main idea:
 - Providing **persistence** for an application should be **as easy as possible**
 - Programmer should not spend too much time and code on these issues
 - Concentrate on more important things!
 - Provide automated support to problems of persistence
 - Persistence handling should be **transparent**
 - Just model your application data the way **you need it**, and not the way your DBMS needs it
 - **Keep your models clean!**





I 3.3 Persistence Frameworks

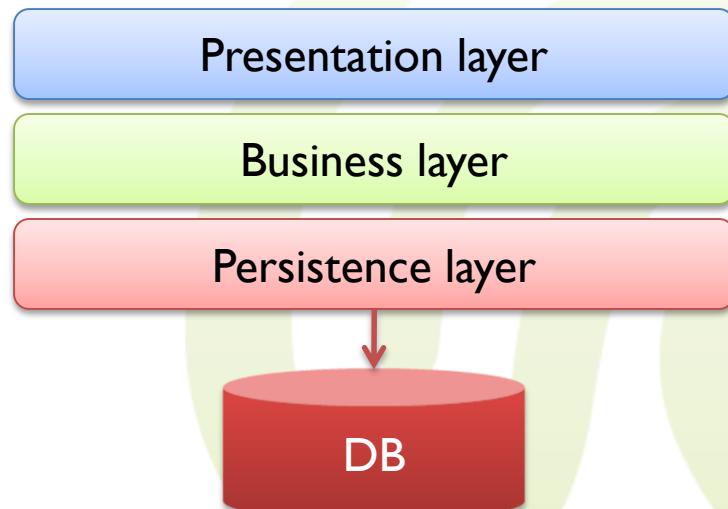
- During software development and maintenance, **complexity is the archenemy**
- A popular approach for reducing software complexity are **layered architectures**
 - Each layer has a **defined responsibility** and communicates with other layers using **clearly defined interfaces**
 - Usually, no code from one layer should spill into another
 - **Implementations** of a layer **may change** without affecting the others
 - Example: Web interface vs. Web service interface
 - Example: Persistence using a RDBMS vs. persistence using XML





I 3.3 Persistence Frameworks

- A layered architecture for applications using persistent object might look like this
 - **Presentation layer:**
Present data to the user and interact with him/her
 - **Business layer:**
Contains the domain data model and the business logic
 - **Persistence layer:**
Stores and retrieves objects of the domain data model





I 3.3 Persistence Frameworks

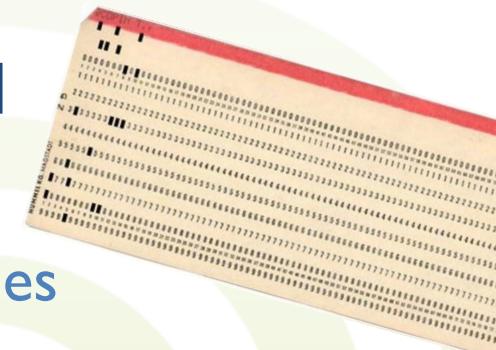
- **Persistence frameworks** aim at providing the **complete persistence layer**
 - In the best case, there are **only very few things** an application programmer needs to do
 - Domain model stays **clean**
 - Usually, information needed by the persistence layer is provided using **meta data**
 - However, you buy these features with a **performance penalty**



I 3.4 Generating IDs

Detour

- As you have seen, identifying data is important
- In the good old days of **punch cards** and **magnetic tapes**, identifying data was easy...
 - Each punchcard had a number that sequentially increased from card to card
 - Card No. 1, Card No. 2, ...
 - The same works for tapes and similar devices
- Using sequential identifiers also worked great within the **hierarchical data model** and **network data model**





I 3.4 Generating IDs

Detour

- The **relational data model** changed the view upon data
 - Data is organized as a **set of tuples**
 - Each tuple is identified by its **value**
 - Explicit identification is not mandatory!
 - A small subset of attributes is selected as **primary key** for easier **identification** and **reference**
- The model was designed to be used with **natural keys**
 - The **key values** are part of the data domain
 - Each student already has a unique matriculation number, so just use it as key
 - Imagine, you have a set of weather stations. Each reading can be uniquely identified by its time, type, and station



I 3.4 Generating IDs

Detour

- But what happens if your data does not provide a useful natural key?
 - Create a synthetic key yourself (so-called **surrogate keys**)
- When performing **object-relational mapping**, this problem frequently arises since each object always needs a **unique ID**
 - “Real” ID usually hidden within the OS or VM
 - Surrogate Keys are needed!



I 3.4 Generating IDs

Detour

- How to create surrogate keys?

There are several approaches:

- **Sequence keys**

- Full sequences keys
- Hi-low sequences

- **UUIDs**

- Time & device-based UUIDs
- Hash-based UUIDs
- Random UUIDs





I 3.4 Sequences

Detour

- **Full sequenced keys**

- A **central authority** (usually the DBMS) provides all keys
- The keys follow a strict sequence
 - Example: 1, 2, 3, 4, 5,
 - Usually, there is **one sequence per table**

- **Pros**

- Keys are short and easy to debug
 - Example: ... WHERE id = 12
- Easy to handle and storage efficient, indexes can grow with less need for restructuring (ordered inserts!)

- **Cons**

- Key is provided centrally by the DBMS → **Bottleneck!**
- For each single key, you need to connect to the DB to retrieve it
→ **Network traffic**





I 3.4 Sequences

Detour

- Most databases supporting this feature use **sequence tables**
 - For each sequence to be generated there is a single-column/single-row table containing an integer that is continuously incremented
- Define a table using an auto-incremented key
 - **Oracle:** The sequence needs to be specified explicitly
 - Use the CREATE SEQUENCE statement
 - When you want to insert a new row with a surrogate key, call the sequence using `<seq_name>.nextval` within the INSERT statement



I 3.4 Sequences

Detour

- **MySQL:** Declare it in the column definition of the CREATE TABLE statement
 - colName dataType NOT NULL AUTO_INCREMENT
- **DB2:** Like MySQL
 - colName dataType NOT NULL GENERATED AS IDENTITY
- **Example DB2:**
 - CREATE TABLE hero (id INTEGER NOT NULL GENERATED AS IDENTITY, name VARCHAR(255))
 - INSERT INTO hero (name) VALUES ('The Hulk')
 - Note that you do not know what key the DB assigned to your new row!





I 3.4 Hi-Low Keys

Detour

- **Hi-low keys**
 - **Central authority** (DBMS) and **application** share responsibility for key creation
 - **Idea:** Key is made up of two parts
 - **High:** Provided by the DBMS
 - **Low:** Provided by the application
 - Every time an application connects to the DBMS, it receives a unique **high** value for that session
 - Usually derived using some sequence within the DBMS
 - The application creates its own sequence for that session and increments it for each needed **low** value
 - Key is created by **concatenating** the **high** and **low** part





– Pros

- Just one DBMS access per user/application session
→ reduces network traffic, rarely any bottlenecks
- Keys are still of manageable size
- **Full key is immediately known to the application**
 - Very important for OR-mapping!

– Cons

- Application is responsible for finally creating the key
(no plain easy auto-incrementing columns)
- Still communication with the DBMS necessary



I 3.4 UUIDs

Detour

- **Universally unique identifiers (UUID)**
 - Standard provided by the Open Software Foundation (**OSF**) for unique surrogate keys and IDs
 - Later: **IETF RFC 4122** and **ISO/IEC 9834-8:2005**
 - **Base idea:**
Generate unique keys **without any central control**
 - Most popular implementations
 - Microsoft Globally Unique Identifiers (**GUID**)
 - **Ext2/Ext3 file system** identifiers





13.4 UUIDs

Detour

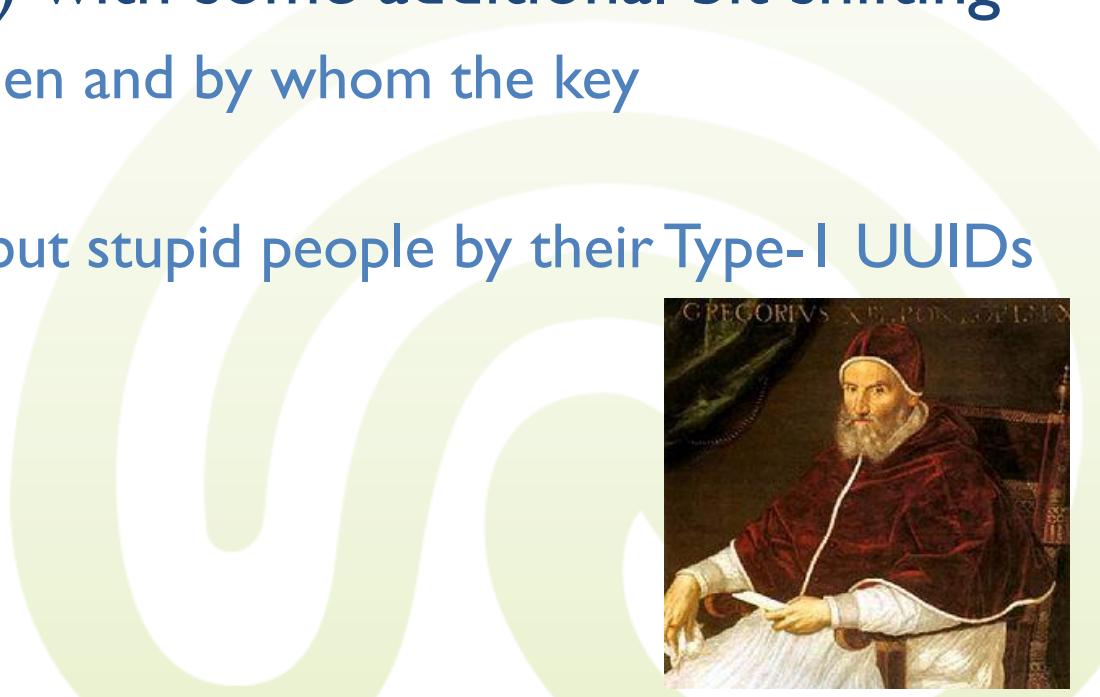
- A UUID is a **128-bit number**
- Usually represented by **32 hex digits** in 5 groups of lengths 8, 4, 4, 4, and 12
 - e.g. **4e84890a-5f12-42fd-b1fe-0d32afb1d9d8**
- There are 5 defined types of UUID algorithms
 - You can identify the used algorithm by inspecting the **first digit** of the third hex block (**in red**)



I 3.4 UUIDs

Detour

- **Type-I UUIDs**
 - UUID is a concatenation of the **MAC address** of the generator host and the **number of 100 ns intervals since February 24, 1582** (introduction of the Gregorian calendar) with some additional bit shifting
 - You can identify when and by whom the key has been generated
 - You can catch evil, but stupid people by their Type-I UUIDs





I 3.4 UUIDs

Detour

– The Melissa worm

- Self-replicating macro virus using vulnerabilities in Outlook
- Melissa virus shut down large parts of the internet in March 1999
- However, the code contained some UUIDs left by its author...
- By reversing the UUID, its creator David L. Smith could be **back-traced and arrested**

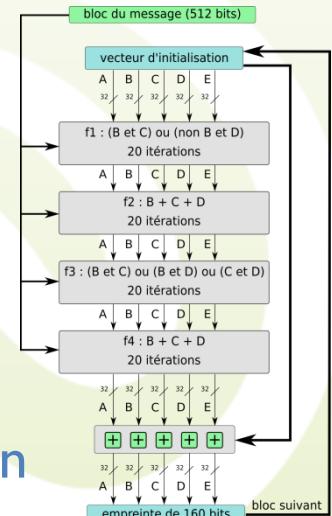




13.4 UUIDs

Detour

- **Type-2 UUID:**
 - Like Type-1, but parts of the timestamp are replaced by the POSIX UID domain and user UID
 - Privacy issues even more severe
- **Type-3 UUID:**
 - Everything but the type digit is generated by a MD5 hash function
 - Input is usually a URL, object identifier, etc...
 - Problems in generating truly unique IDs
- **Type-5 UUID:**
 - Like type 3, but uses SHA-1 hashing
 - SHA-1 is a stronger and more efficient hash function

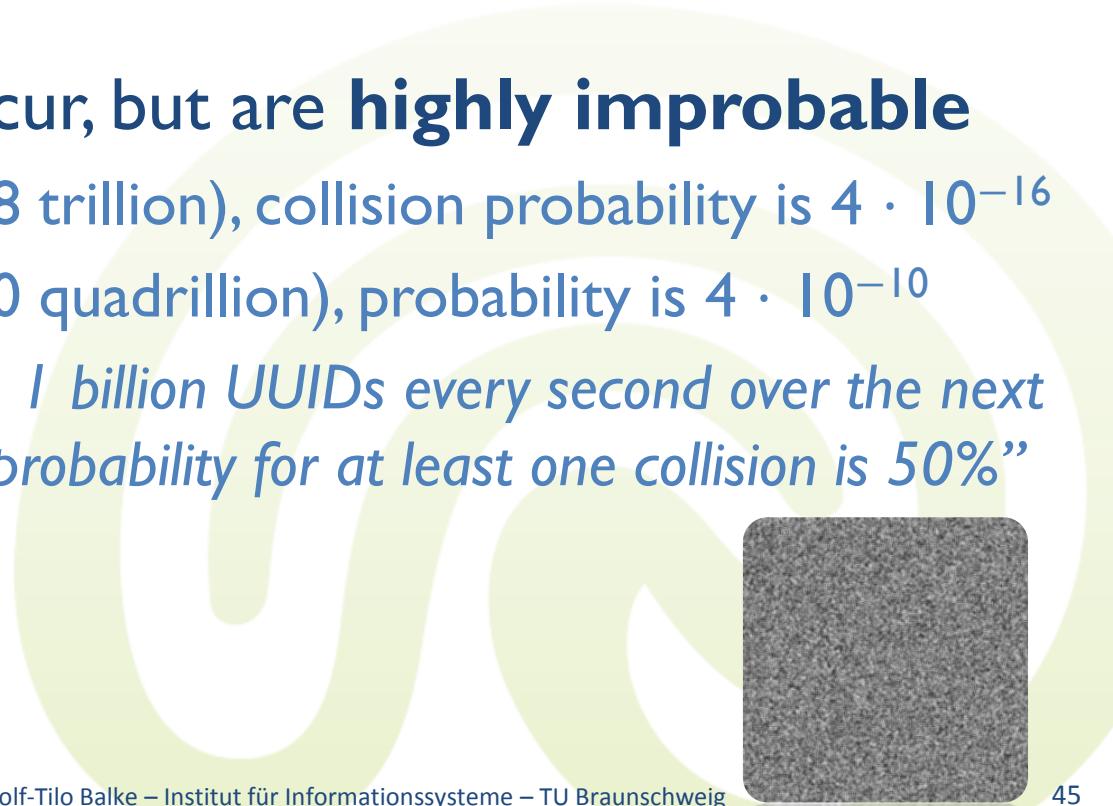




I 3.4 UUIDs

Detour

- **Type-4 UUID**
 - All digits but the type digit are **randomly generated**
 - A cryptographically secure random number generator is needed
 - **Collisions** may occur, but are **highly improbable**
 - For 2^{36} UUIDs (~68 trillion), collision probability is $4 \cdot 10^{-16}$
 - For 2^{46} UUIDs (~70 quadrillion), probability is $4 \cdot 10^{-10}$
 - “When you generate 1 billion UUIDs every second over the next 100 years, then the probability for at least one collision is 50%”





I 3.4 UUIDs

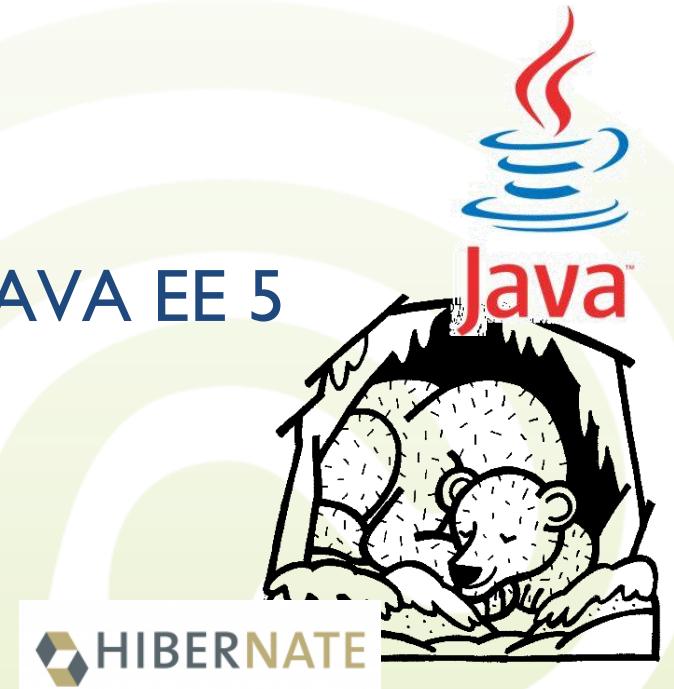
Detour

- Is using UUIDs a good idea?
 - Depends...
 - **Good**
 - UUIDs can be generated very easily without central control and without support from the DBMS
 - When using UUIDS, you can easily integrate data from different data sources
 - No key collisions
 - Performance of queries not affected if used correctly
 - **Bad**
 - UUIDs are horrible to debug
 - `SELECT ... FROM ... WHERE id = 8ac7fb3d4f4047419c7f7d22d1802fe3`
 - Usually, more storage space is needed



I3.5 Persistence Frameworks

- Currently, there are many products available
- We will focus on
 - **JPA (Java Persistence API):** Default API provided by SUN JAVA EE 5
 - **Hibernate:** Most popular Java persistence framework





I 3.5 JPA

- Before the introduction of JPA (**Java Persistence API**), there were a multitude of persistence frameworks
 - Hibernate
 - Apache OJB
 - Apache Castor
 - JPOX
 - XORM
 - Persist
 - Oracle Toplink
 - ...
- Each of these frameworks had their own **proprietary API** and **metadata format**



Castor

XORM

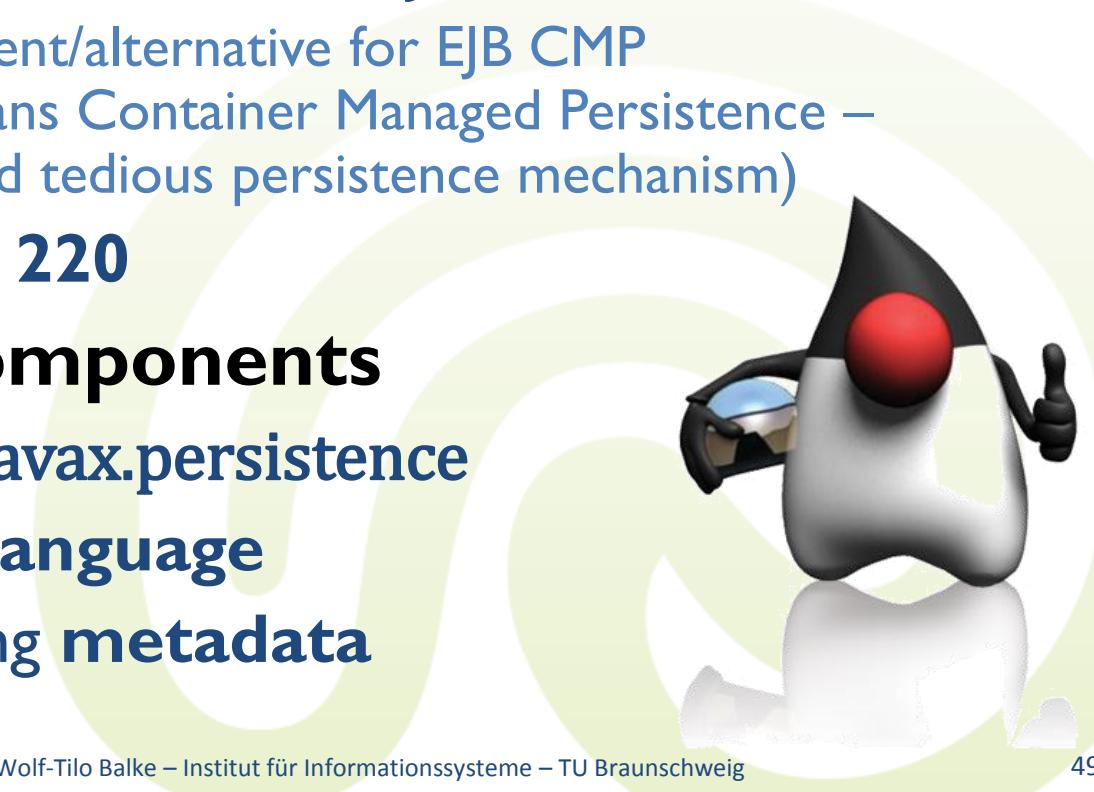
JPOX
Java Persistent Objects

HIBERNATE



I 3.5 JPA

- JPA provides a common interface for all relational persistence frameworks written in Java
 - Released in **May 2006** by Sun Microsystems
 - Part of **EJB 3.0 specification** of Java EE 5.0 standard
 - Unofficial replacement/alternative for EJB CMP (Enterprise Java Beans Container Managed Persistence – a really complex and tedious persistence mechanism)
 - Also known as **JSR 220**
- Consist of three **components**
 - The actual **API** in `javax.persistence`
 - The **JPQL query language**
 - Facilities for handling **metadata**





I 3.5 JPA

- JPA is a **standardized API**
 - Sun was able to convince most persistence framework vendors to adopt the API
 - If you want to use it, you have to get a compliant framework
 - The current reference implementation is Oracle TopLink Essential
 - Now open source and supported by Eclipse, called **EclipseLink**
http://www.eclipse.org/projects/project_summary.php?projectid=rt.eclipselink
- **Resources**
 - Overview:
 - <http://java.sun.com/javaee/technologies/persistence.jsp>
 - API doc:
 - <http://java.sun.com/javaee/5/docs/api/javax/persistence/package-summary.html>



I 3.5 JPA

- **Main idea:**
 - Use of **POJOs** as the application's domain data model
 - **POJO: Plain old Java objects**, that is, JavaBeans without any complex stuff in them
- **Metadata** describes how the POJOs are mapped to the relational DB
 - Metadata as **annotations**: The POJO is enriched using JSR 175 annotations (Java 5 and beyond)
 - Very easy, but domain model needs to be directly annotated
 - Metadata as **XML**: In addition to the POJO domain model, there are XML files describing the OR mapping
 - More complex, more effort needed
 - Strong separation of business and persistence layer



I 3.5 JPA

- Using **annotations**,
the development **workflow** looks like this:
 1. Annotate all **persistent entities**
 - That is: Those classes that are supposed to be stored in an own table
 2. Annotate either all **attributes** or all **getter/setters** of your persistent entities
 - Define what attributes contain **UIDs**
 - Define **relationships**
 - Define special **persistence behavior**
 3. Provide an **XML document** describing your **persistence units**
 - Which entities should end up in which database?
 4. Use the **JPA EntityManager** to create/read/update/delete your persistent objects



I3.5 JPA

Detour

- **Beans**

Villain.java

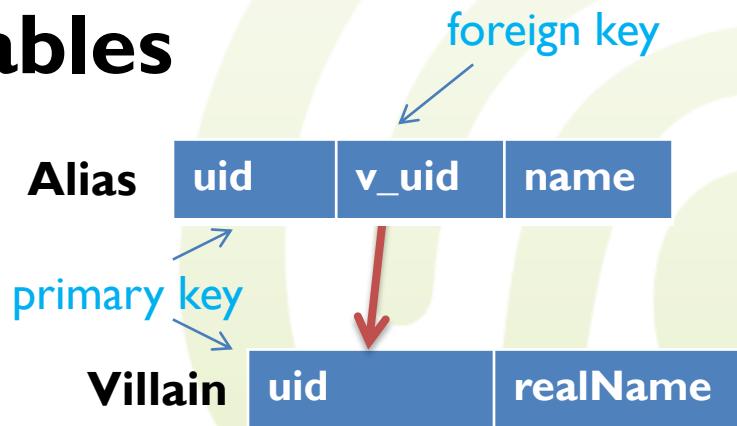
```
public class Villain {  
    String realName;  
    Set<Alias> aliases;  
    //...put additional getter/setters here  
}
```

POJOs

Alias.java

```
public class Alias {  
    String name;  
    //...put getter/setters here  
}
```

- **Relational tables**





Annotate persistent entities

Villain.java

```
@Entity  
public class Villain {  
    String realName;  
    Set<Alias> aliases;  
    //...put additional getter/setters here  
}
```

Alias.java

```
@Entity  
public class Alias {  
    String name;  
    //... put getter/setters here  
}
```

Add and annotate UID columns

Villain.java

```
@Entity  
public class Villain {  
    @Id  
    @GeneratedValue(strategy =  
        GenerationType.SEQUENCE)  
    int uid;  
    String realName;  
    Set<Alias> aliases;  
}
```

Alias.java

```
@Entity  
public class Alias {  
    @Id  
    int uid;  
    String name;  
}
```



I 3.5 JPA

Detour

Annotate attributes and relationships

Villain.java

```
@Entity  
public class Villain {  
    @Id  
    @GeneratedValue(strategy =  
        GenerationType.SEQUENCE)  
    int uid;  
    String realName;  
    @OneToMany(fetch = FetchType.LAZY)  
    Set<Alias> aliases;  
}
```

Alias.java

```
@Entity  
public class Alias {  
    @Id  
    int uid;  
    String name;  
    @Transient  
    String nonsense;  
}
```

“Ignore this attribute”

“Use lazy loading”

- All simple attributes are **persisted by default** unless annotated with `@Transient`
- All collections need an annotation defining the relationship
(one-to-one, one-to-many, many-to-one, many-to-many)
 - Provide the desired name of the **foreign key column**
 - Optionally, also provide **loading and updating strategies**



- Usually, the JPA implementation **automatically designs the relational schema** as it likes
- That is fine, if **only one application** uses the persistent entities or all applications use the same JPA meta data
- However, if you have to **use an existing relational schema** or other existing applications have to use the tables, it might become tricky
 - Detailed control of **object-relational mapping** necessary
 - Use **advanced annotations** for this task





- **Advanced annotations** cover the following:
 - **Indexing**
 - **Structural design**
 - Naming, embedding vs. linking, splitting and merging of object structures, etc.
 - **IDs**
 - Natural ids, UUIDS, sequence-based ids, etc.
 - **Custom data types**
 - **Update and delete behavior**
 - Cascading, Restricting, etc.
 - **Constraints**
 - **Computed values**
 - **Sorting**
 - **Locking**
 - **Caching**
 - **Filtering**
 - etc.





I 3.5 JPA

Detour

- Also, all this **mapping meta data** can be provided by an **XML** file instead of annotations
- Beside mapping meta data, you need to define your **DBMS connection**
 - Use **persistence.xml** file in the **META-INF** folder

```
<persistence>
    <persistence-unit name="villains">
        <provider>YOUR JPA IMPLEMENTATION NAME GOES HERE</provider>
        <class>Villains</class>
        <class>Alias</class>
        <properties>
            <!-- Properties for your JDBC driver, username, password, URL, etc. --!>
        </properties>
    </persistence-unit name="villains">
</persistence>
```



I 3.5 JPA

Detour

- After persistence properties are defined, the **JPQL query** language may be used to retrieve objects
 - Similar to SQL, but uses persistent entities and their attributes instead of tables
- You can query and load/store/change objects using the JPA Entity Manager
 - Entity Manager is created using a persistence unit definition



I 3.5 JPA

Detour

```
factory = Persistence.createEntityManagerFactory("villains");
EntityManager em = factory.createEntityManager();
// query
Query q = em.createQuery("SELECT v FROM villains AS v WHERE
    v.realName='Galan'");
List<Villain> villain = (List<Villain>) q.getResultList();
//create a new villain
em.getTransaction().begin();
Villain newVillain = new Villain();
newVillain.setRealName("Gala");
newVillain.addAlias(new Alias("Galactus"));
newVillain.addAlias(new Alias("Eater-of-Worlds"));
em.persist(newVillain);
// change Galactus
newVillain.setRealName("Galan");
em.persist(newVillain);
//
em.getTransaction().commit();           em.close();
```





I 3.6 Hibernate

- Hibernate is the **most popular** Java persistence framework
 - Initially developed by **Gavin King**
 - **Open source**
 - Became really popular in 2003
 - However, it is not really known when the first version has been released...
 - Like most early frameworks, Hibernate provided proprietary mapping **meta data definitions**, a **query language** (HQL), and an **entity manager**
 - Supports **JPA** since version 3.0
 - Acquired by **JBoss Inc.**





I 3.7 Object Databases

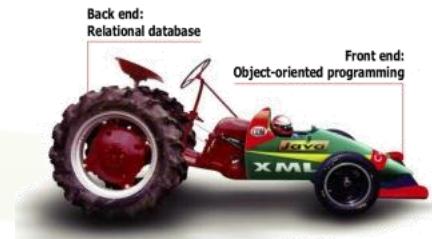
- **Objects databases (ODBMS)** came up for the first time around the **mid-80s**
 - Driven by the increasing popularity of **object-oriented programming languages**
 - **Promise:** Get rid of the annoying object-relational **impedance mismatch**
 - **Store objects in all their complexity, do not match them to tables, etc.**
 - In theory, extremely high performance possible
 - Most programmers loved that idea





13.7 Object Databases

- Object databases directly **interacted** with the **programming language**, thus developing applications should become very easy
- The first wave of commercial products in the mid-90s was **extremely hyped**
 - Gemstone, O2, Versant, Jasmine, Matisse, Objectivity, ObjectStore, Caché, etc.
- Standard committees proposed various **ODBMS standards**
 - Object Database Management Group (ODMG)
 - Object-extensions in SQL-99
 - ...



Objectivity.





I 3.7 Object Databases

- Unfortunately, most ODBMS **spectacularly failed**
 - Products unfinished and unpolished
 - Crappy performance due to misuse
 - Obscure and highly proprietary APIs, standards, and query languages
 - Integration with legacy systems was very hard
 - Most vendors went out of business...
 - ODBMS closed in 2001





I 3.7 Object Databases

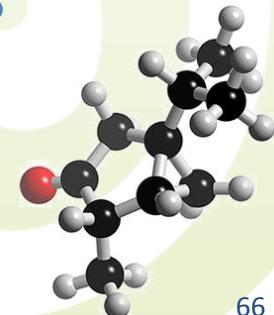
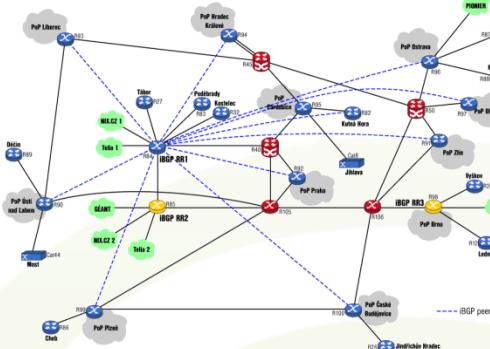


- However, in the last few years, ODBMS gained momentum again
 - Software is more and more developed by smaller companies in **smaller projects** within shorter time
 - **Agile programming** on the rise
 - Object-relational mapping becomes **increasingly expensive**
 - **Consumes too much time** in development and production
 - Newer ODBMS mainly focus on persisting the object state and don't include the behavioral aspects



13.7 Object Databases

- New generation of ODBMS does not aim to replace RDBMS, but provide alternatives for certain areas
 - **Embedded databases**
 - **Mobile databases**
 - **Real-time systems**
 - Telecommunications
 - **Scientific databases** (and all other databases storing highly **complex data structures**)
 - Databases in **Physics, biology, chemistry, etc.**
 - Currently, the **world's largest database** is a ODBMS (several petabyte, produced by the Stanford Linear Accelerator Experiment)





I 3.7 Object Databases

- Objects databases are **not a great choice** when
 - Your data needs to **be accessed by** other applications which are not specially devolved for ODBMS
 - Today's ODBMS are still quite proprietary
 - Data is just stored “somehow” (**encapsulation**)
 - Your data has **tabular nature**
 - True for most accounting data
 - In short: Everything you could (in theory) easily do in a huge Excel sheet does not belong into an ODBMS
 - You demand that your DBMS computes **statistics** or produces complicated **aggregations**
 - ODBMS are all-or-nothing: You will get only those objects you put into it earlier. Nothing else.



13.7 Object Databases

“... we'll simply observe that object databases haven't been widely adopted and that it doesn't appear likely that they will be in the near future.”

Gavin King, 2007

- This might or might not be true, **but** object databases are still a very interesting and maybe even a very good alternative to consider
 - **It depends** on your task and your data if they are a good idea
 - If your use case is suitable for a ODBMS, you can achieve spectacular performance paired with easy development and management



I 3.7 db4objects

Detour

- **db4objects** is native Java object database
 - Developed since 2000 as open source sponsored by Veritas and Sun
 - <http://www.db4o.com>
 - Sold to Versant in 2008
 - Planned target market: Mobile user data management
- What does “native” mean?
 - Written in **Java**, only using Java and nothing else
 - **Object data model** (no mapping or whatsoever)
 - **No embedded query languages**
 - Queries are also performed via Java objects





I 3.7 db4objects

Detour

- Handling of Java POJOs is extremely easy thanks to the native support
 - No modifications or meta data necessary
 - No unique identifiers necessary

Villain.java

```
public class Villain {  
    String realName;  
    Set<Alias> aliases;  
    //...put additional getter/setters here  
}
```

A real POJO!

Creating and persisting a villain

```
ObjectContainer db =  
    Db4o.openFile(dbFile);  
    Villain v = new Villain("Norrin Radd");  
    v.addAlias(new Alias("Silversurfer"));  
    v.addAlias(new Alias("The Herald"));  
    db.store(v);
```



I 3.7 db4objects

Detour

- Provides two query “languages”
- **QBE** (Query by example)
 - Just provide a template resembling the type of object you like to have
- **SODA**
 - Simple Object Database Access
 - Query by criteria
 - Query is an object structure of different query objects





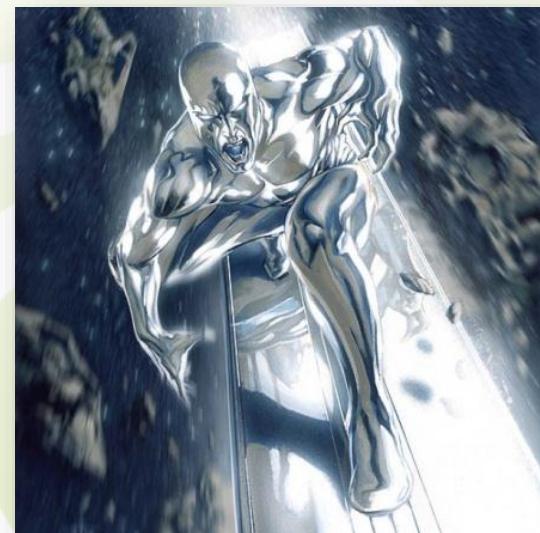
I 3.7 db4objects

Detour

- **Example:**
 - Query for all villains having the alias “Silversurfer”
 - **QBE**
 - Create a prototypical object matching the query description
 - All uninteresting attributes/relationships are left out or set to null

```
Villain proto = new Villain(null);  
proto.addAlias(new Alias("Silversurfer"));
```

```
List<Villain> result= (List<Villain>)  
db.queryByExample(proto);
```





I 3.7 db4objects

Detour

- **Example:**

- Query for all villains having the alias “Galactus”

- **SODA**

- Create a query expression by composing query objects
- Provide all conditions that have to be fulfilled by result objects

```
Query q = db.query();
q.constrain(Villain.class);
q.descend("aliases").descend("name").
    constrain("Galactus").equal();
result = (List<Villain>) q.execute();
```





I 3.7 db4objects

Detour

- Benchmarks
 - Uses PolePosition open source DB benchmark
 - <http://polepos.sourceforge.net/>
 - Note that **you are not allowed to publish benchmarks for most commercial DBMS...**
 - However, they do not perform well compared to the very strong MySQL
 - Barcelona: **Transactions** (100 selects, 30000 objects with 5 levels)



Barcelona Benchmarks	read	write	query	delete
Native/db4o 6.4	1.0	1.0	1.0	1.0
Hibernate/hsqldb	15.8	3.7	2,583.1	4.3
Hibernate/mysql	48.0	26.1	14.4	26.9
JDBC/MySQL	40.8	19.5	9.3	15.8
JDBC/JavaDB	27,843.7	20.5	47,993.1	17.7
JDBC/HSQldb*	1.9	1.1	2,554.4	0.5
JDBC/SQLite	8.8	519.1	1.1	362.1



* JDBC/HSQldb not ACID transaction safe



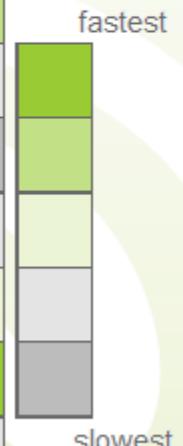
I 3.7 db4objects

Detour

- Benchmark

- Melbourne - writes 100,000 unstructured flat objects of one kind in bulk mode
- Bahrain - writes 30,000 simple flat objects individually
- Sepang - writes an object tree of depth 14
- Imola - retrieves 300,000 objects by native ID

Other Benchmarks	Melbourne	Bahrain	Sepang	Imola
Native/db4o 6.4	1.0	1.0	1.0	1.0
Hibernate/hsqldb	3.1	1.3	2.7	5.9
Hibernate/mysql	7.4	4.1	10.9	32.2
JDBC/MySQL	3.8	2.6	7.2	18.3
JDBC/JavaDB	1.8	2.1	407.0	5.7
JDBC/HSQLDB*	0.2	0.2	1.0	0.4
JDBC/SQLite	100.3	62.5	154.8	5.8



* JDBC/HSQLDB not ACID transaction safe



Next Lecture

- Active databases
 - Integrity constraints
 - Triggers
 - UDFs and stored procedures
- Basic security
 - Access control
 - SQL injection

